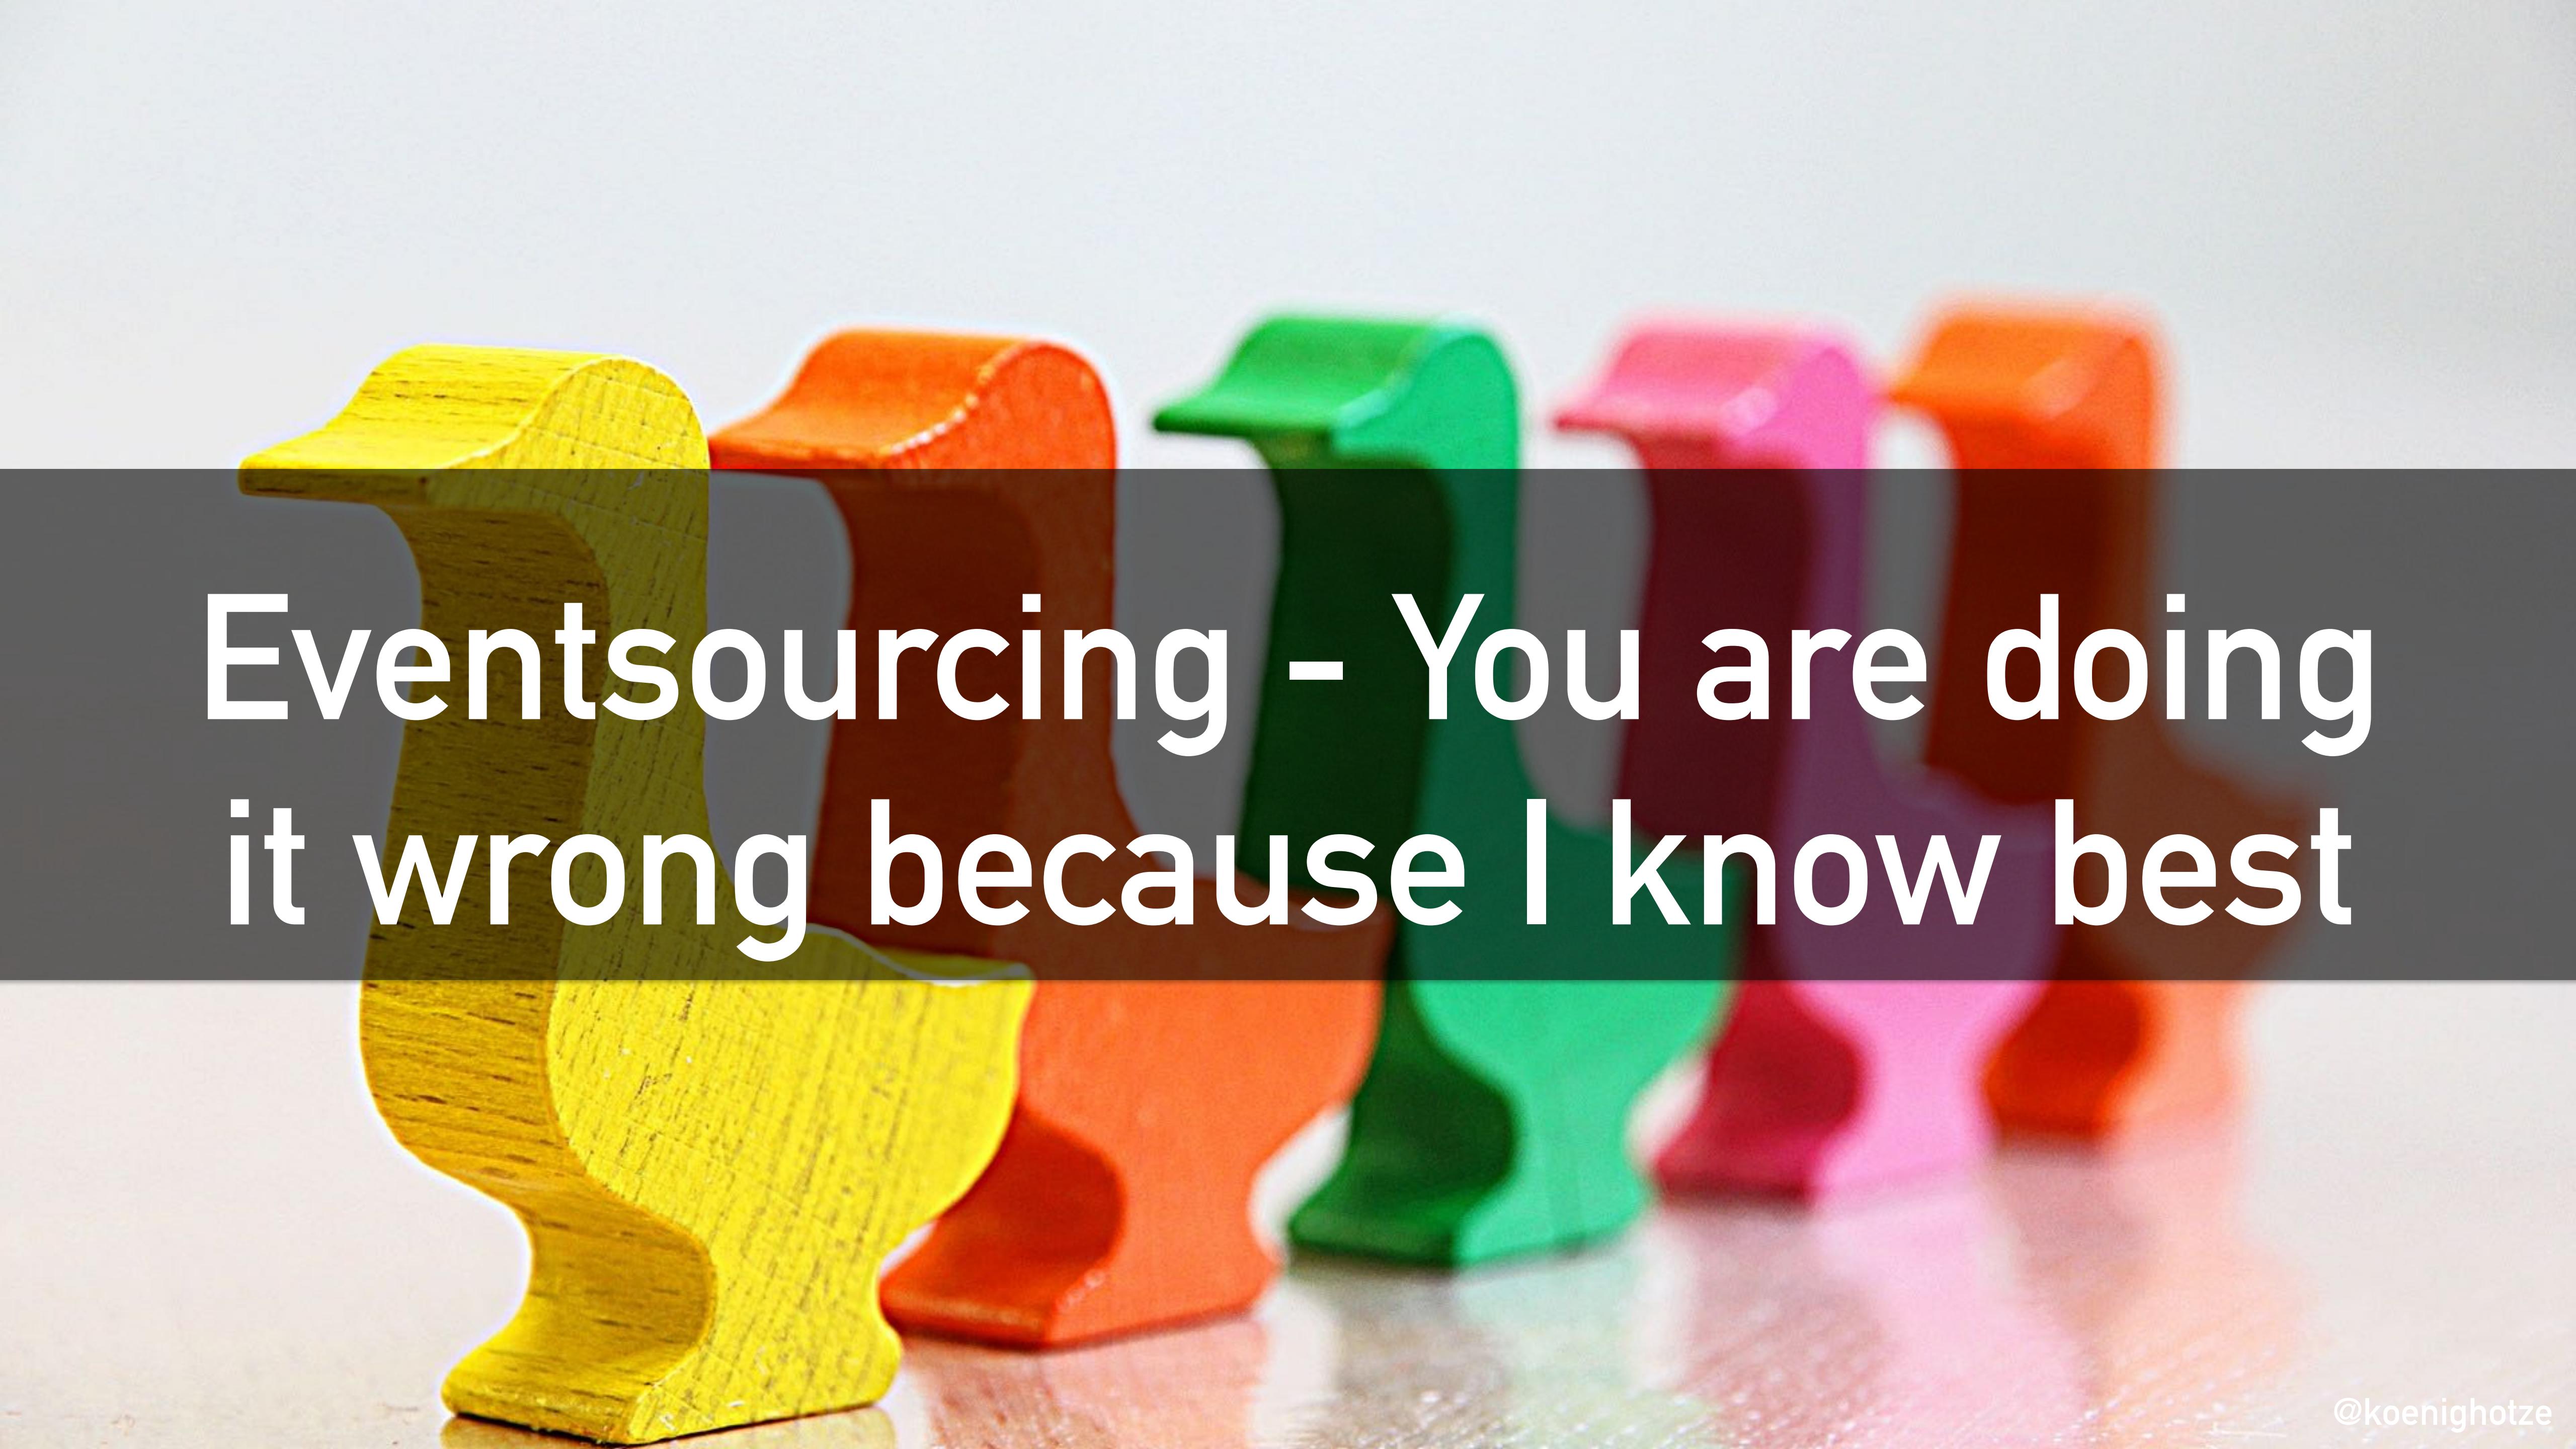
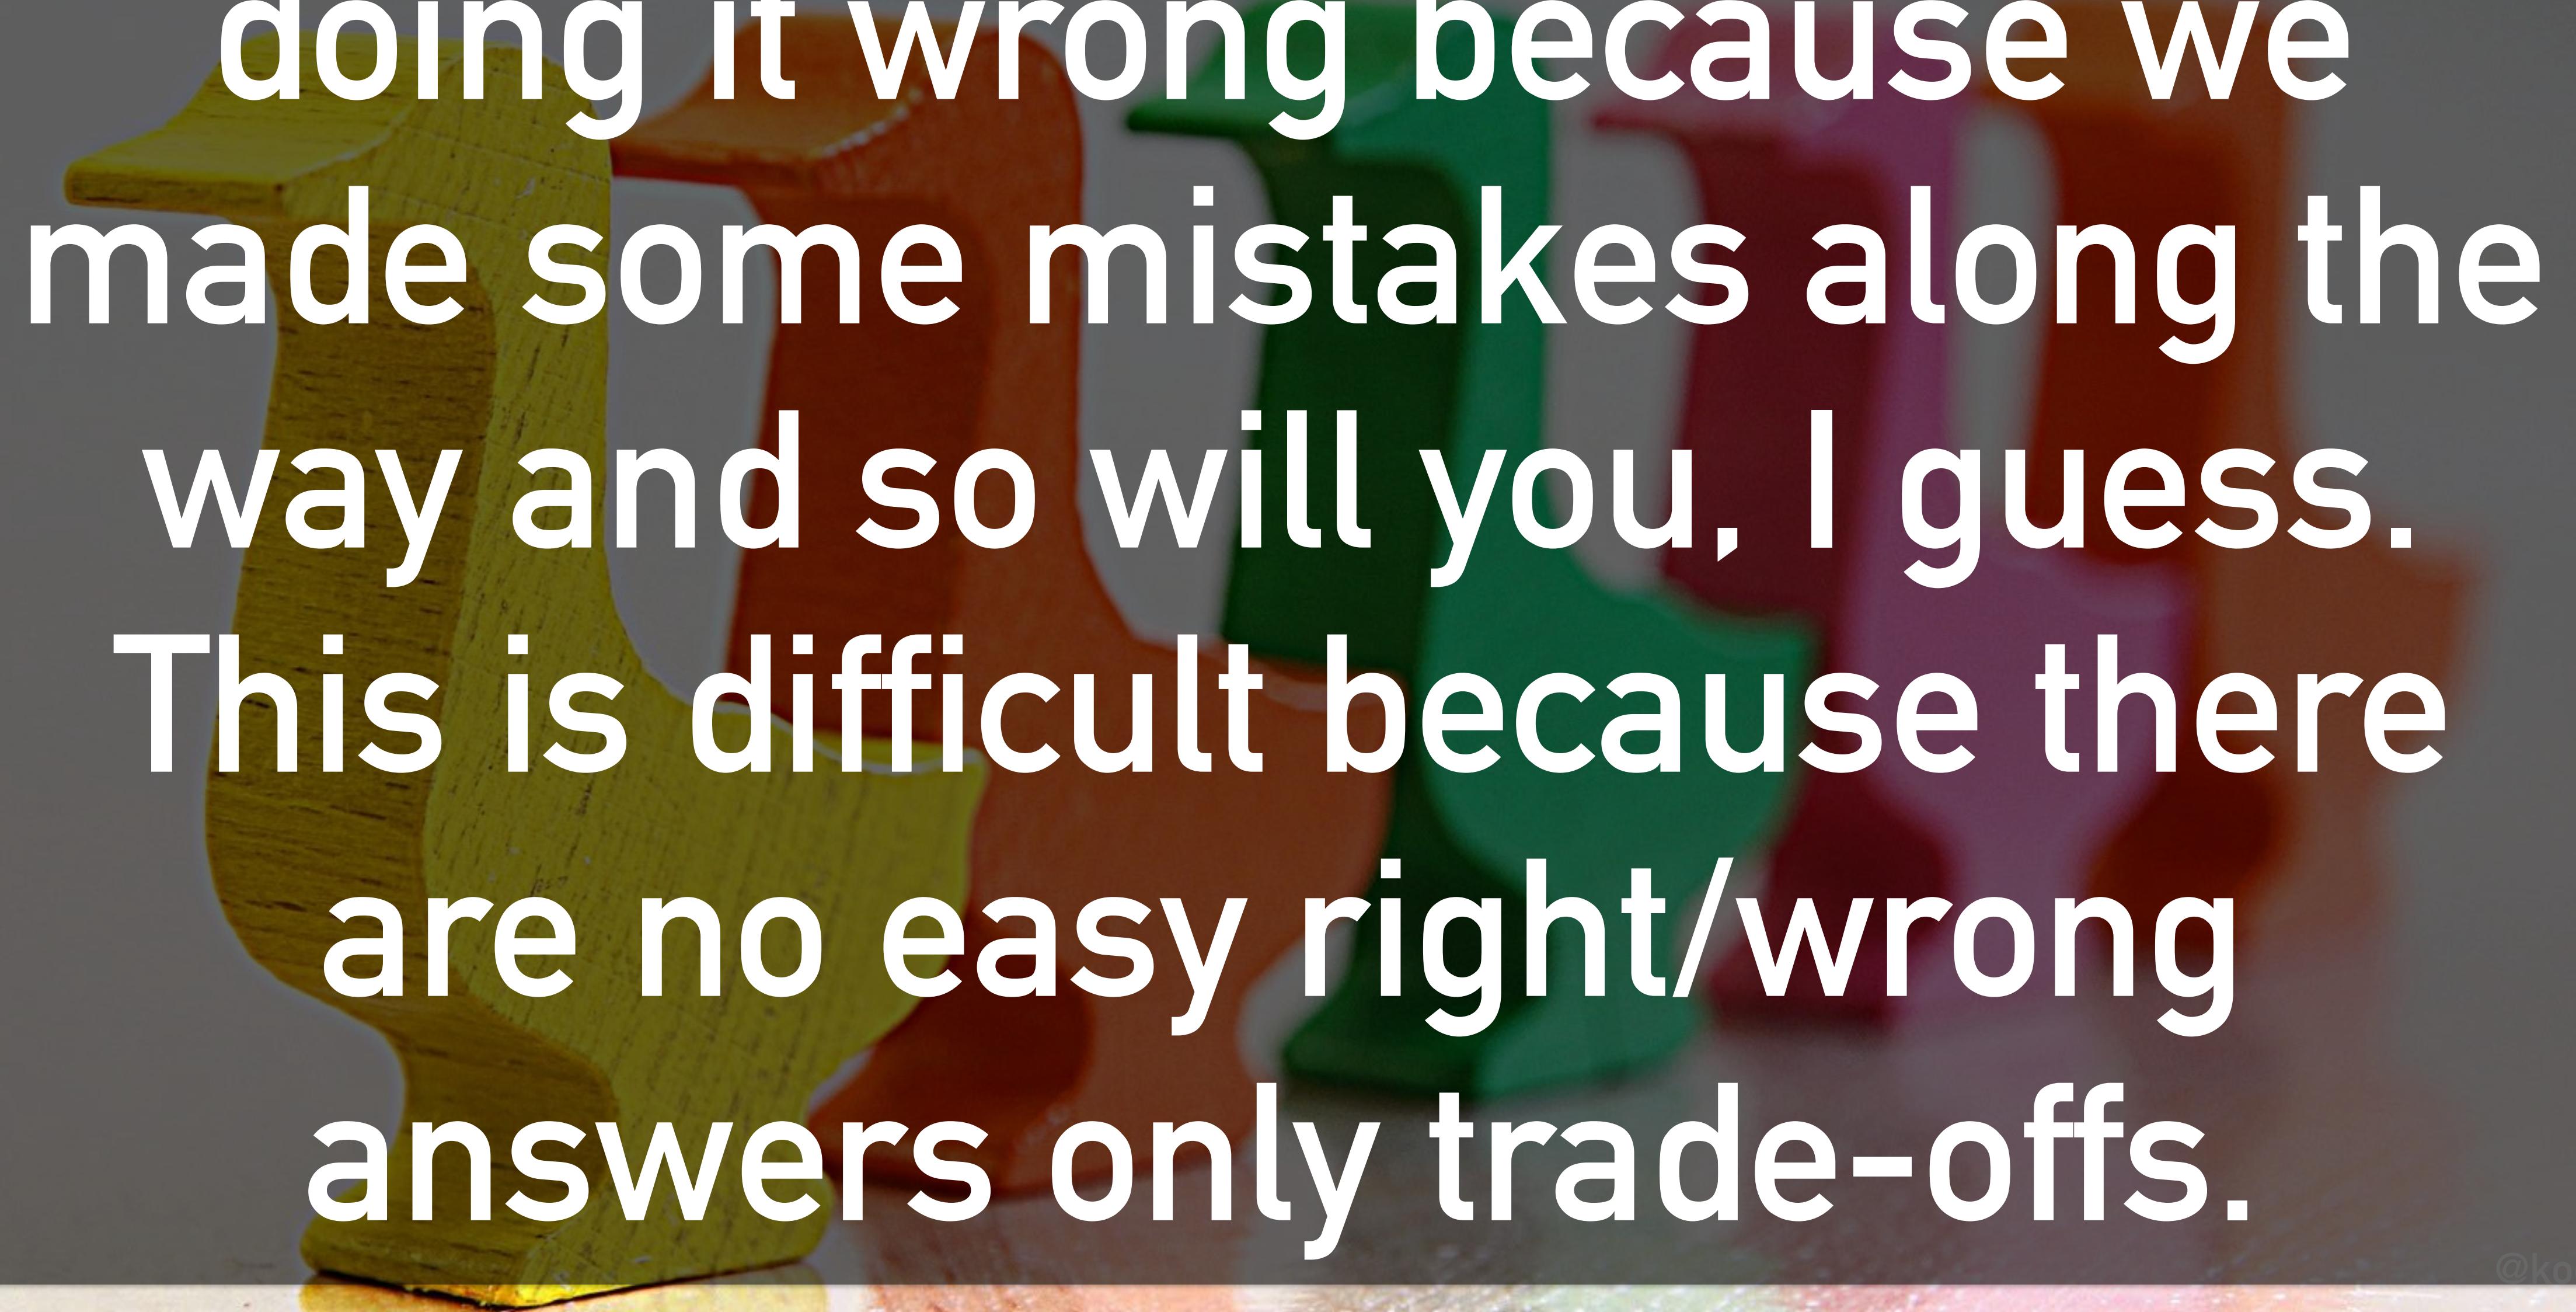




Eventsourcing - You are
doing it wrong



Eventsourcing - You are doing
it wrong because I know best



Eventsourcing - You are maybe doing it wrong because we made some mistakes along the way and so will you, I guess. This is difficult because there are no easy right/wrong answers only trade-offs.



Eventsourcing - You are
Probably doing it wrong

A black and white photograph of a canal scene. In the foreground, a portion of a boat is visible on the left. The middle ground shows a canal with a bridge in the background. On the right, there is a large, multi-story brick apartment building with many windows and balconies. The sky is overcast with dramatic clouds.

David Schmitz

@koenighotze

Senacor Technologies

Are **YOU** building microservices?
Are **YOU** doing Domain Driven Design?
Are **YOU** applying eventsourcing?
Are **YOU** using Kafka as an eventstore?

how
who
what
when
why
where



Typical misconceptions

Patterns “we” found useful

Traps to avoid

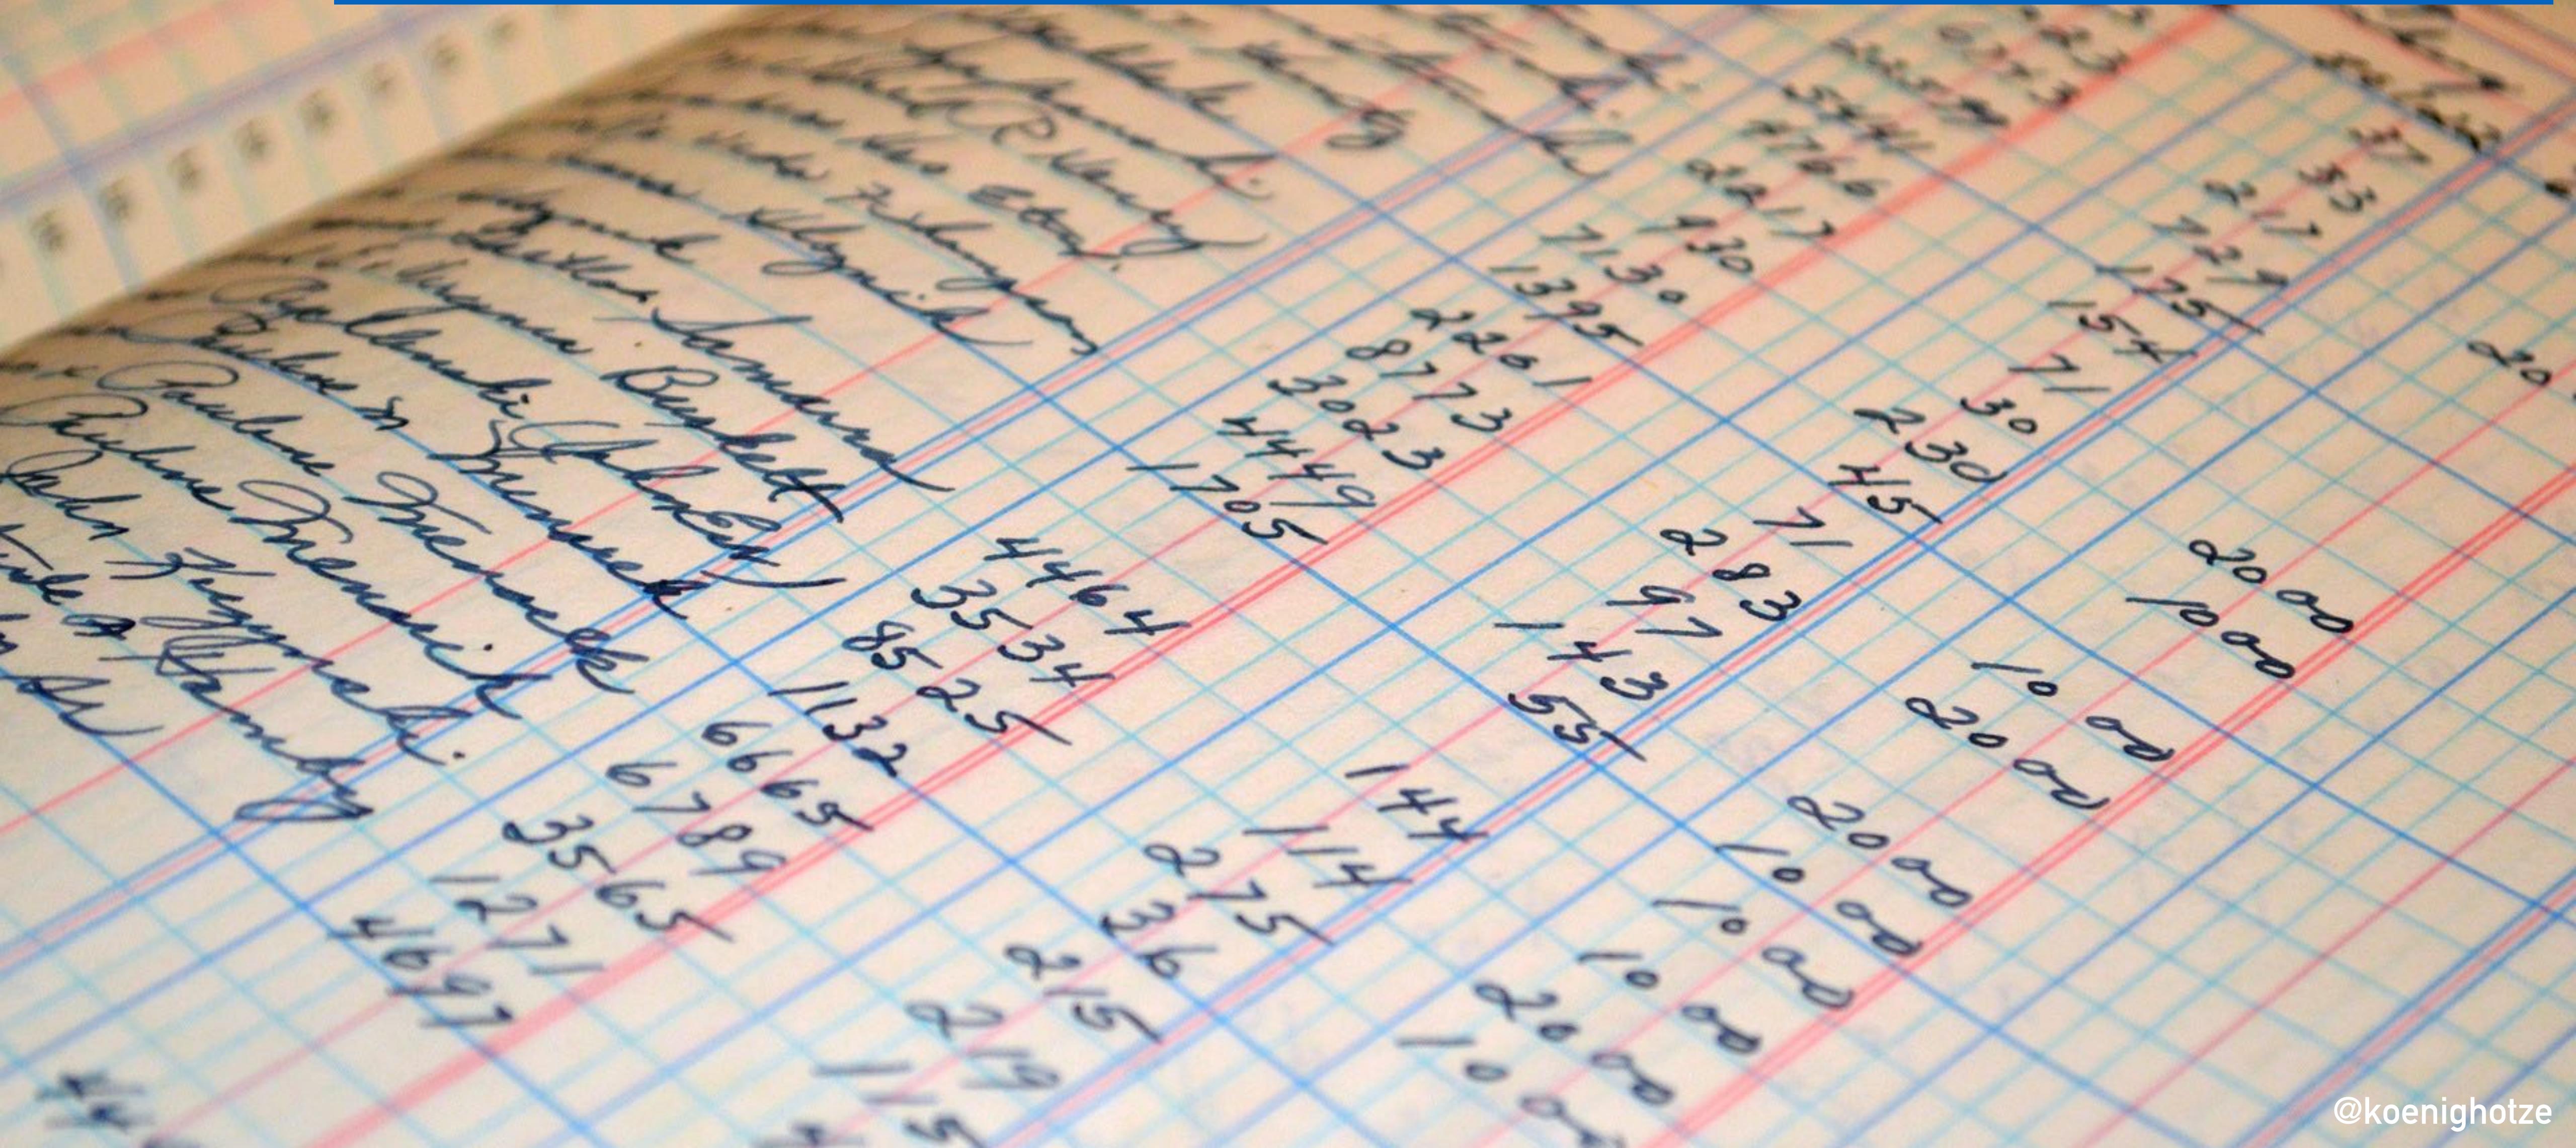
Not a Kafka-rant

What works for us, might not work for
you and the other way around

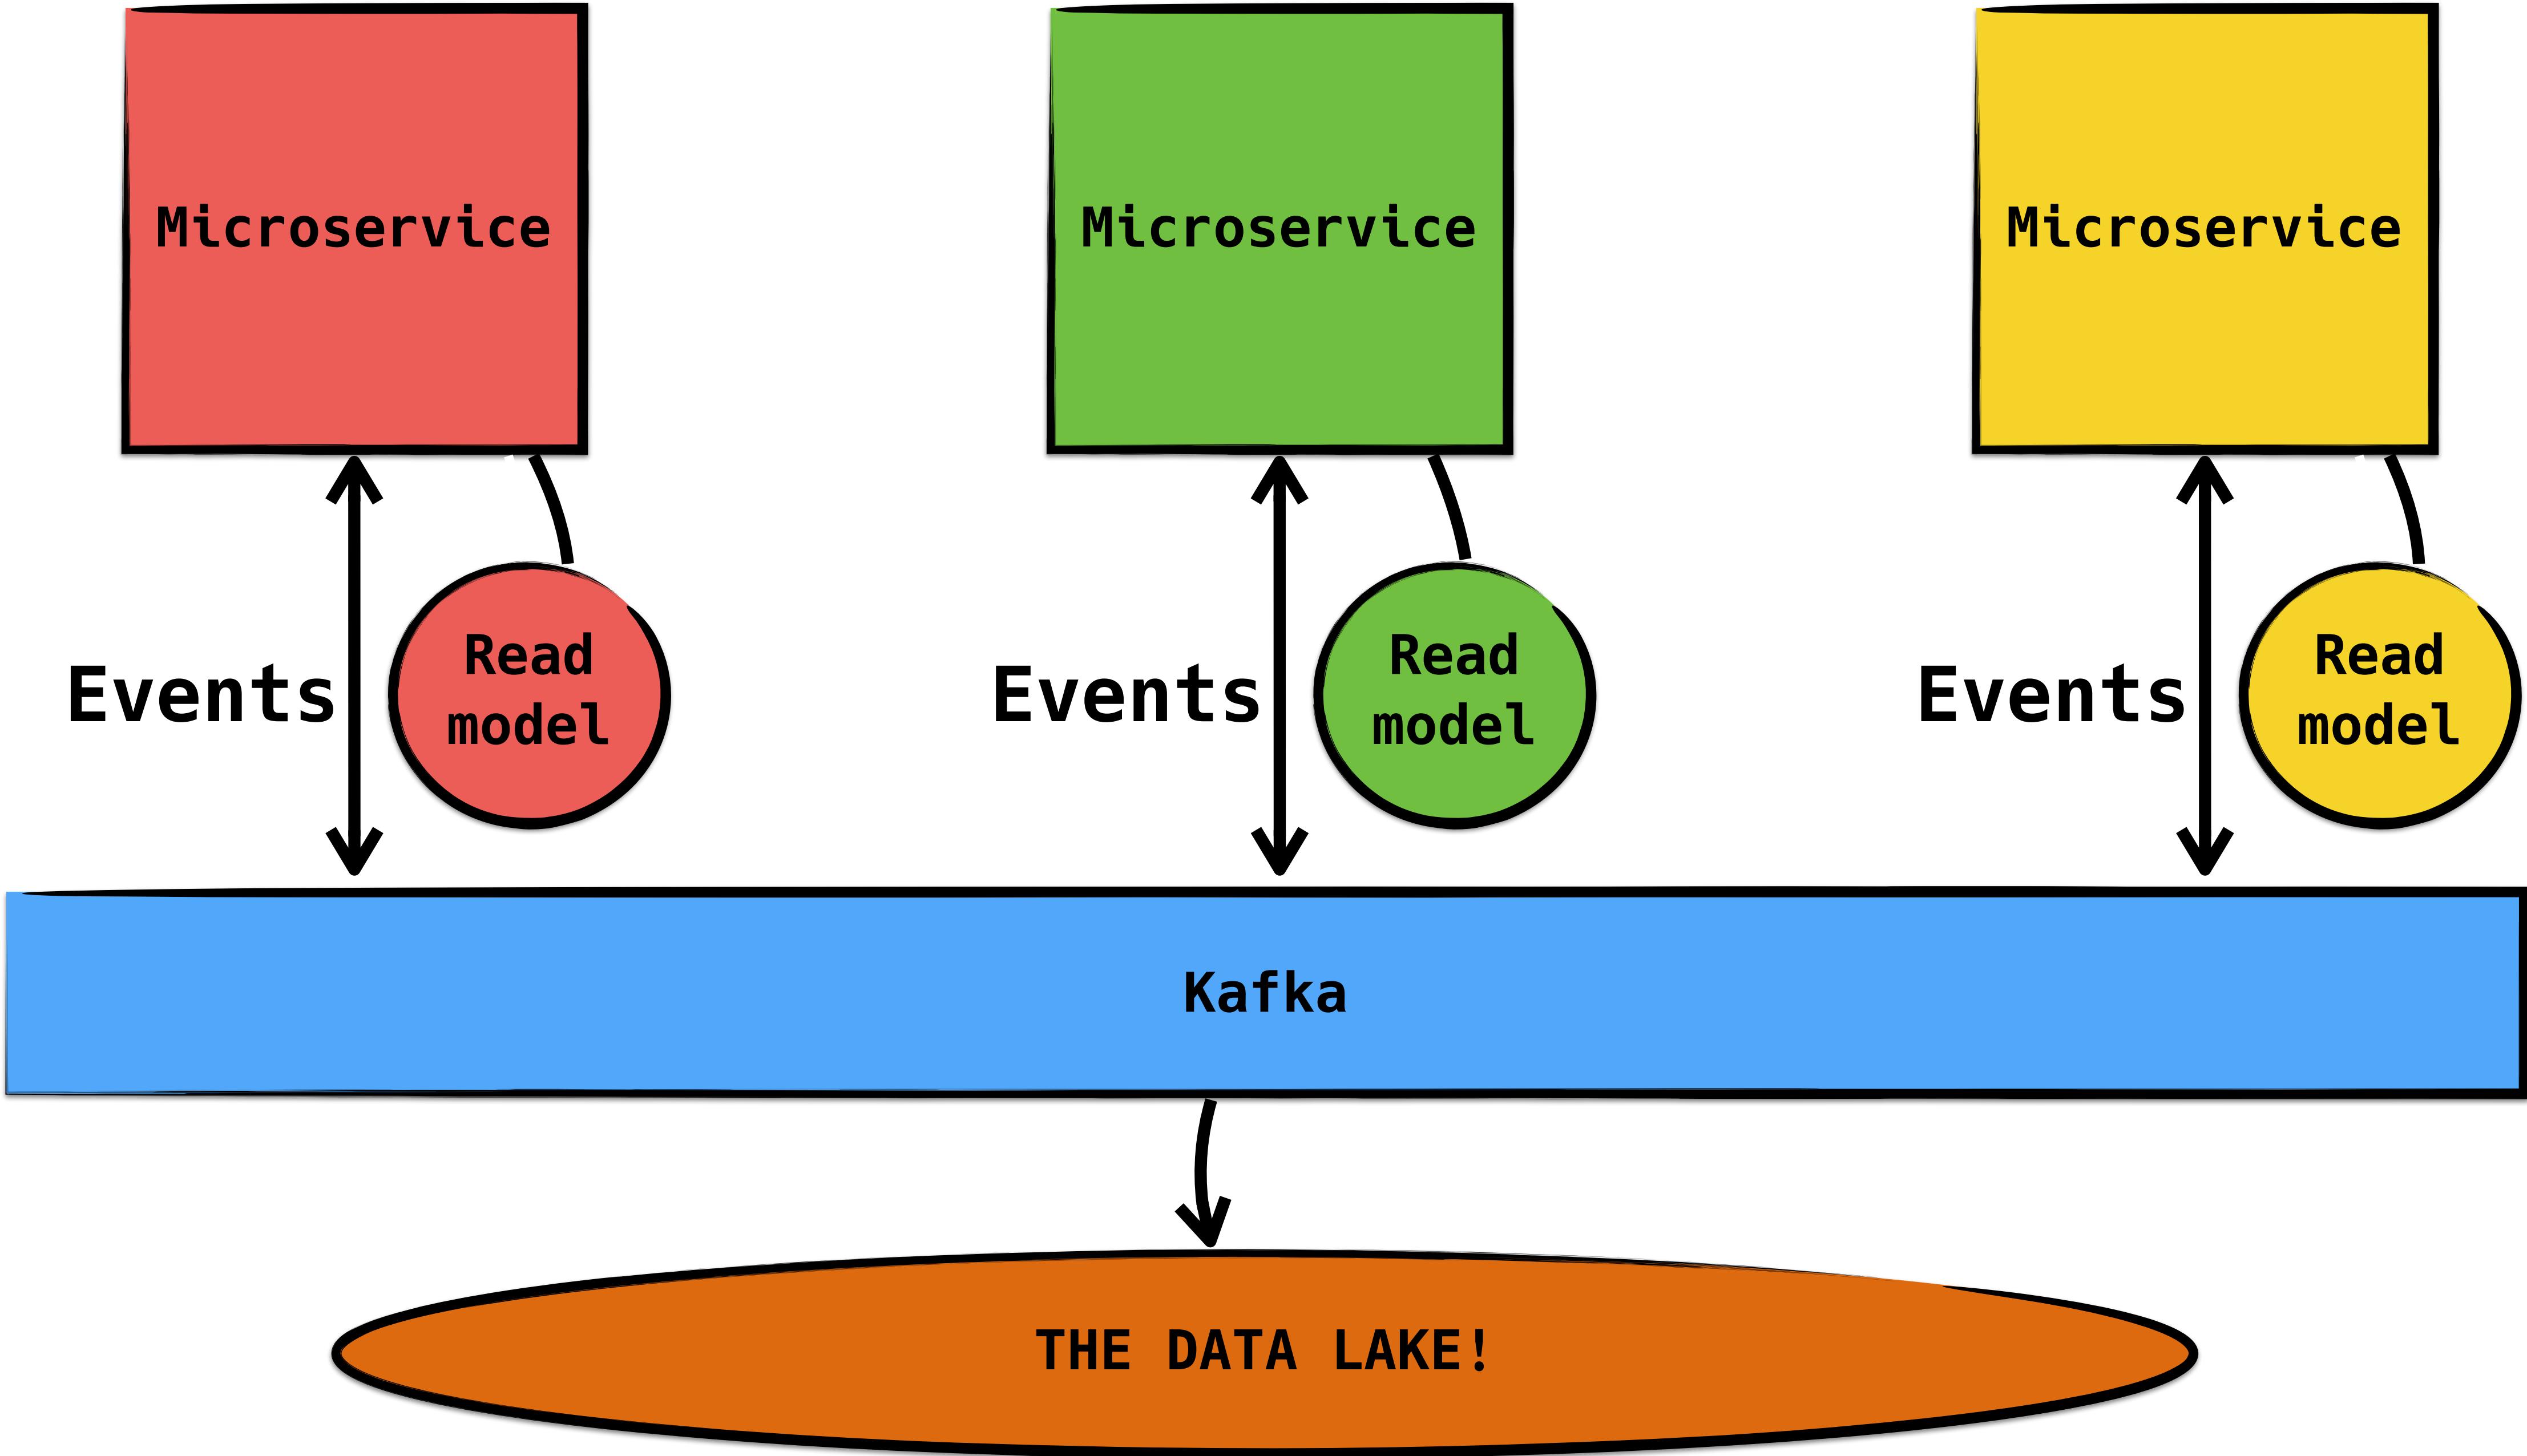


What we'll cover!

Eventsourcing bootcamp



Eventsourcing from an architects perspective





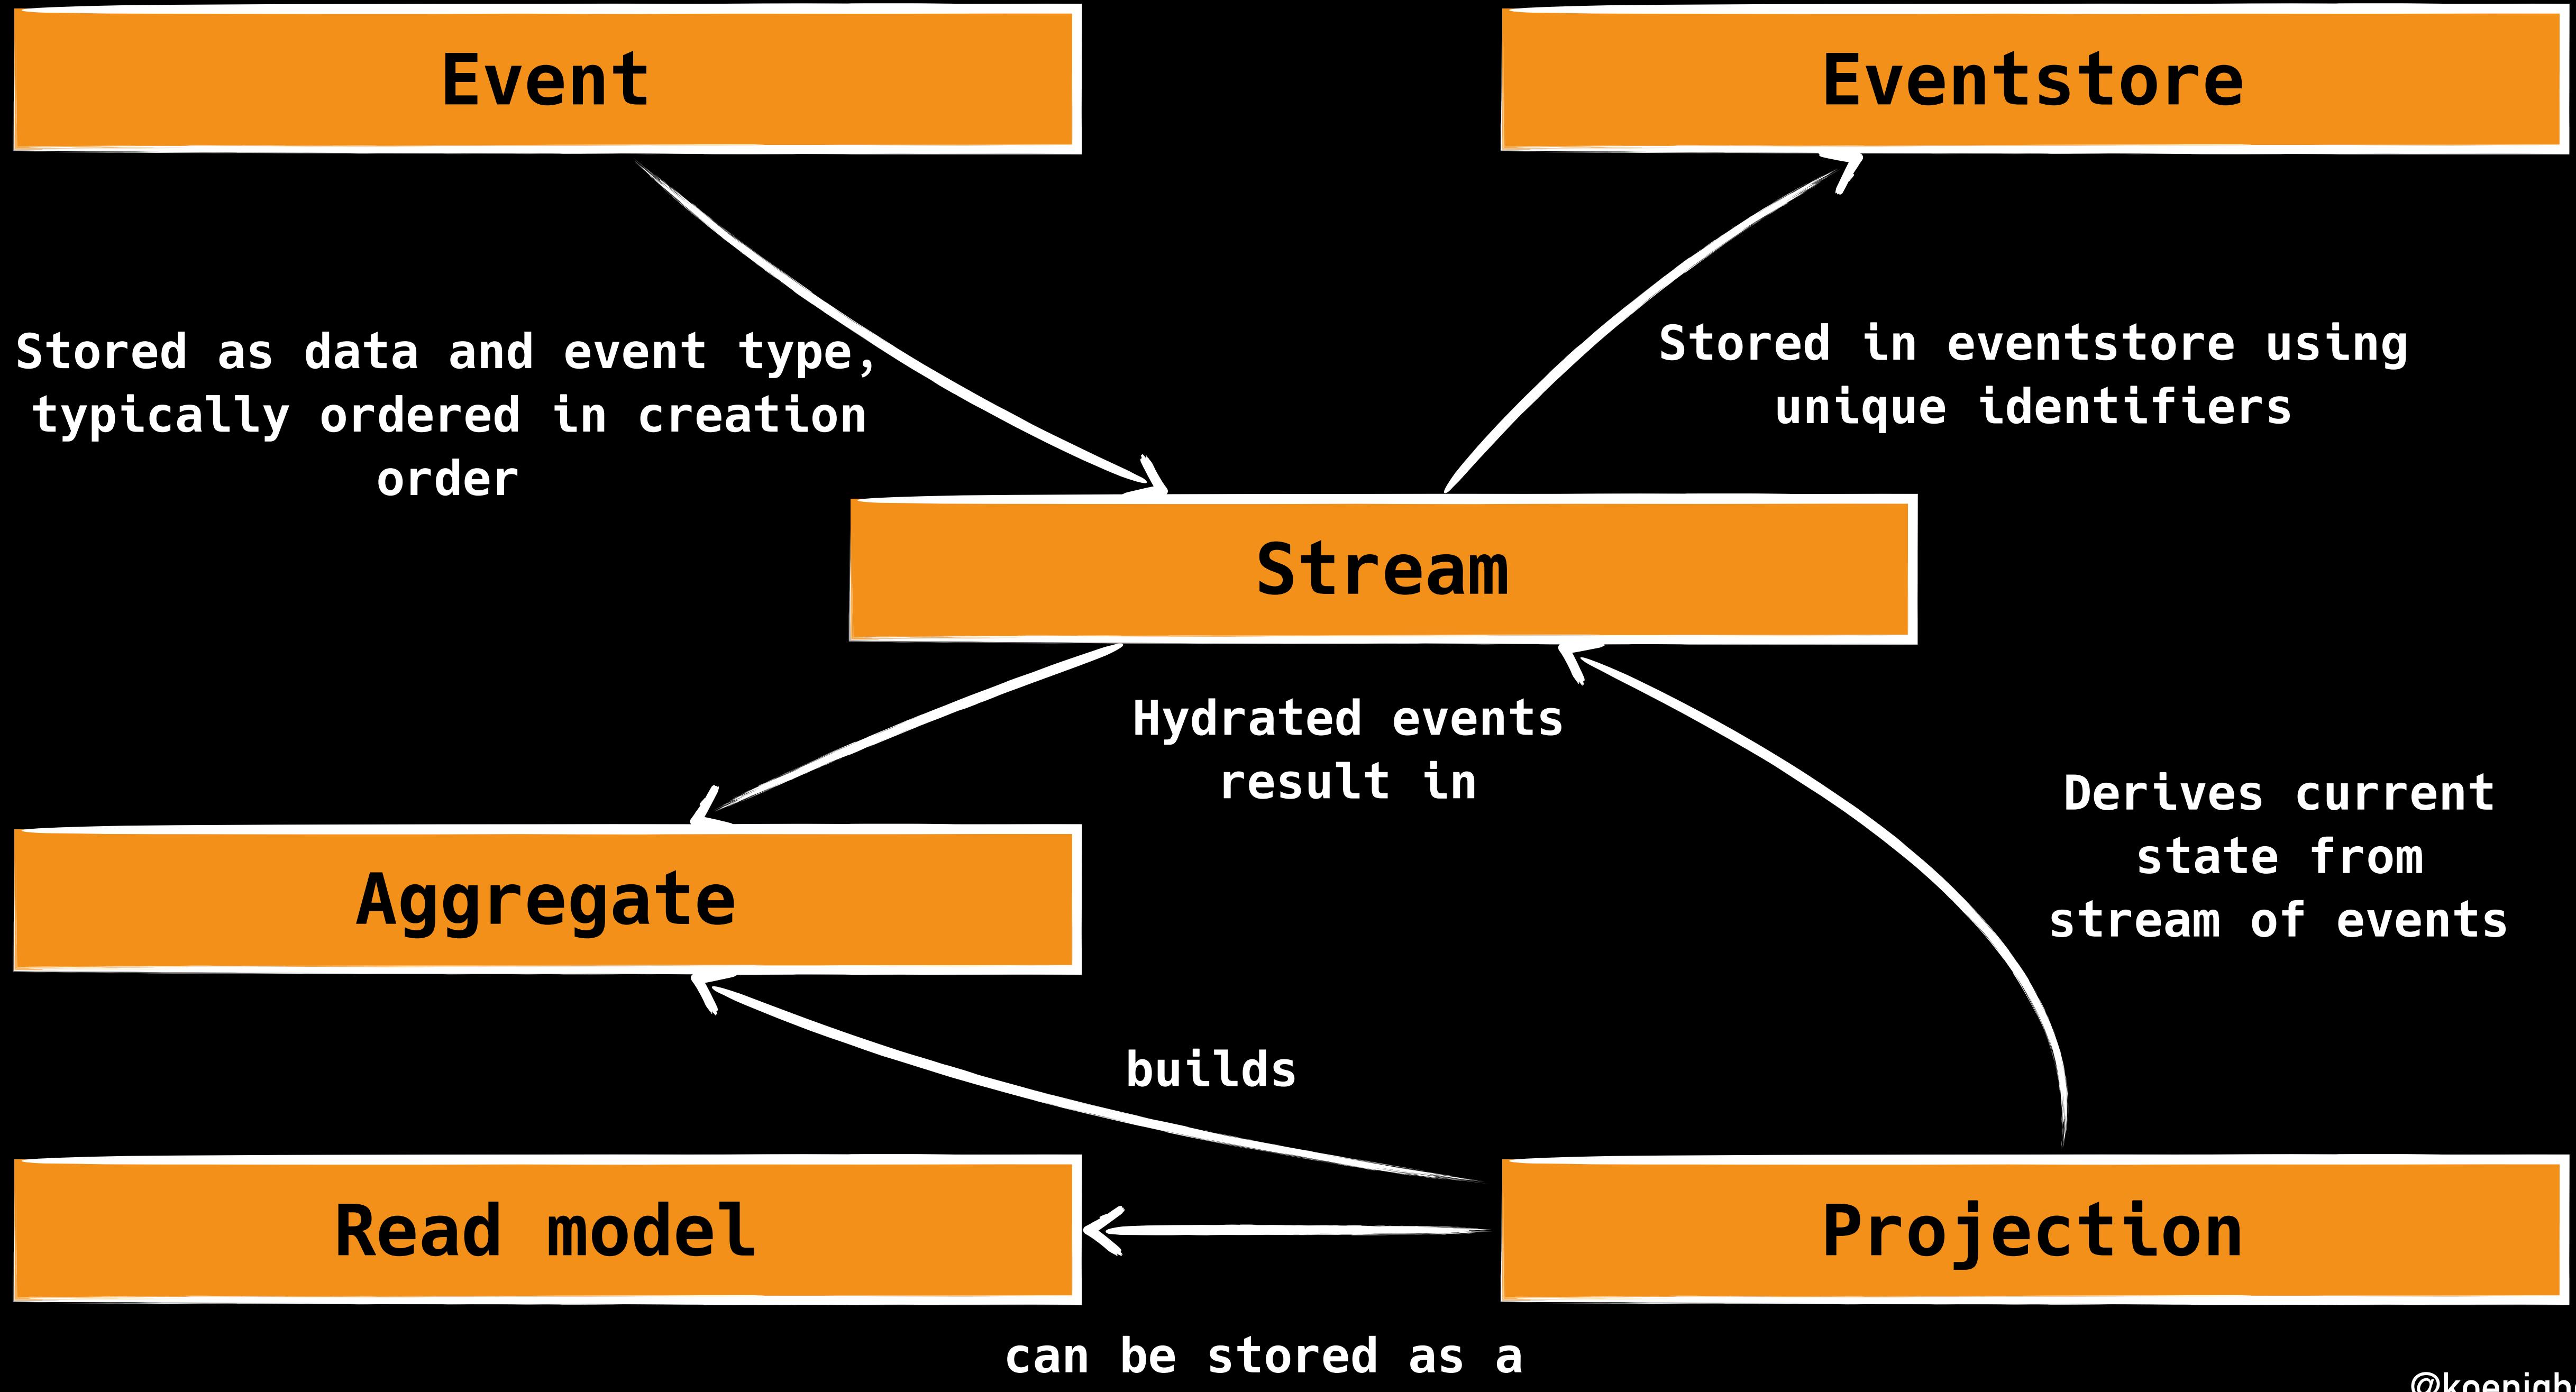


“Just” in Action

The idea in basic terms

Domain driven design
Event driven architecture
Distributed systems

Start with your domain
Find the bounded contexts
Find the root aggregates
Find events in your
ubiquitous language



User-Stream

UserRegistered

userId: 9714de5c...

UserOnboarded

email: foo@bar.de
address: ...

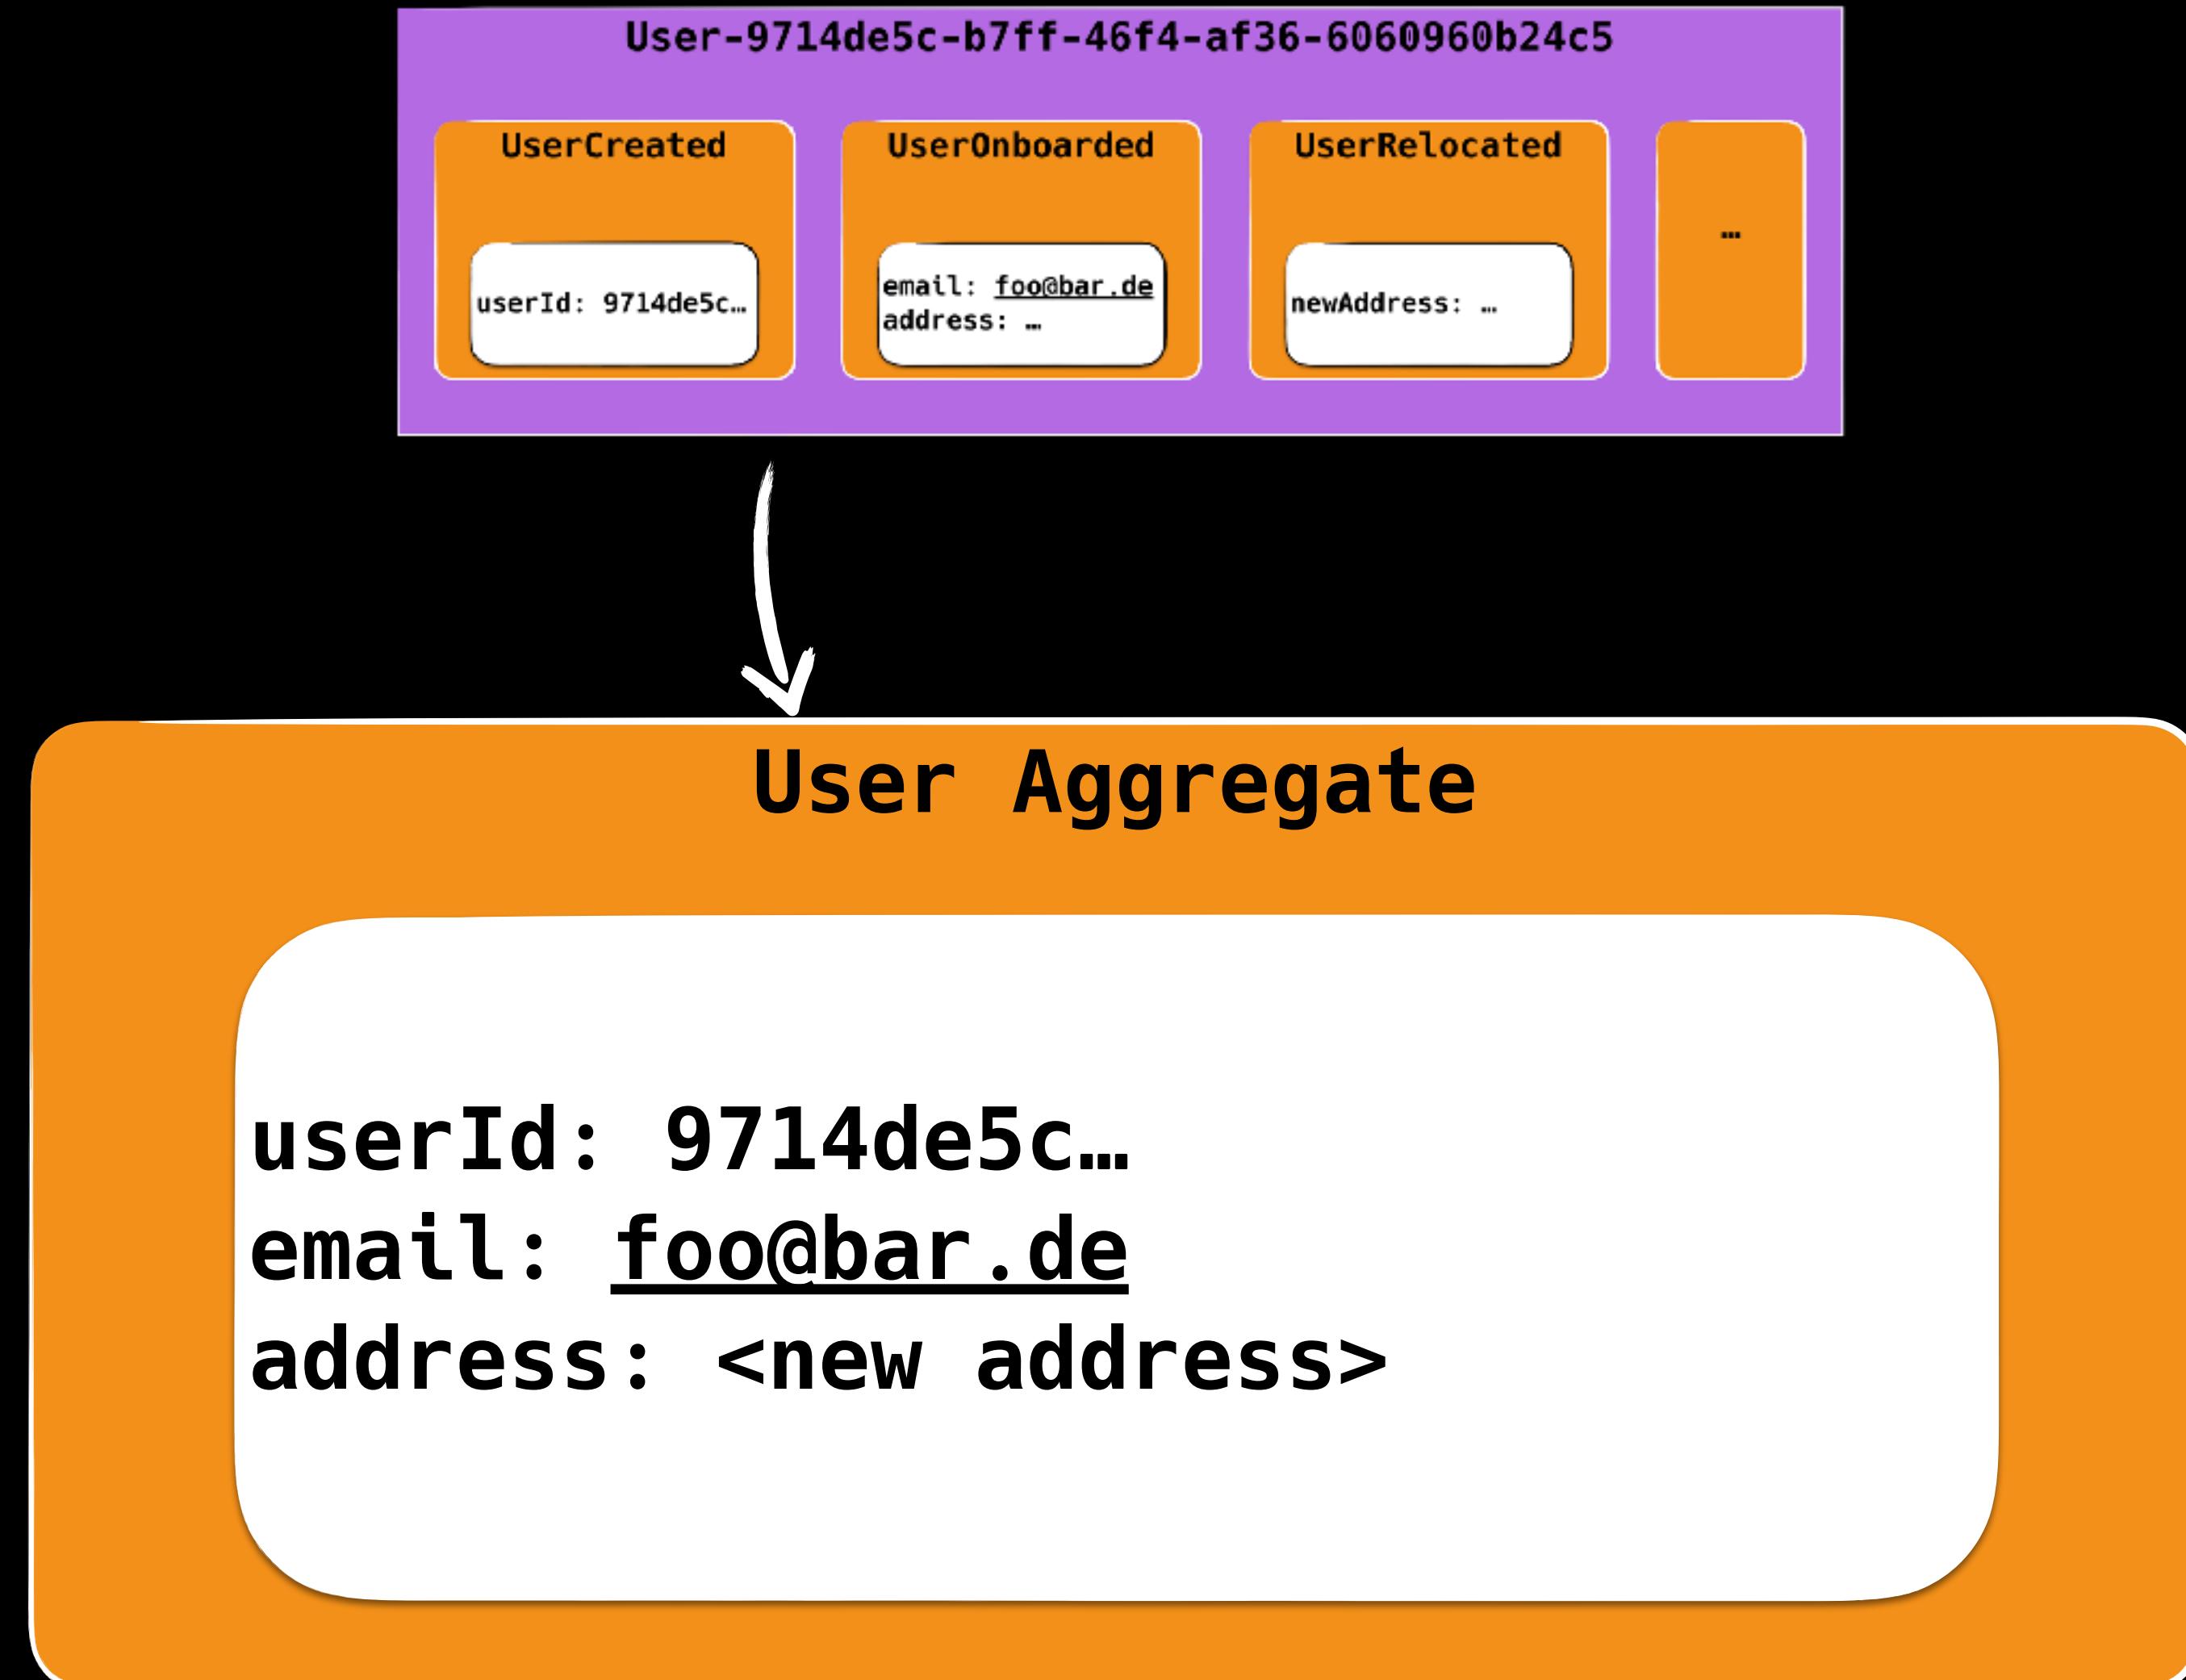
UserRelocated

newAddress: ...



Direction of time

Represents
the user at
the time of
the last
event read



Eventsourcing helps answering the question
of dealing with **data in distributed systems** in
a scalable way.

It makes the **dynamics** of your systems
explicit as **first class concepts**



Read models

...and why you may not need to them (initially)

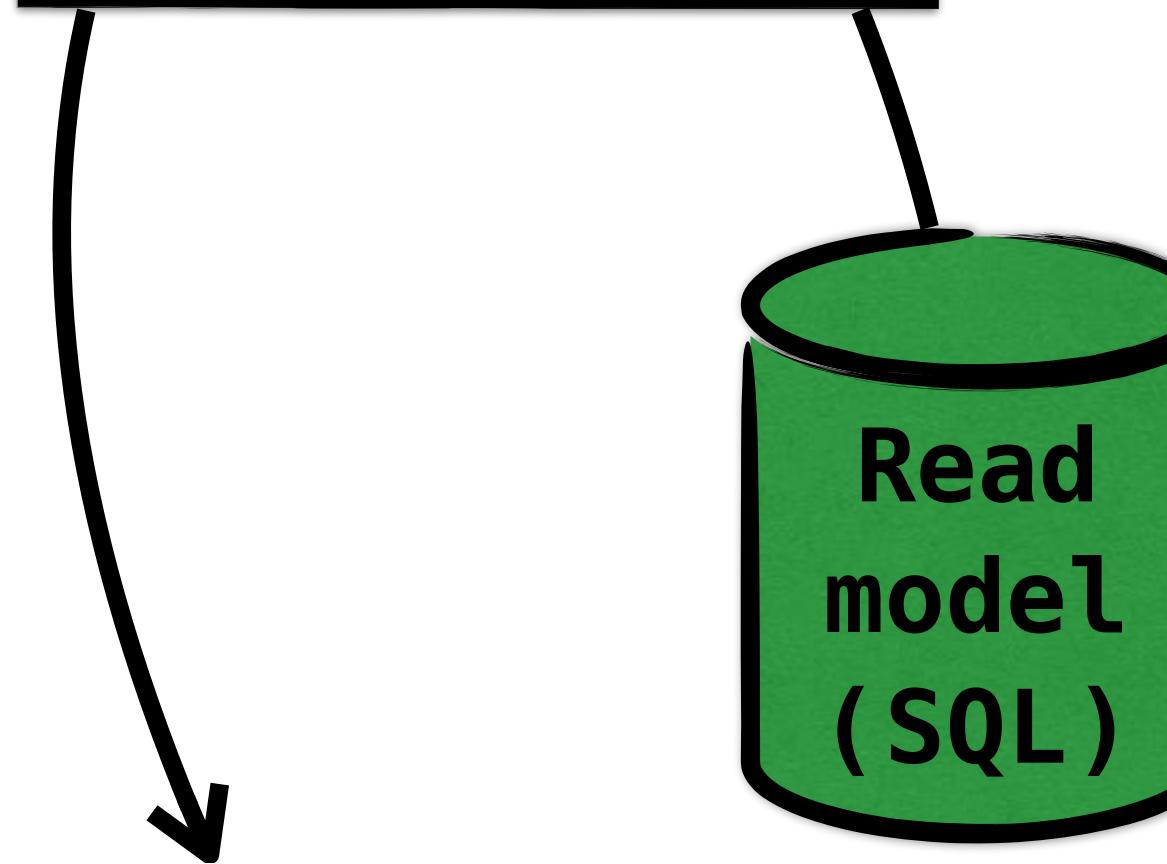
**Read models are local hydrates of
the events stored in specific streams**

```
readModel = Stream.of(events)
    .leftFold(handlers)
```

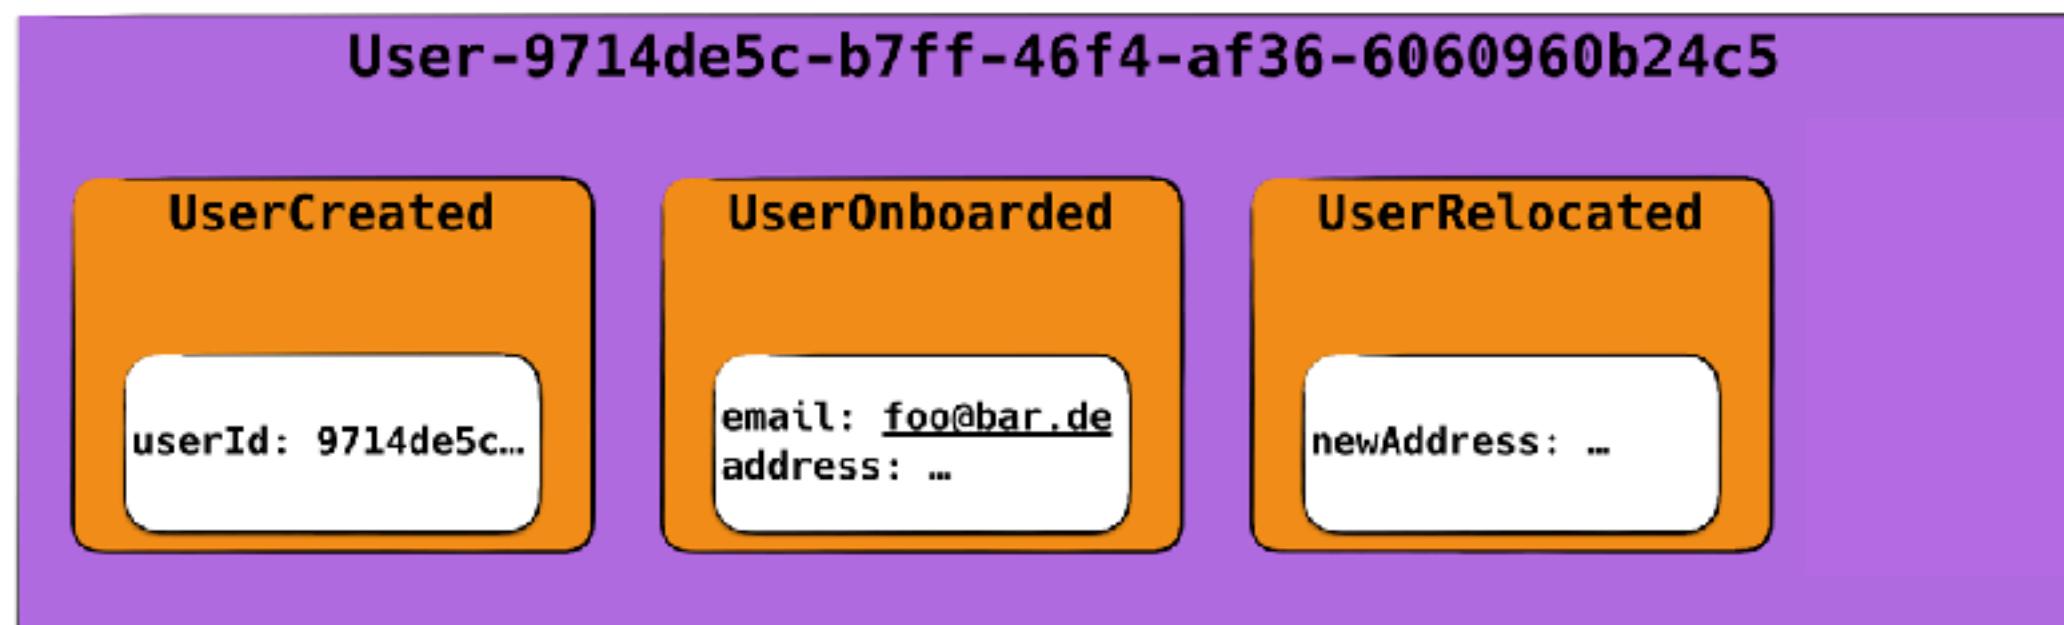
How can we handle read models?

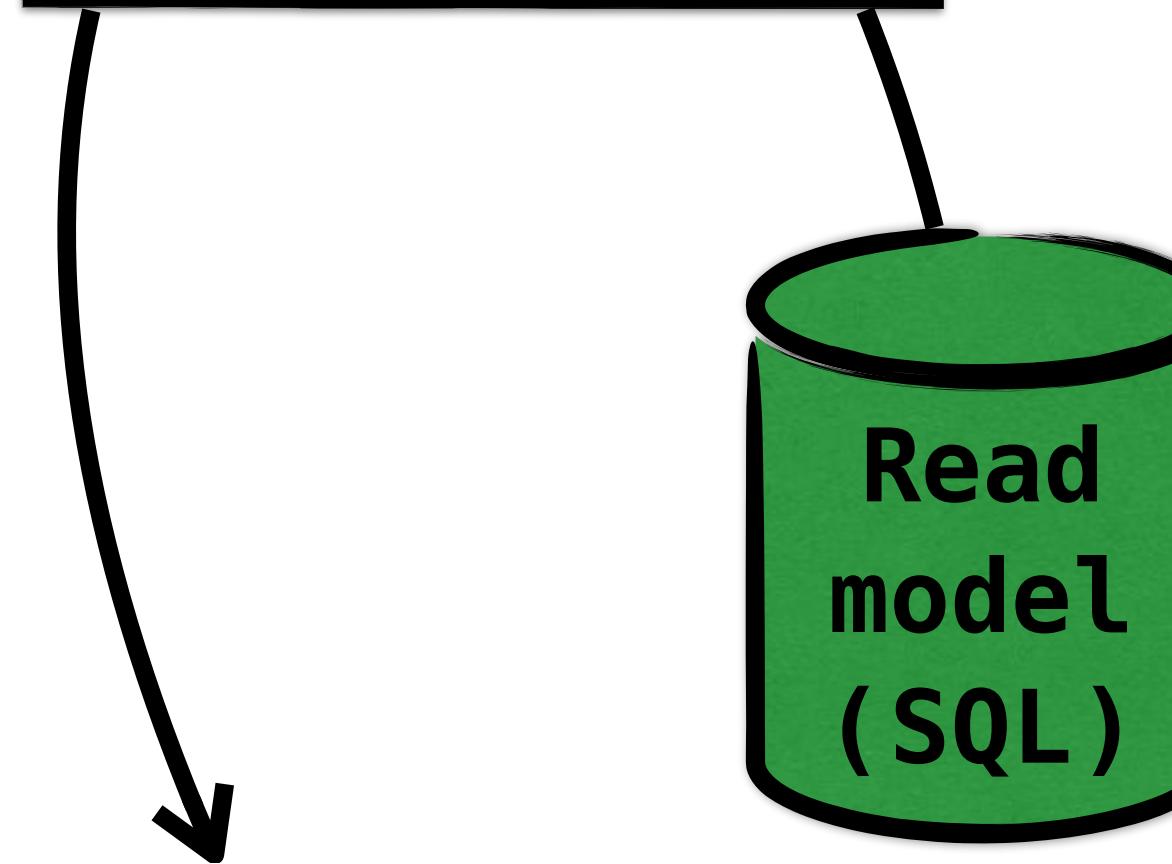
**“Just” use a local
database**



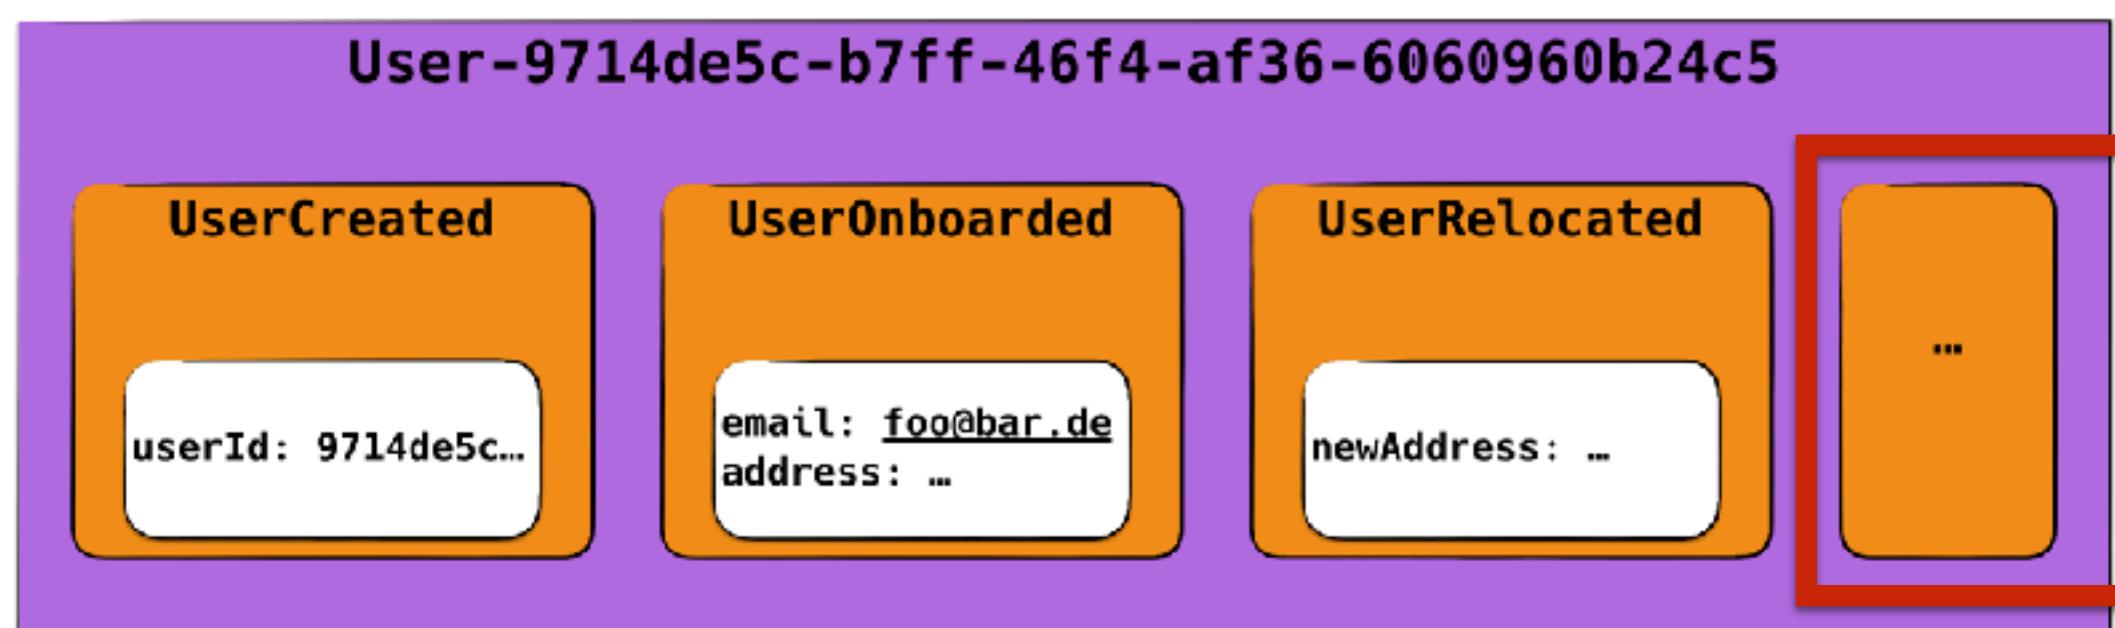


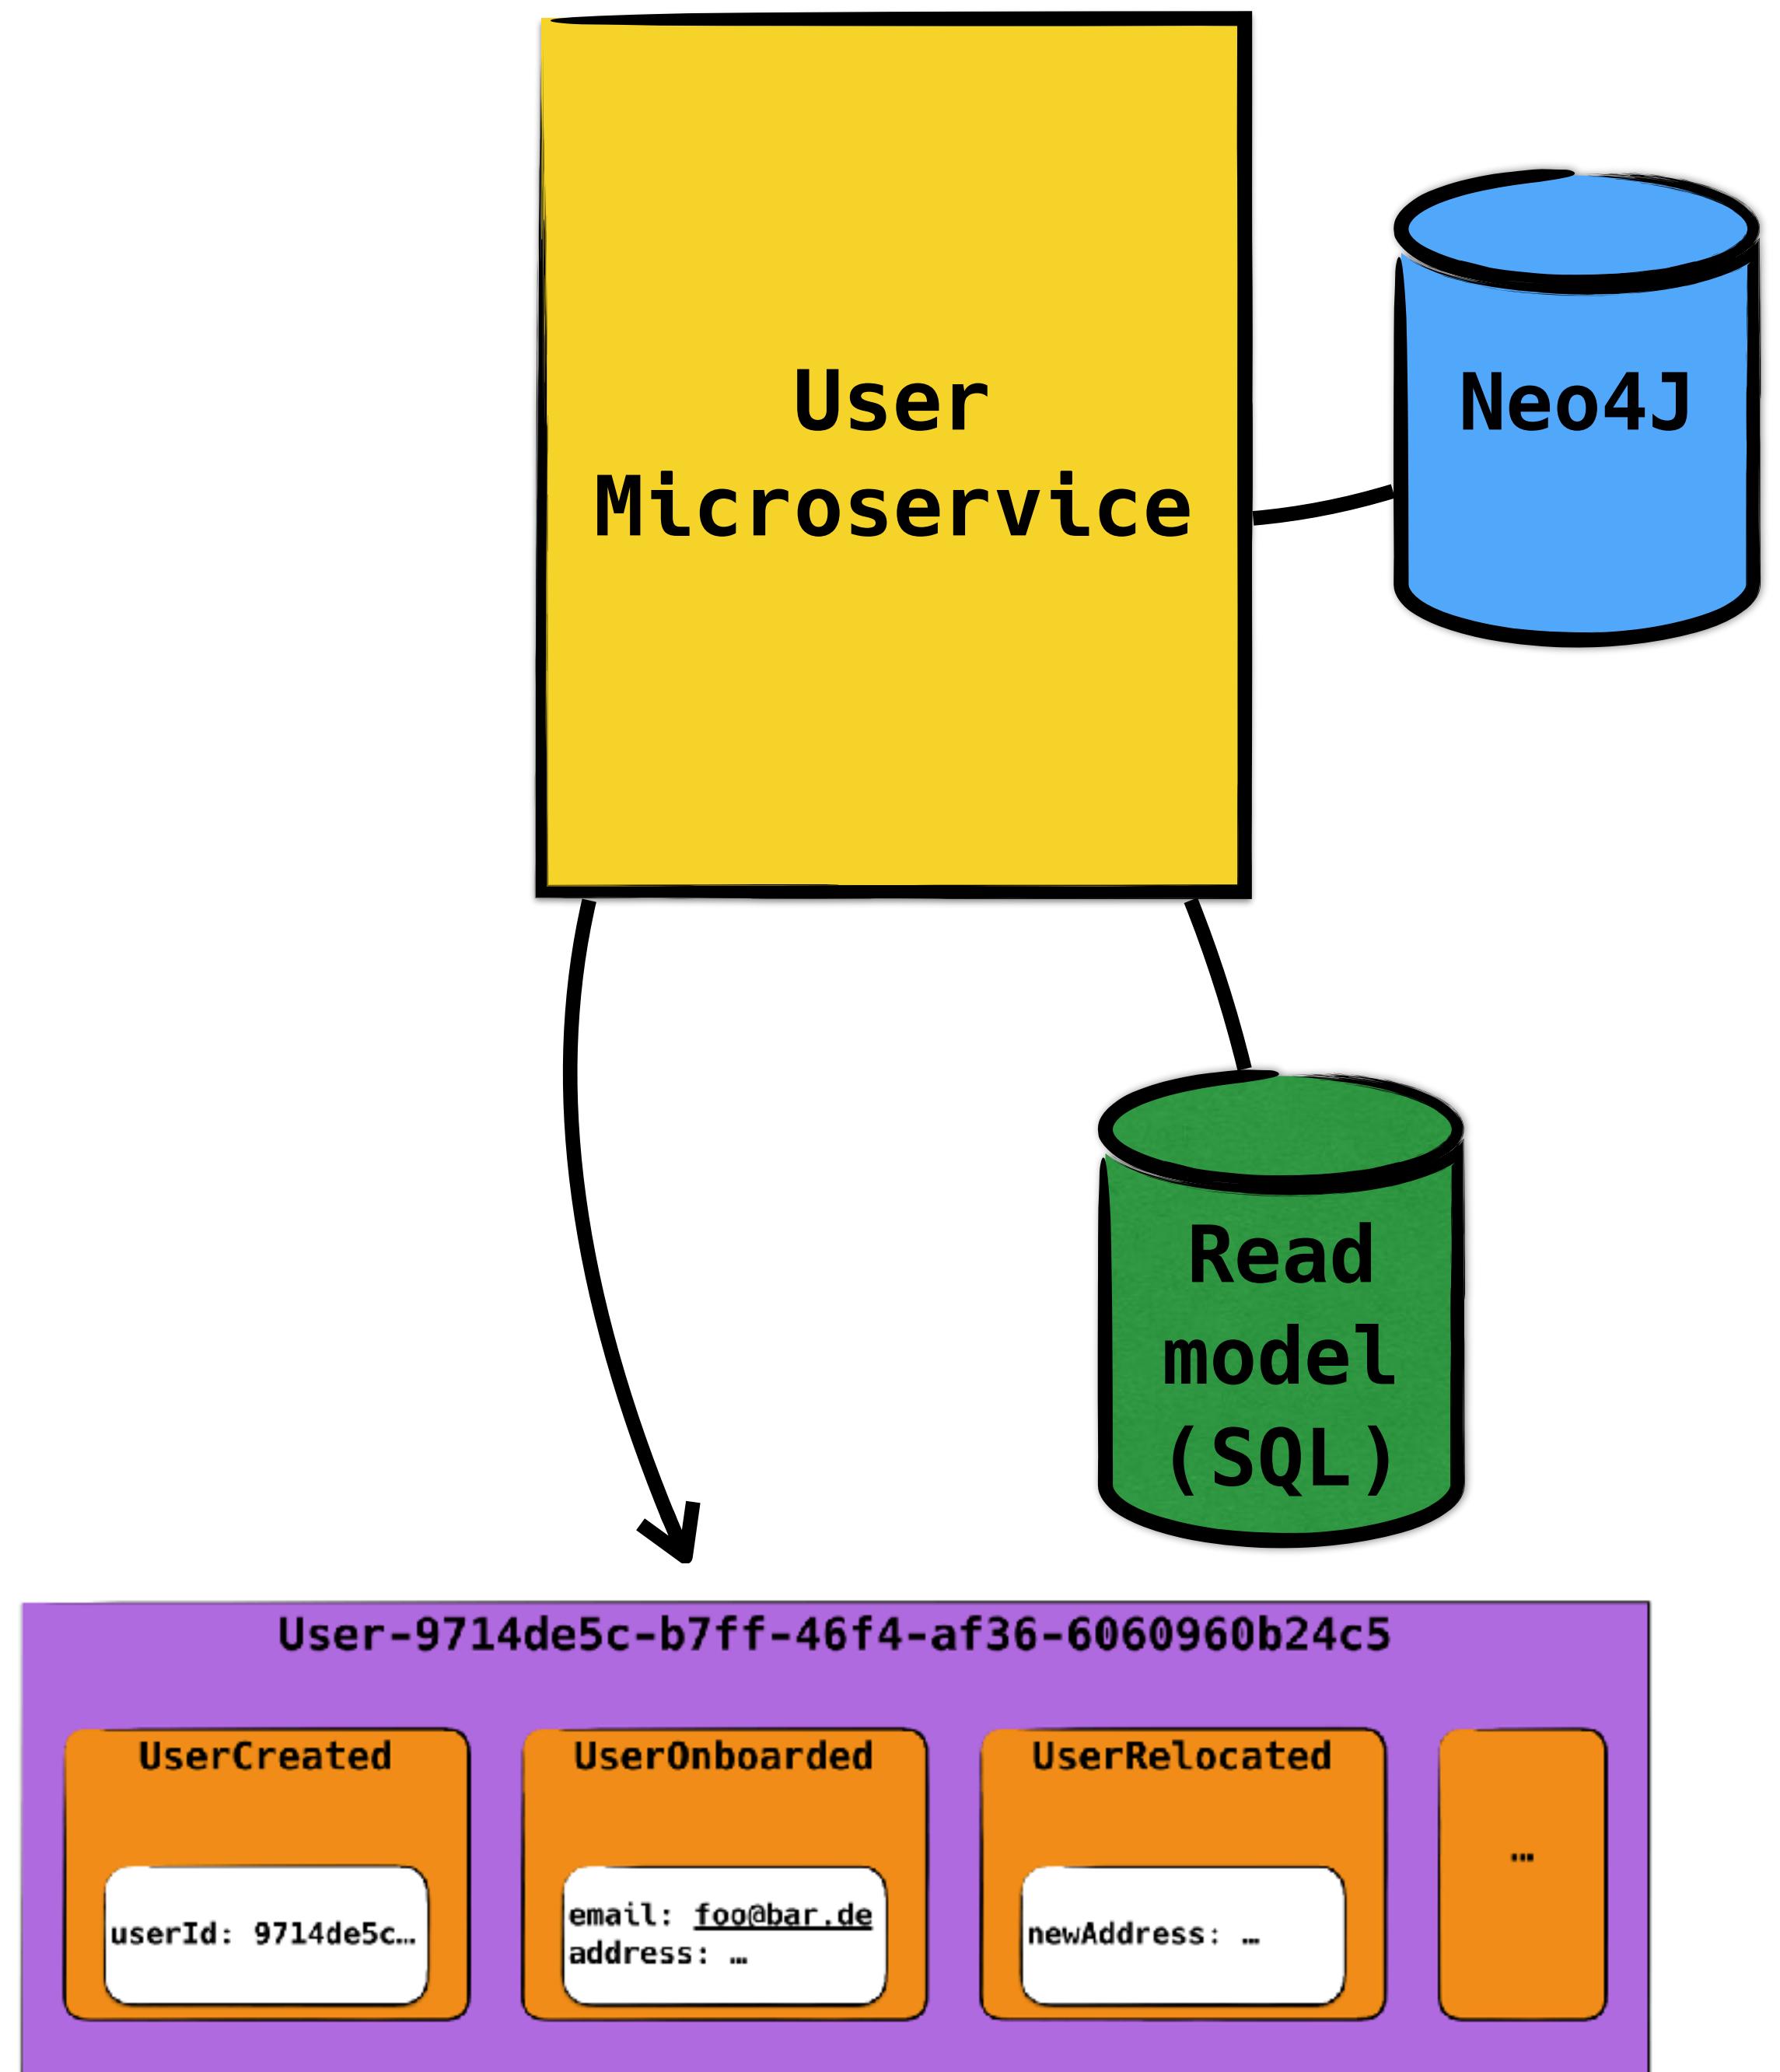
user id	last event id	user name	email	street
9741...	3	David	foo@bar.de	...
4532...	7	Martin	null	...





user id	last event id	user name	email	street
9741...	4	David	<u>qux@ba.ze</u>	...
4532...	7	Martin	null	...





Challenges?

Eventual consistency

Replays

Re-deliveries

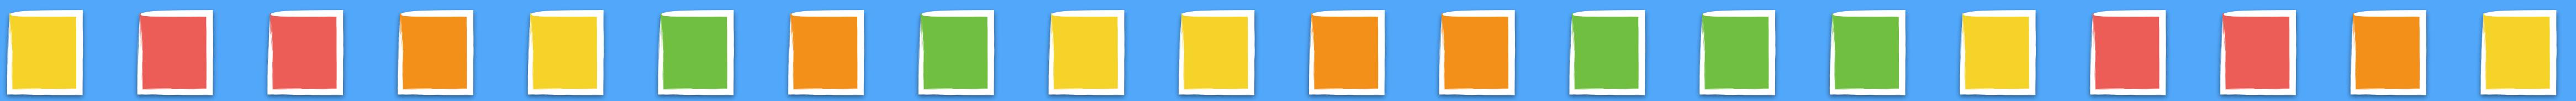
Operational complexity

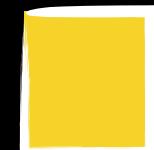
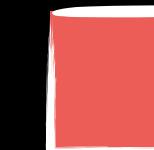
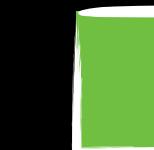
You may **not** need a read model

Typical strategies for storing events

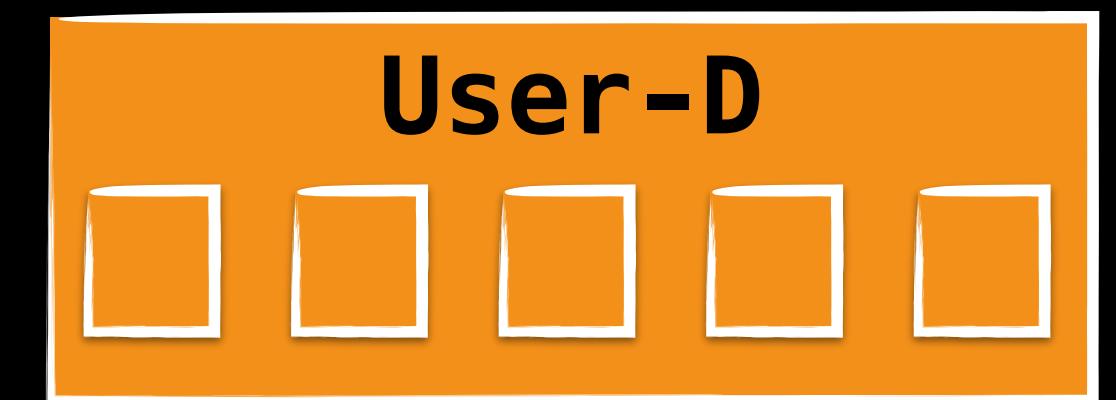
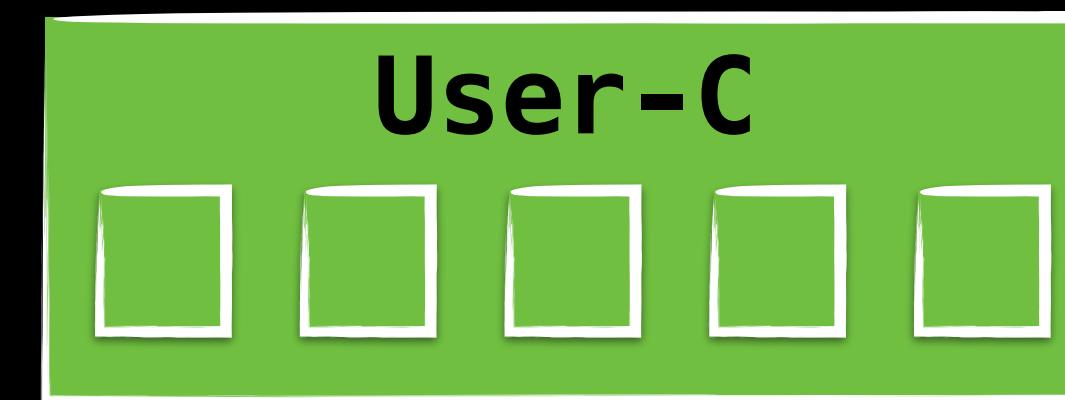
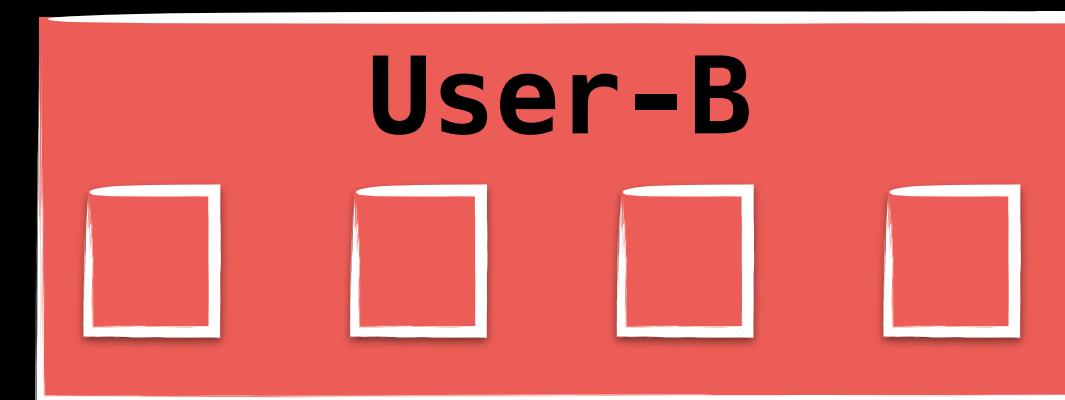
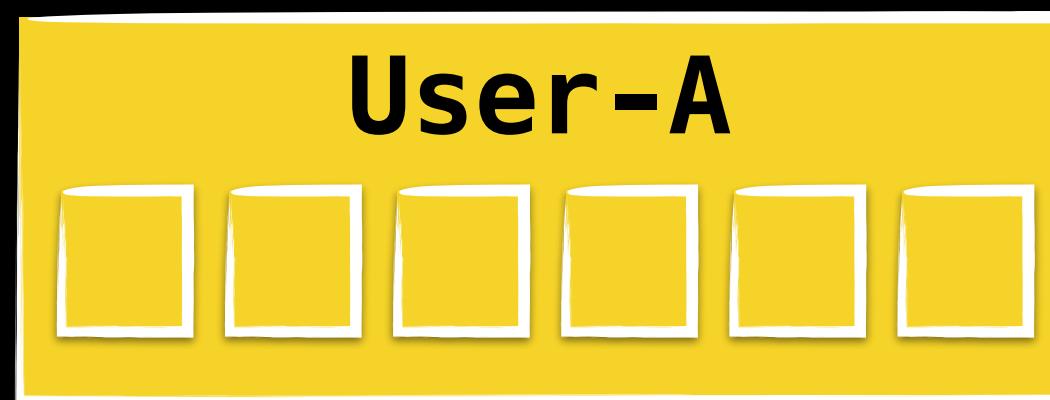
Maybe all events of an aggregate
type in a **single stream**?

Users

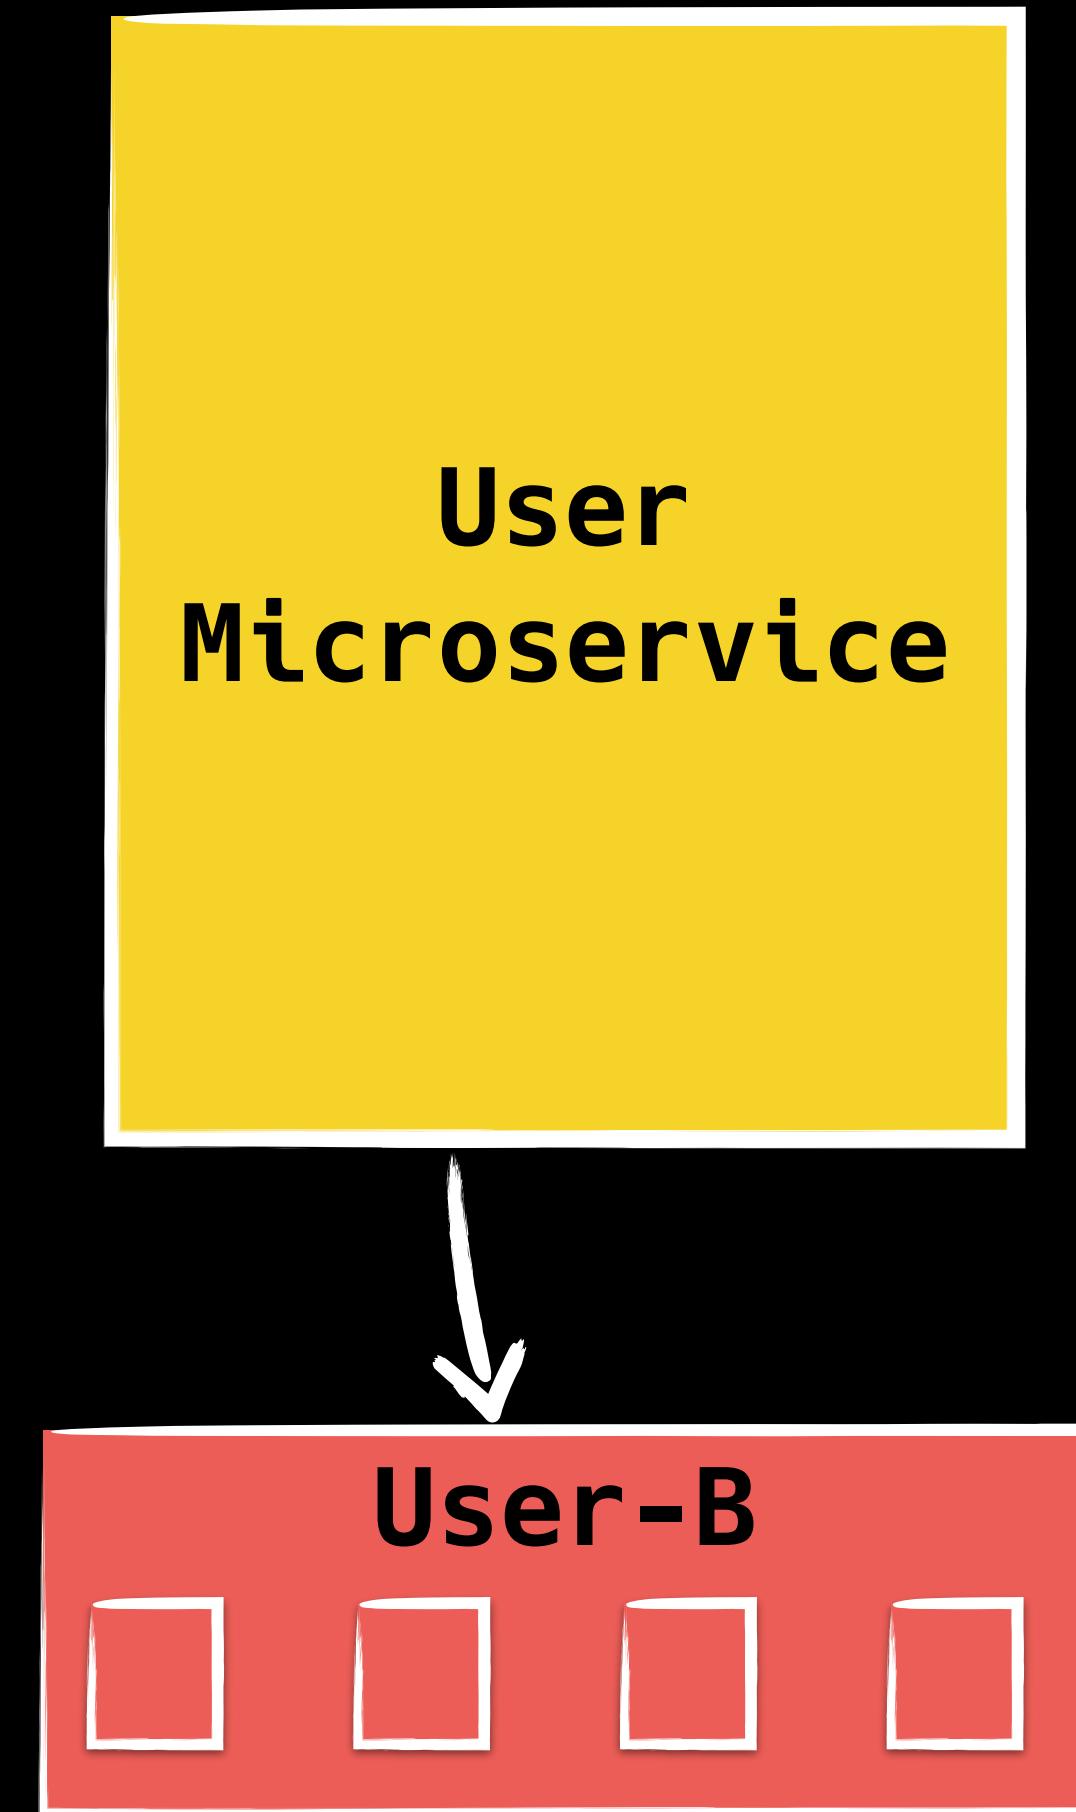


 User A  User B  User C  User D

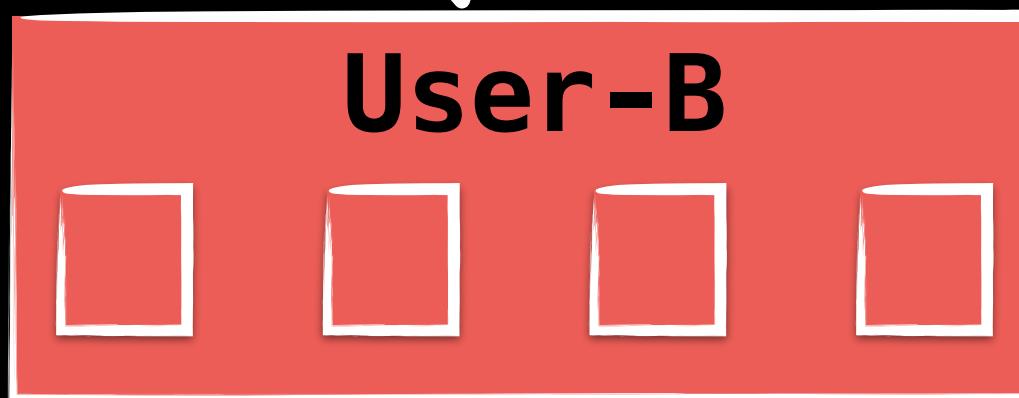
Better: One stream per aggregate



GET /users/B

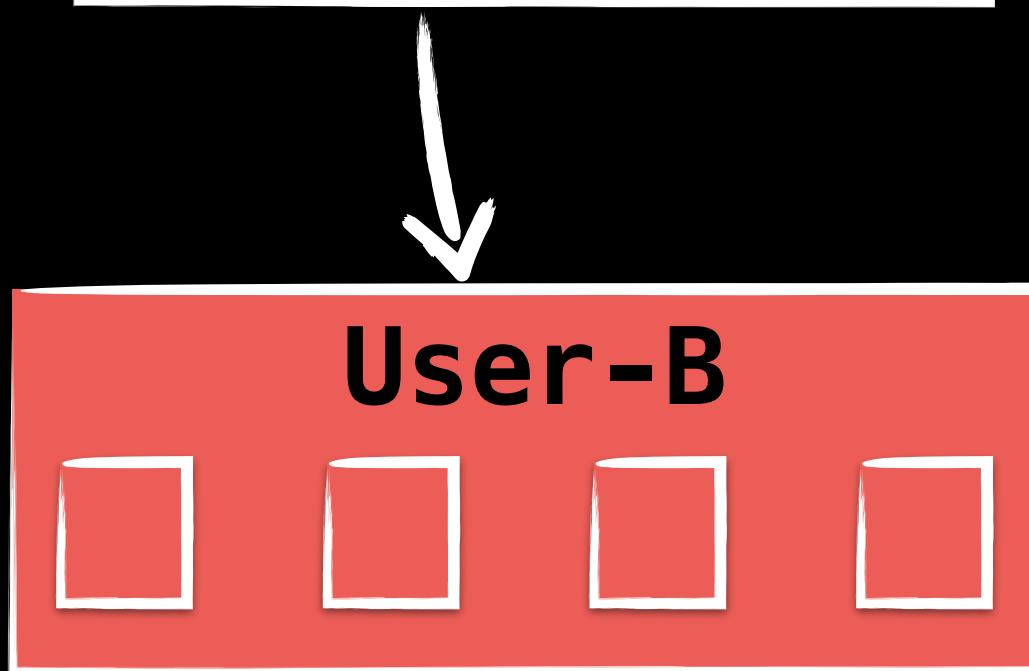


GET /users/B



```
● ● ●  
handlers = {  
  'UserCreated': (current, event) => {},  
  'UserOnboarded': (current, event) => {},  
  'UserDeleted': ...  
}  
  
aggregate = readAggregateFromStream(  
  'user',  
  'B',  
  {},  
  fromStartOfStream,  
  handlers  
)
```

GET /users/B



● ● ●

```
handlers = {
  'UserCreated': (current, event) => {},
  'UserOnboarded': (current, event) => {},
  'UserDeleted': ...
}
```

```
aggregate = readAggregateFromStream(
  'user',
  'B',
  {},
  fromStartOfStream,
  handlers
)
```

Consistent read

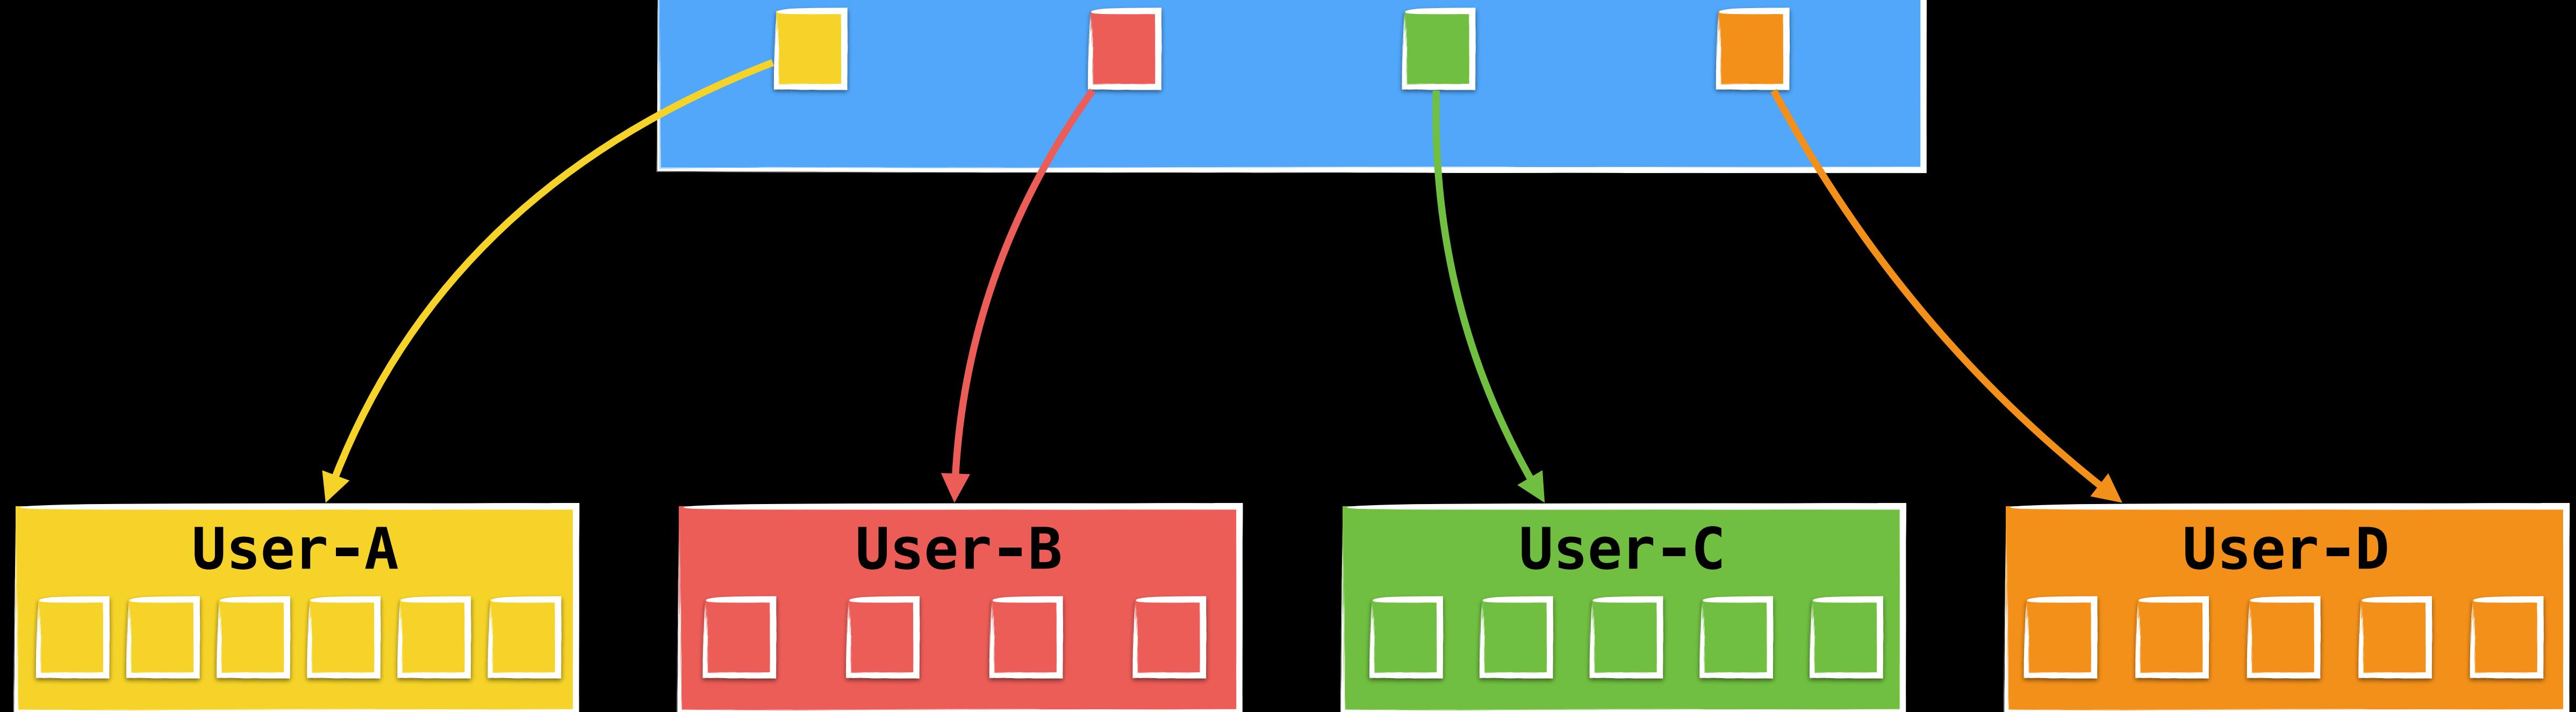
No operational overhead

Super-simple programming logic

But, what about speed?

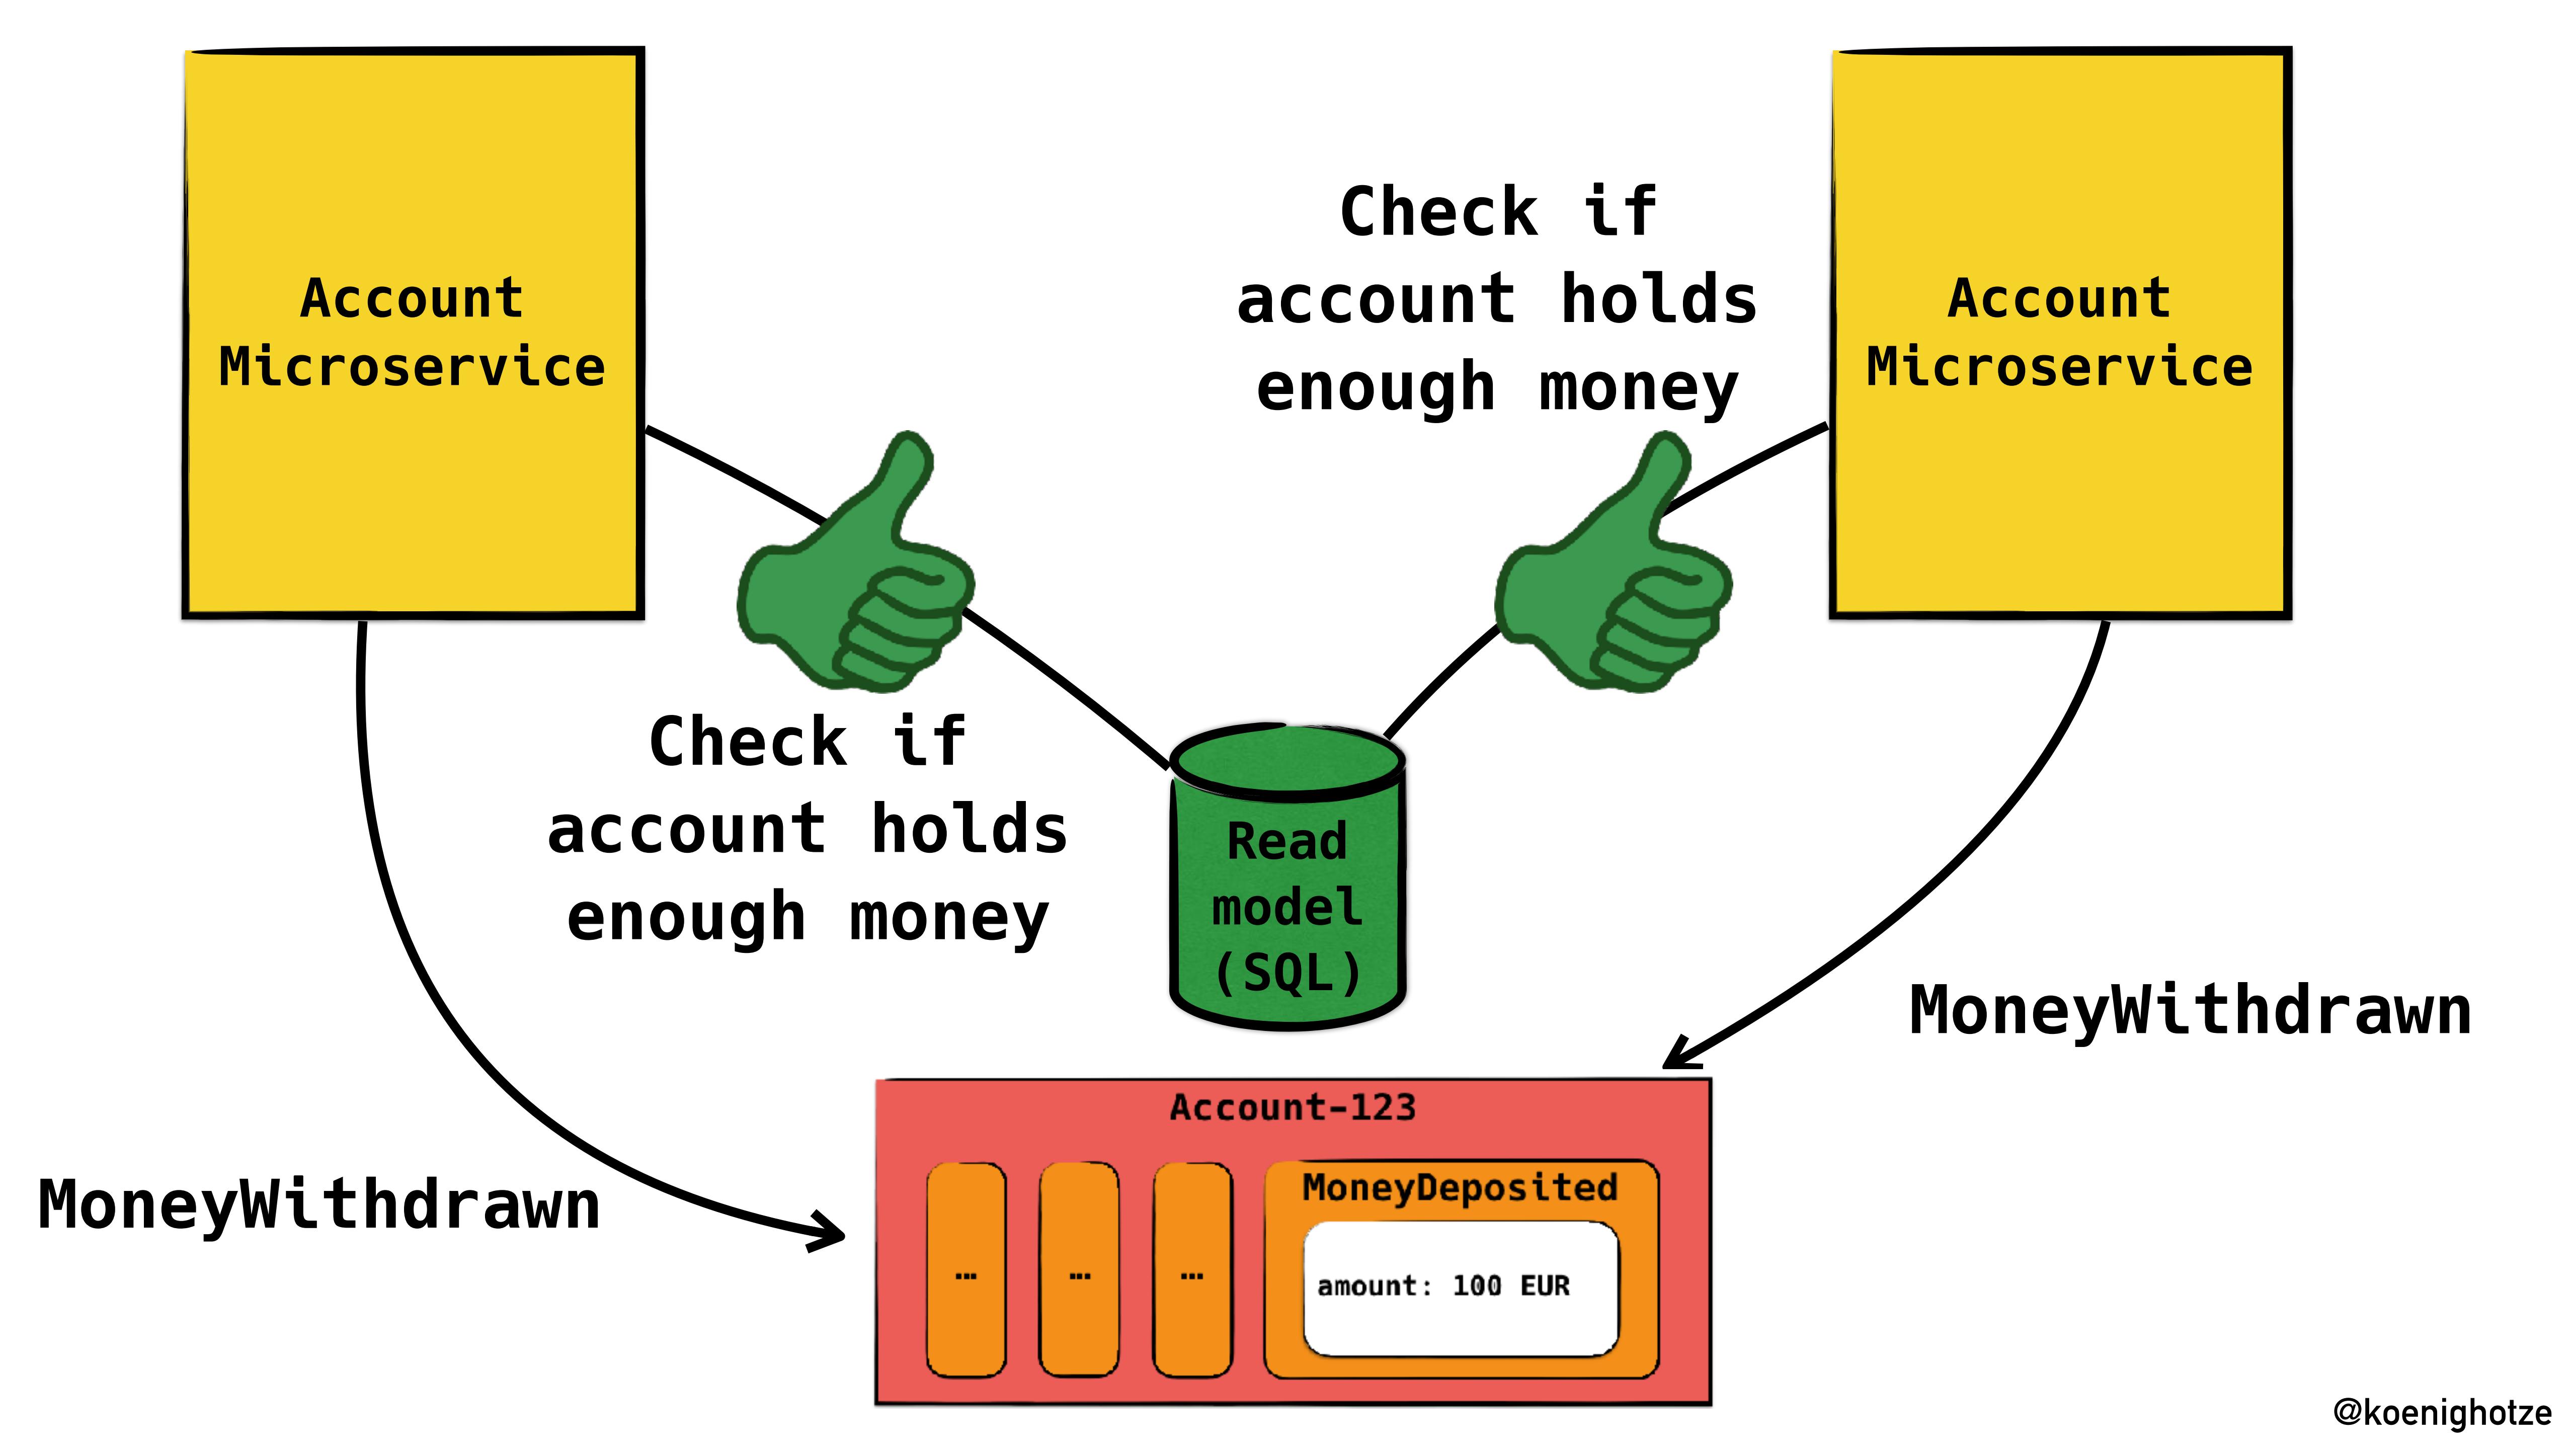
What about queries like 'all
users with age > 18'?

Adult-Users-Projection



But I can use a readmodel
for fast validation, can I?

Maybe not



Validation against a read model is prone to
inconsistencies

Prefer validating against the eventstore itself

Most event handlers are neither CPU
or IO intensive

Prefer small aggregates, if it makes
sense in your domain

Measure and introduce persistent
read models only if needed

Transactions, concurrency and your eventstore

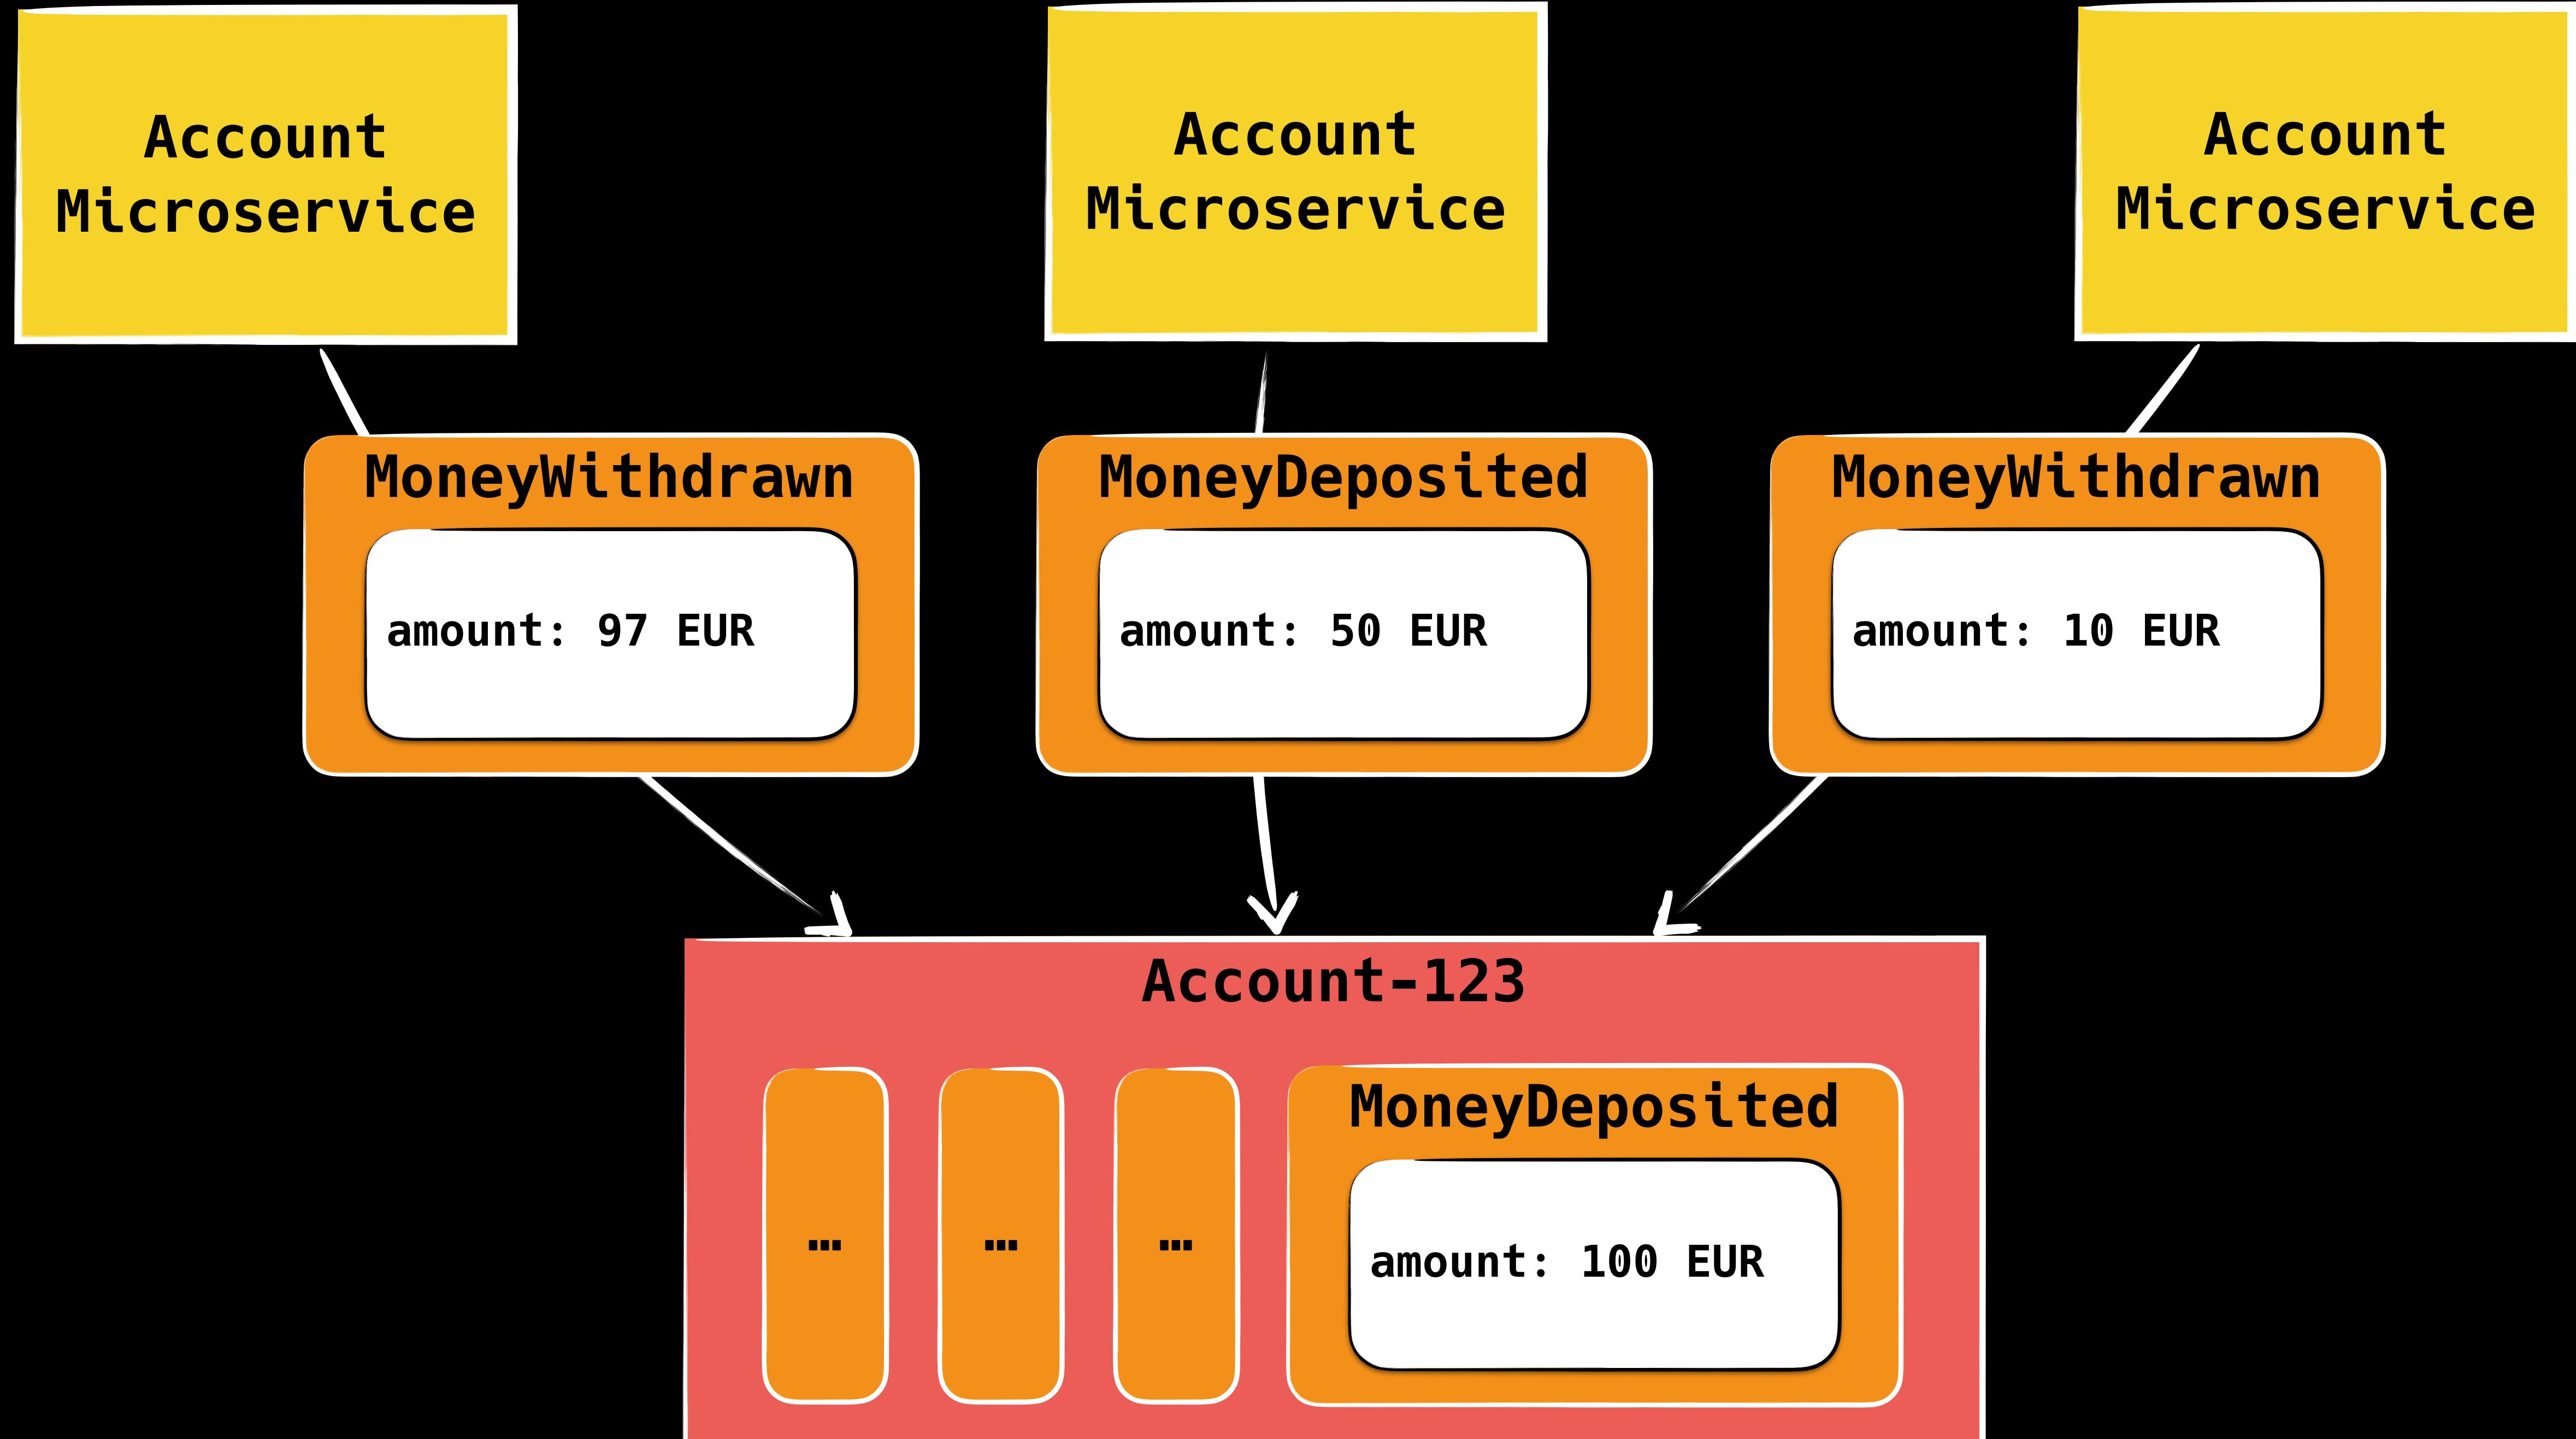


How can we guarantee
correctness when writing?

“Only withdraw money, if the bank account holds enough money!”*

*Actually, a real bank would not want such a business rule. They earn money if you overdraw your account. An overdraft fee is one of the most expensive fees banks charge. Just saying..

Are **YOUR** systems single-threaded?



The correctness of the
business result
depends on the order
of events

MoneyDeposited

amount: 100 EUR

100 EUR

MoneyWithdrawn

amount: 10 EUR

90 EUR

MoneyWithdrawn

amount: 97 EUR

-7 EUR



MoneyDeposited

amount: 50 EUR

140 EUR

Let's shuffle

MoneyDeposited

amount: 100 EUR

100 EUR

MoneyWithdrawn

amount: 10 EUR

90 EUR

MoneyWithdrawn

amount: 97 EUR

140 EUR

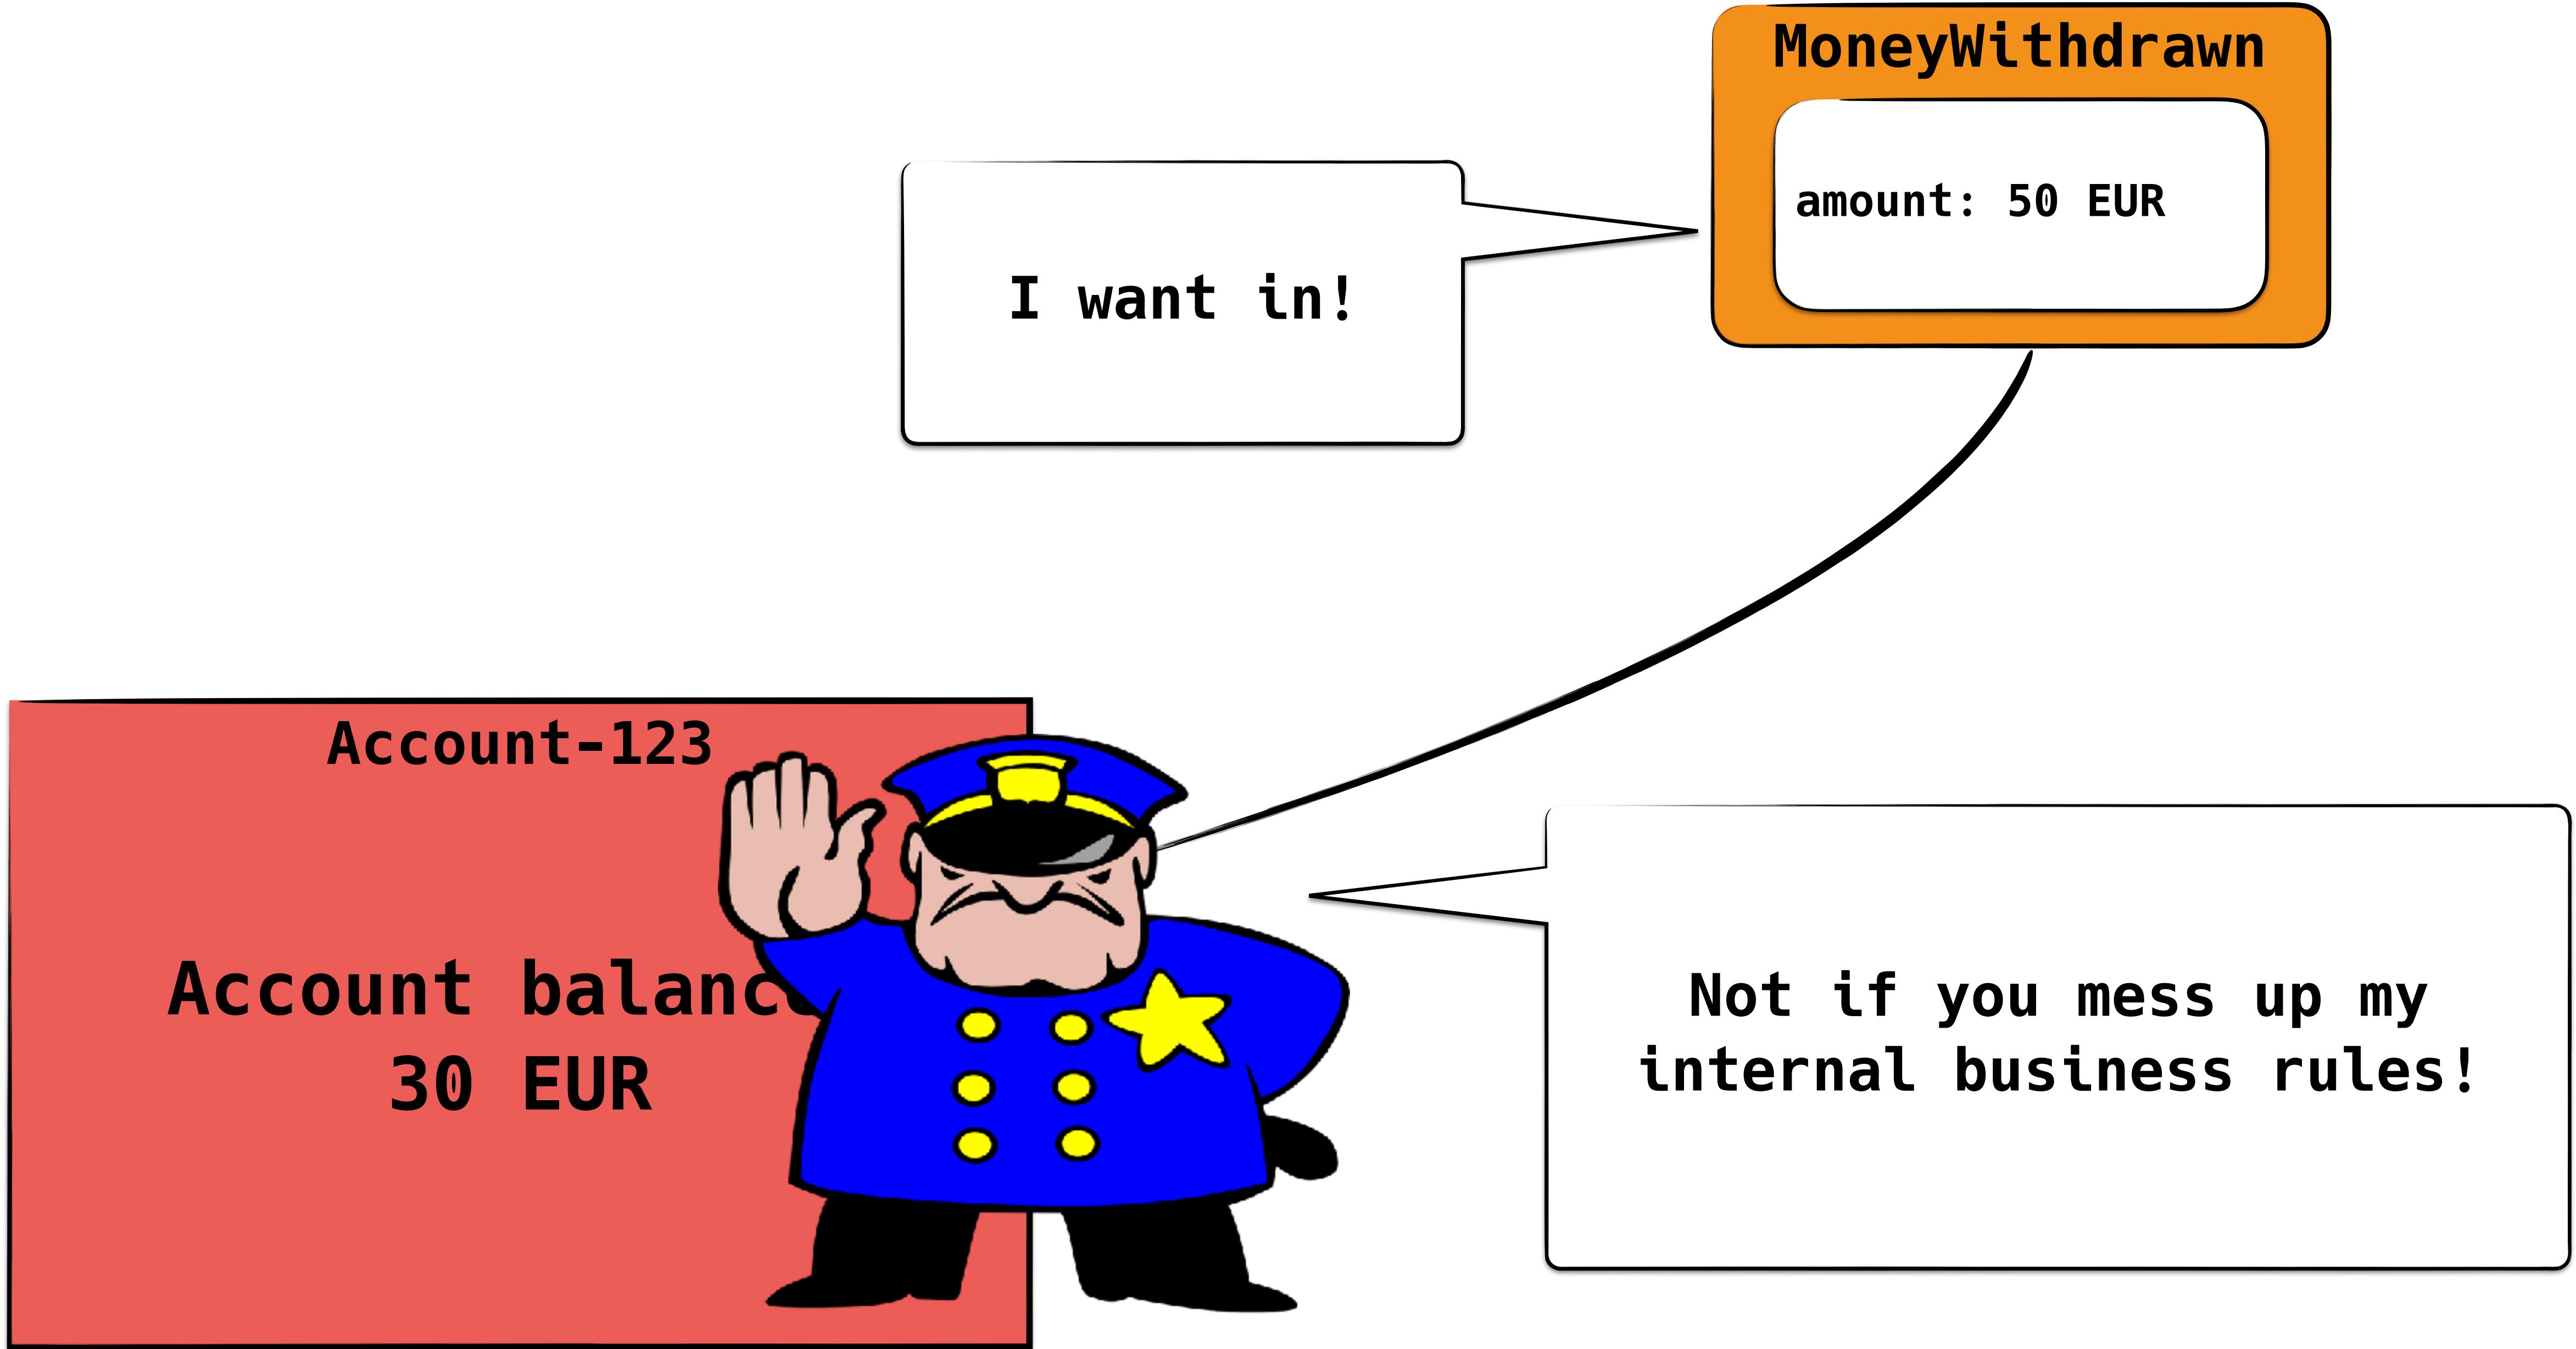
MoneyDeposited

amount: 50 EUR

43 EUR

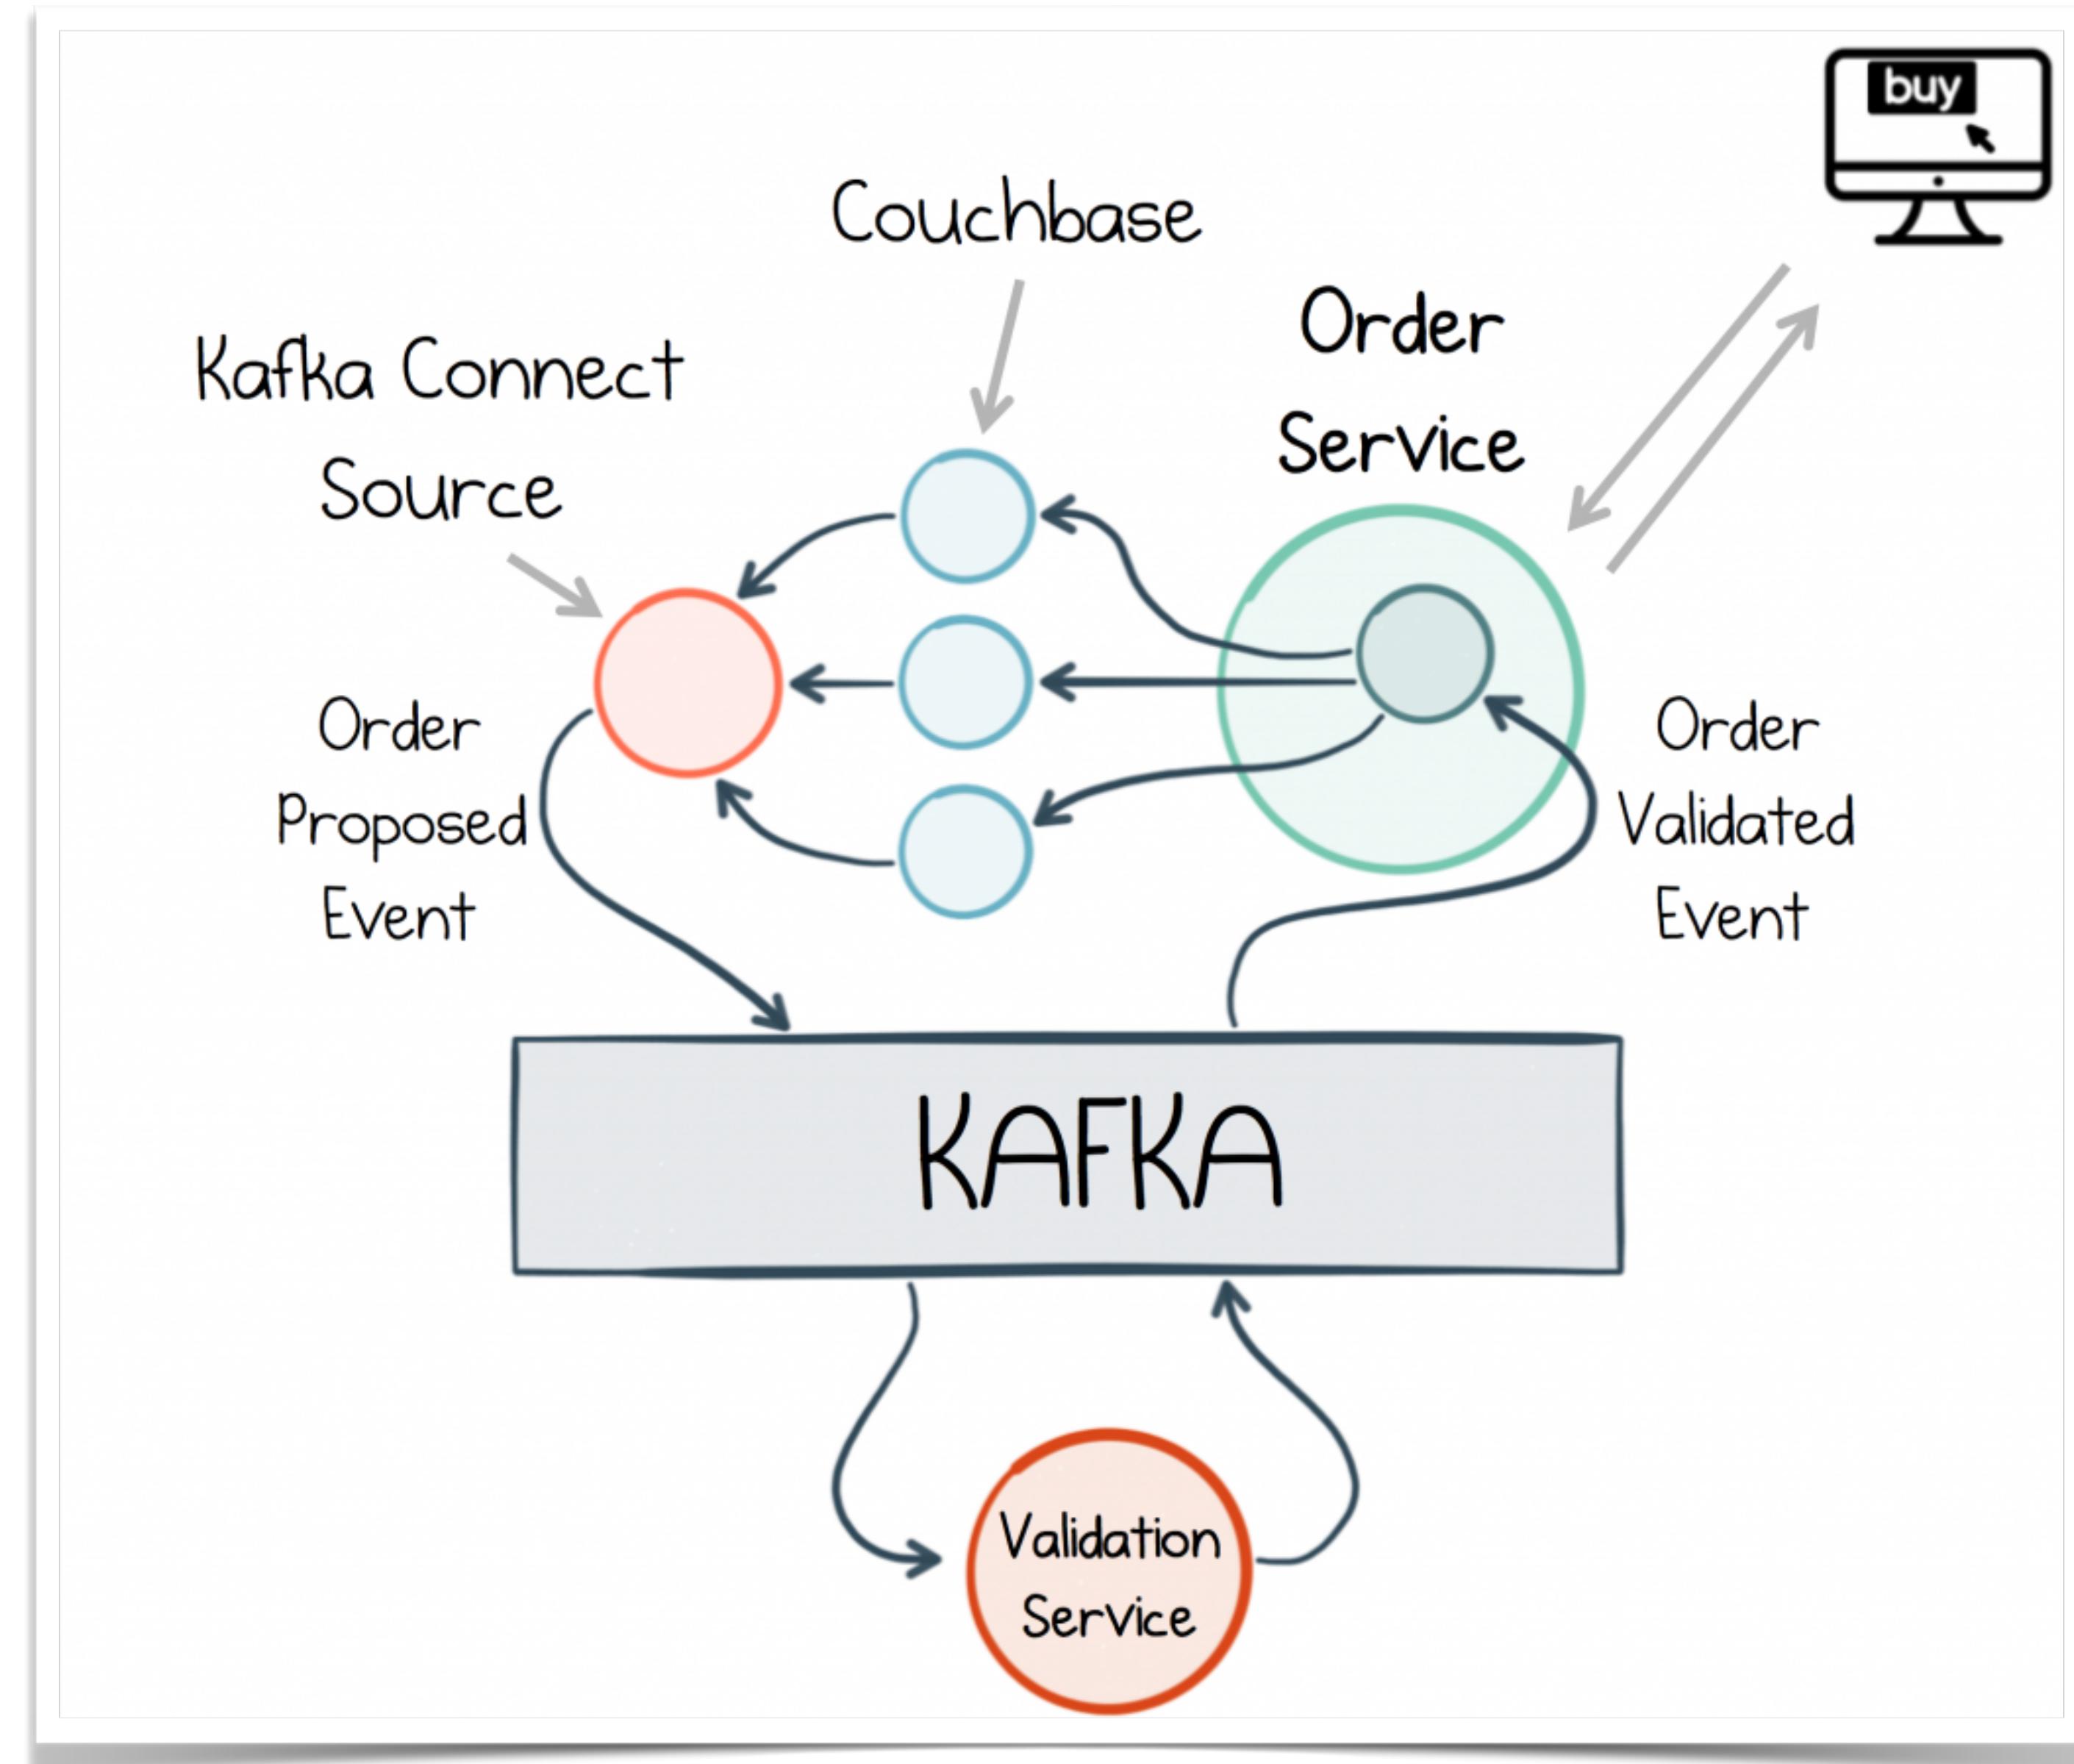


The aggregate is responsible for
enforcing **business invariants**



**“Just” add a database
in front of your
eventstore!**





<https://www.confluent.io/blog/messaging-single-source-truth/>

@koenighotze

Quick tip for finding friends in ops:

Ask them to “just” install and maintain
production ready Kafka and
Couchbase installations in the Cloud

Optimistic concurrency control

From Wikipedia, the free encyclopedia

Optimistic concurrency control (OCC) is a [concurrency control](#) method applied to transactional systems such as [relational database management systems](#) and [software transactional memory](#). OCC assumes that multiple transactions can frequently complete without interfering with each other. While running, transactions use data resources without acquiring locks on those resources. Before committing, each transaction verifies that no other transaction has modified the data it has read. If the check reveals conflicting modifications, the committing transaction rolls back and can be restarted.^[1] Optimistic concurrency control was first proposed by [H.T. Kung](#) and John T. Robinson.^[2]

https://en.wikipedia.org/wiki/Optimistic_concurrency_control

Optimistic concurrency control

From Wikipedia, the free encyclopedia

OCC assumes that multiple transactions can frequently complete without interfering with each other.

transaction has modified the data it has read. If the check reveals conflicting modifications, the committing transaction rolls back and can be restarted.^[1] Optimistic concurrency control was first proposed by H.T. Kung and John T. Robinson.^[2]

https://en.wikipedia.org/wiki/Optimistic_concurrency_control

Optimistic concurrency control

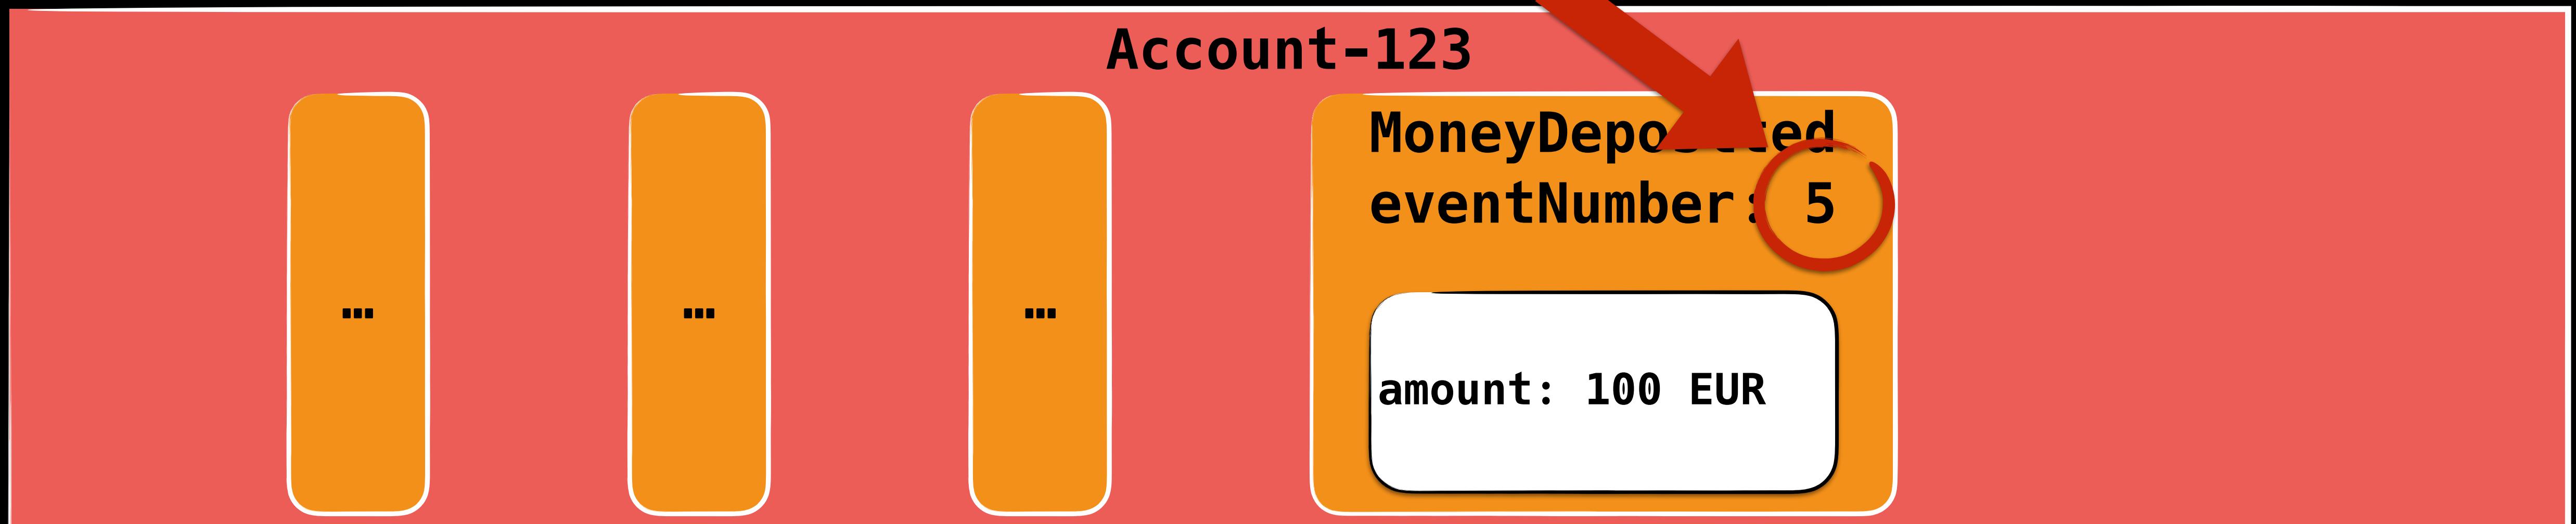
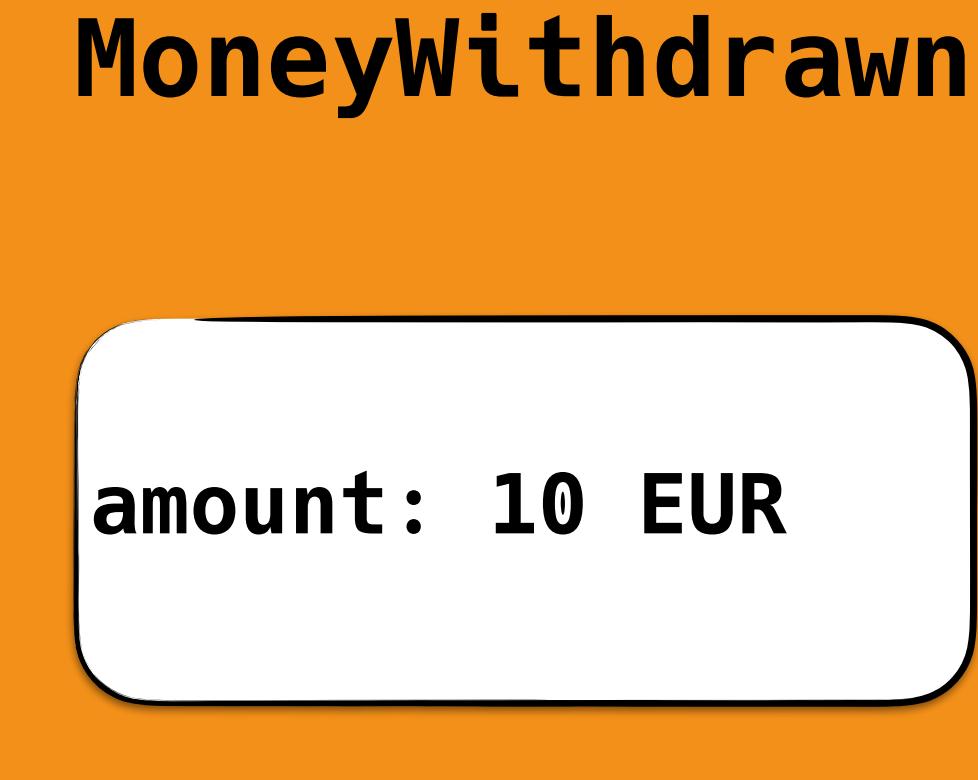
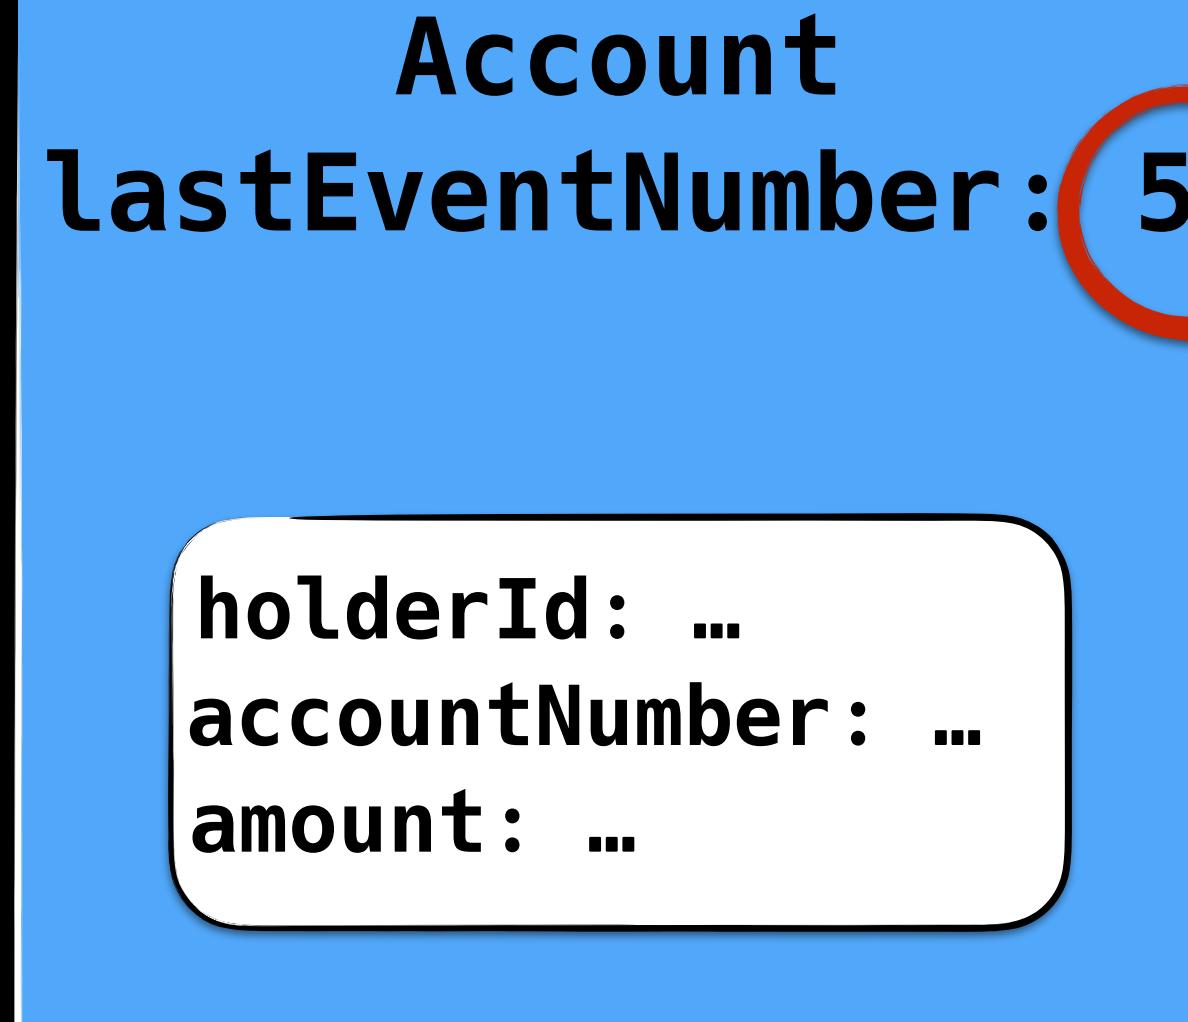
From Wikipedia, the free encyclopedia

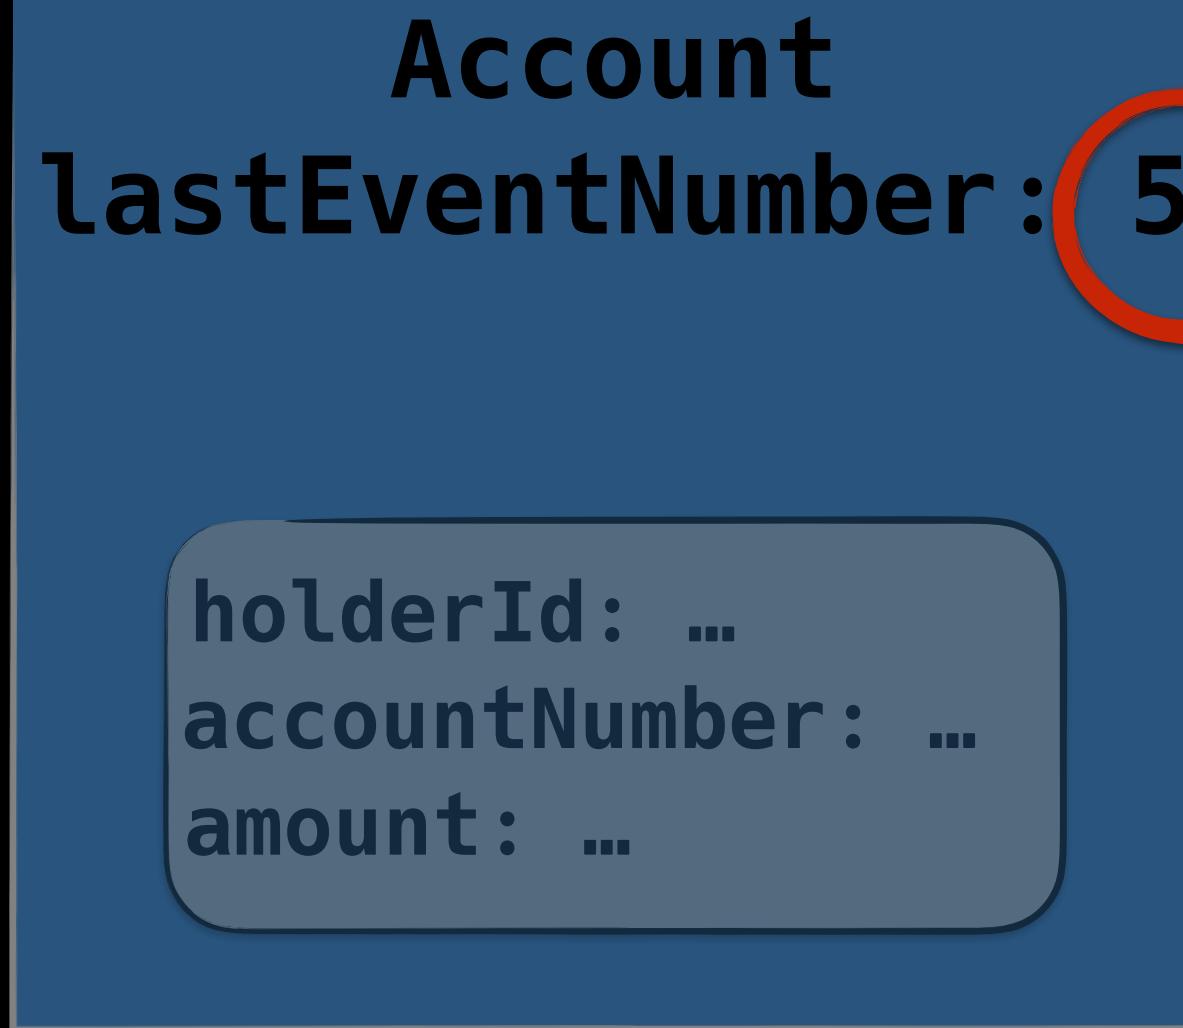
..each transaction verifies that no other transaction has modified the data it has read...

transaction has modified the data it has read. If the check reveals conflicting modifications, the committing transaction rolls back and can be restarted.^[1] Optimistic concurrency control was first proposed by H.T. Kung and John T. Robinson.^[2]

https://en.wikipedia.org/wiki/Optimistic_concurrency_control

The happy path





MoneyWithdrawn

amount: 10 EUR

Account.lastEventNumber(5)

==

Stream.lastEventNumber(5)



Account

lastEventNumber: 5

holderId: ...

accountNumber: ...

amount: ...

MoneyWithdrawn

amount: 10 EUR

Account-123

...

...

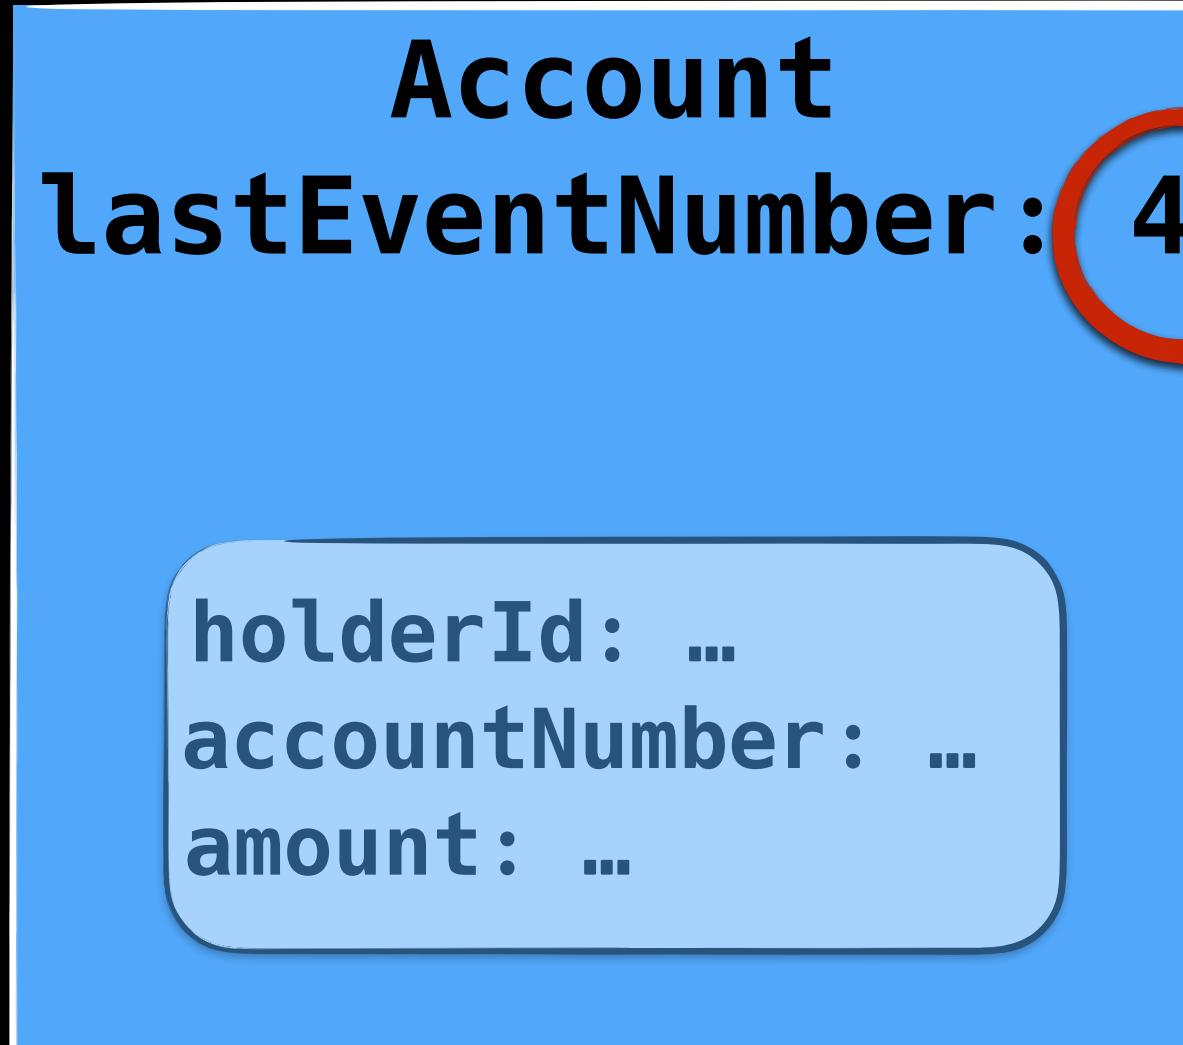
...

MoneyDeposited
eventNumber: 5

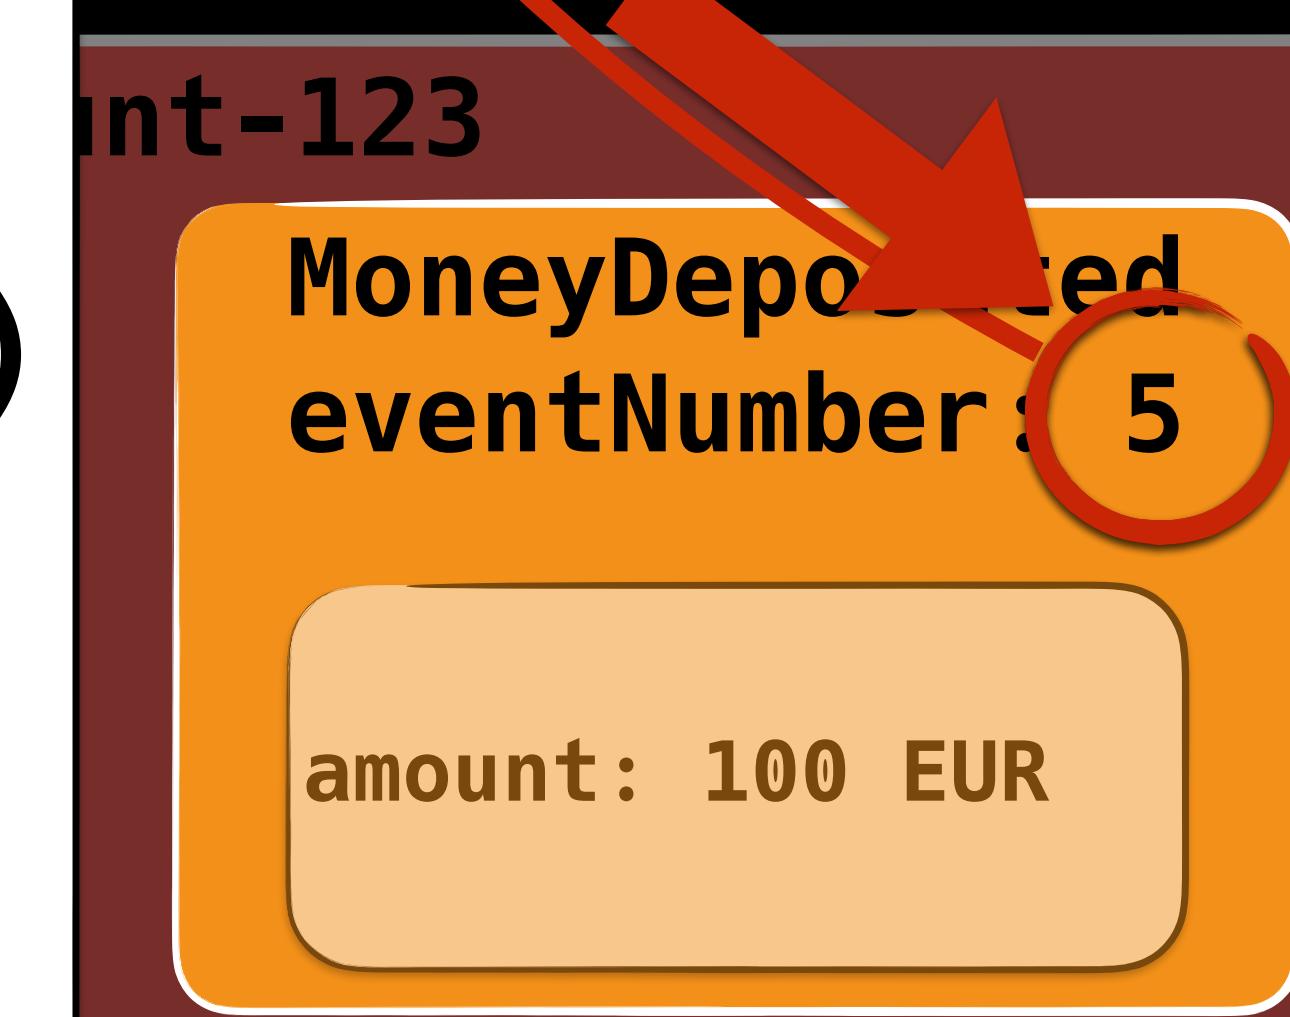
amount: 100 EUR

eventNumber: 6

The not-so-happy path



Account.lastEventNumber(4)
==
Stream.lastEventNumber(5)

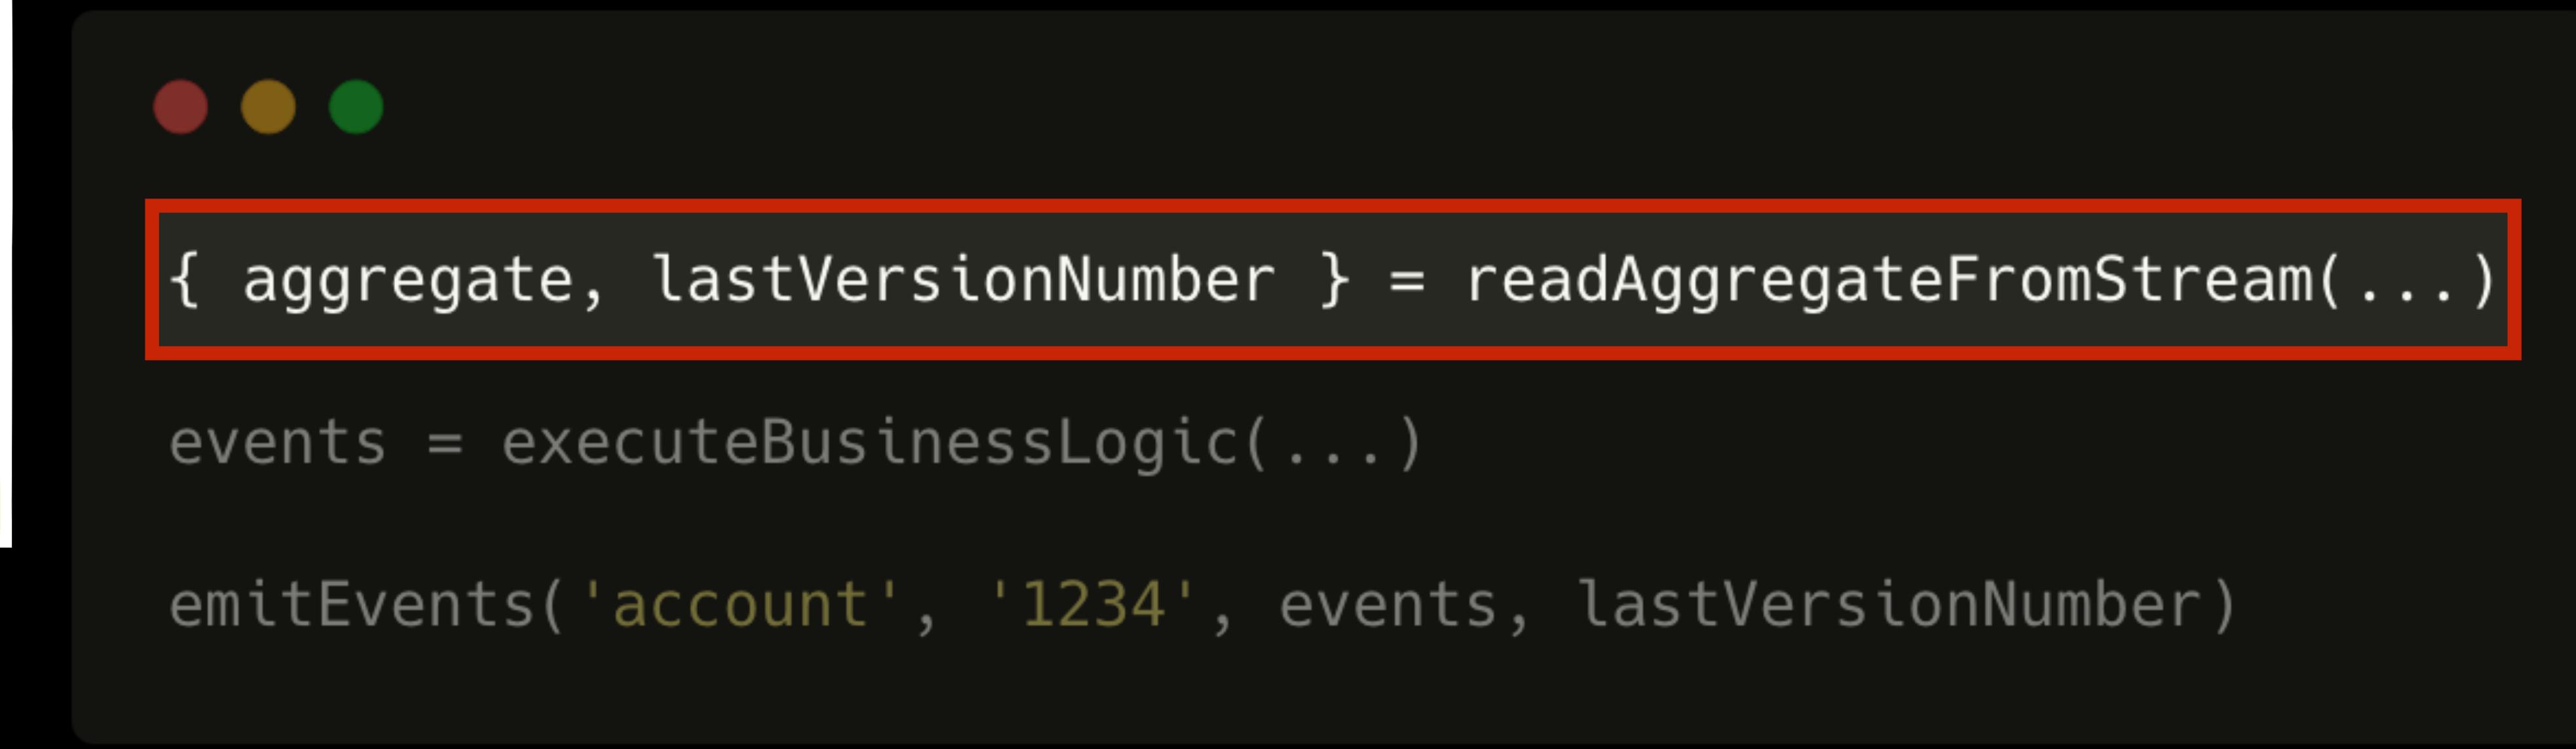


PUT /account/1234



```
{ aggregate, lastVersionNumber } = readAggregateFromStream( ... )  
  
events = executeBusinessLogic( ... )  
  
emitEvents( 'account', '1234', events, lastVersionNumber )
```

PUT /account/1234



PUT /account/1234



```
{ aggregate, lastVersionNumber } = readAggregateFromStream( ... )
```

```
events = executeBusinessLogic( ... )
```

```
emitEvents( 'account', '1234', events, lastVersionNumber )
```



And Kafka?



Kafka / KAFKA-2260

Allow specifying expected offset on produce

Details

Type:	Improvement	Status:	OPEN
Priority:	Minor	Resolution:	Unresolved
Affects Version/s:	None	Fix Version/s:	None
Component/s:	producer		
Labels:	None		

Description

I'd like to propose a change that adds a simple CAS-like mechanism to the Kafka producer. This update has a small footprint, but enables a bunch of interesting uses in stream processing or as a commit log for process state.

▼ [Andy Bryant](#) added a comment - 27/Jul/18 04:14

👍 Would prove very handy in event source based designs

▼ [Russell Ferriday](#) added a comment - 9 hours ago

This would enable full-on eventsourcing on Kafka, without having to restrict to single-thread designs.

One example of a great (>250 github star) FOSS project being held back by this:

<https://github.com/johnbywater/eventsourcing/issues/108>

Can we see this soon?



Kafka / KAFKA-2260

Allow specifying expected offset on produce

Details

Type:

Improvement

Status:

OPEN

Priority:

Minor

Resolution:

Unresolved

Affects Version/s:

None

Fix Version/s:

None



... added a comment - 27/Jul/18 04:14

👍 Would prove very handy in event source based designs

Andy Bryant added a comment - 27/Jul/18 04:14

👍 Would prove very handy in event source based designs

Russell Ferriday added a comment - 9 hours ago

This would enable full-on eventsourcing on Kafka, without having to restrict to single-thread designs.

One example of a great (>250 github star) FOSS project being held back by this:

<https://github.com/johnbywater/eventsourcing/issues/108>

Can we see this soon?

@koenighotze



Kafka / KAFKA-2260

Allow specifying expected offset on produce

Details

Type:

Improvement

Status:

OPEN

Priority:

Minor

Resolution:

Unresolved

Affects Version/s:

None

Fix Version/s:

None

Type:

Improvement

Status:

OPEN

Priority:

Minor

Resolution:

Unresolved

Affects Version/s:

None

Fix Version/s:

None

▼ Andy Bryant added a comment - 27/Jul/18 04:14

👍 Would prove very handy in event source based designs

▼ Russell Ferriday added a comment - 9 hours ago

This would enable full-on eventsourcing on Kafka, without having to restrict to single-thread designs.

One example of a great (>250 github star) FOSS project being held back by this:

<https://github.com/johnbywater/eventsourcing/issues/108>

Can we see this soon?

Advantages of OCC?

Scalability and no locks
Consistency
Design choice
Super-simple programming logic

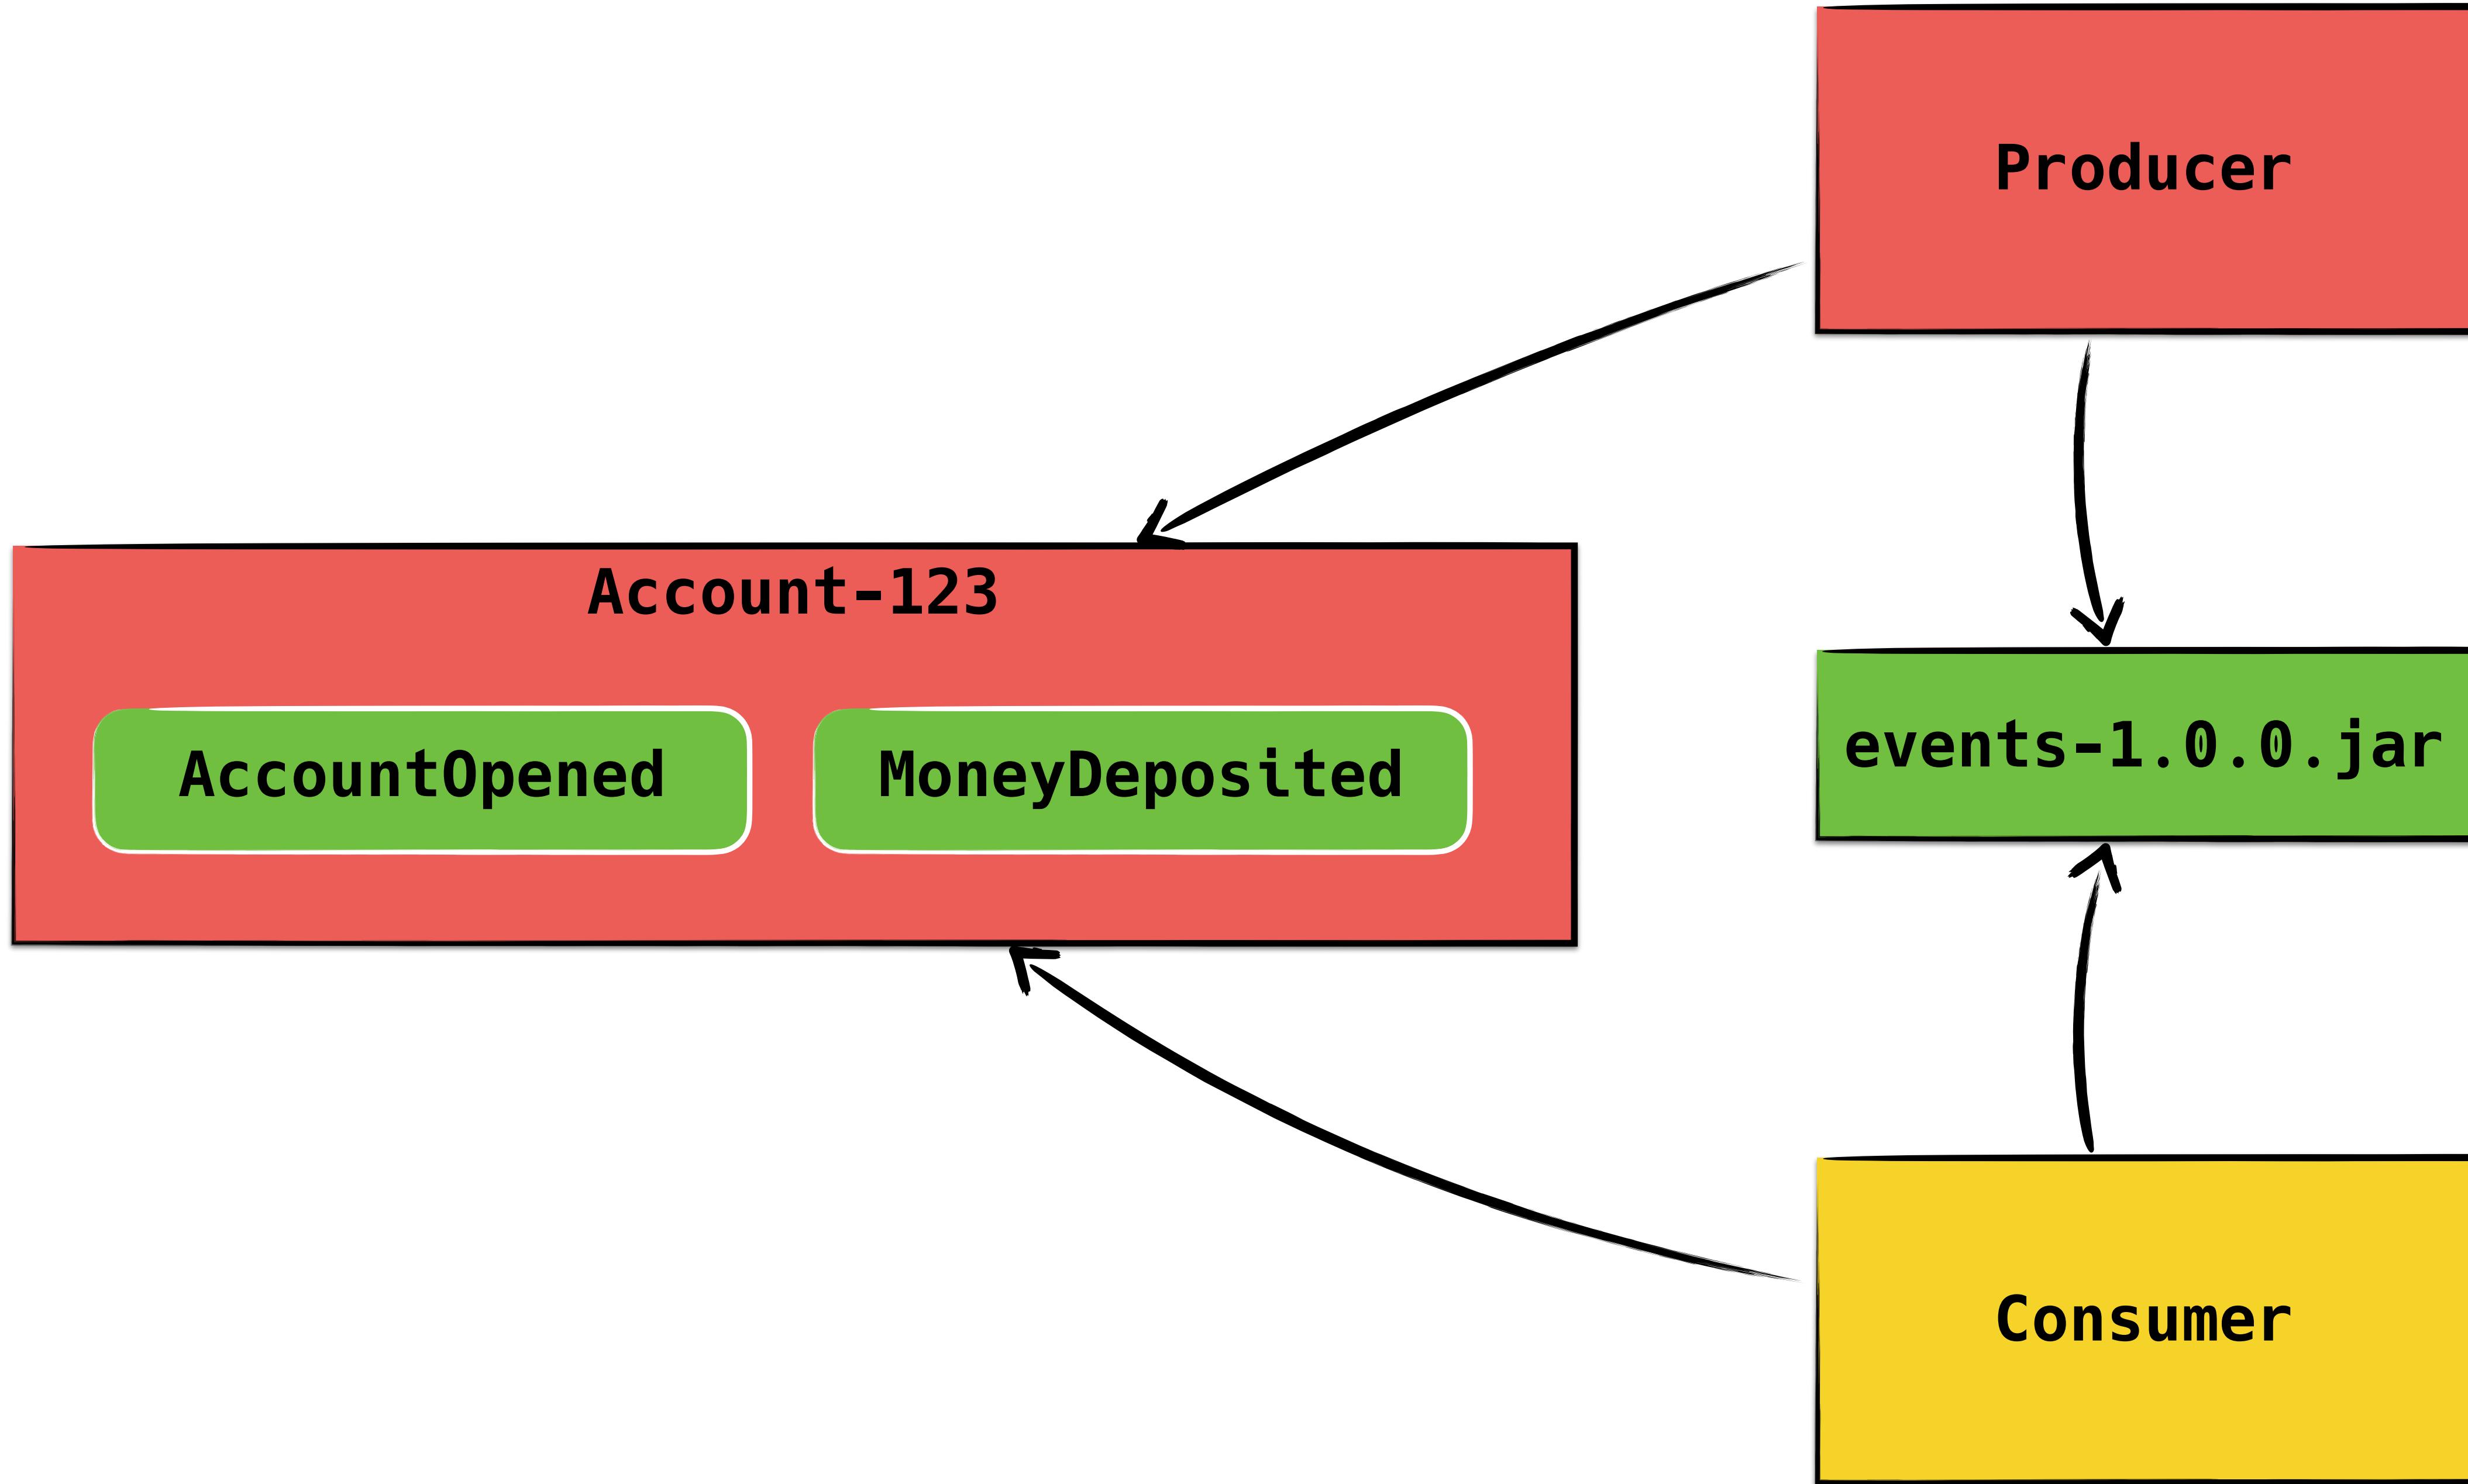


Versions, up-front-design and
breaking things down the road

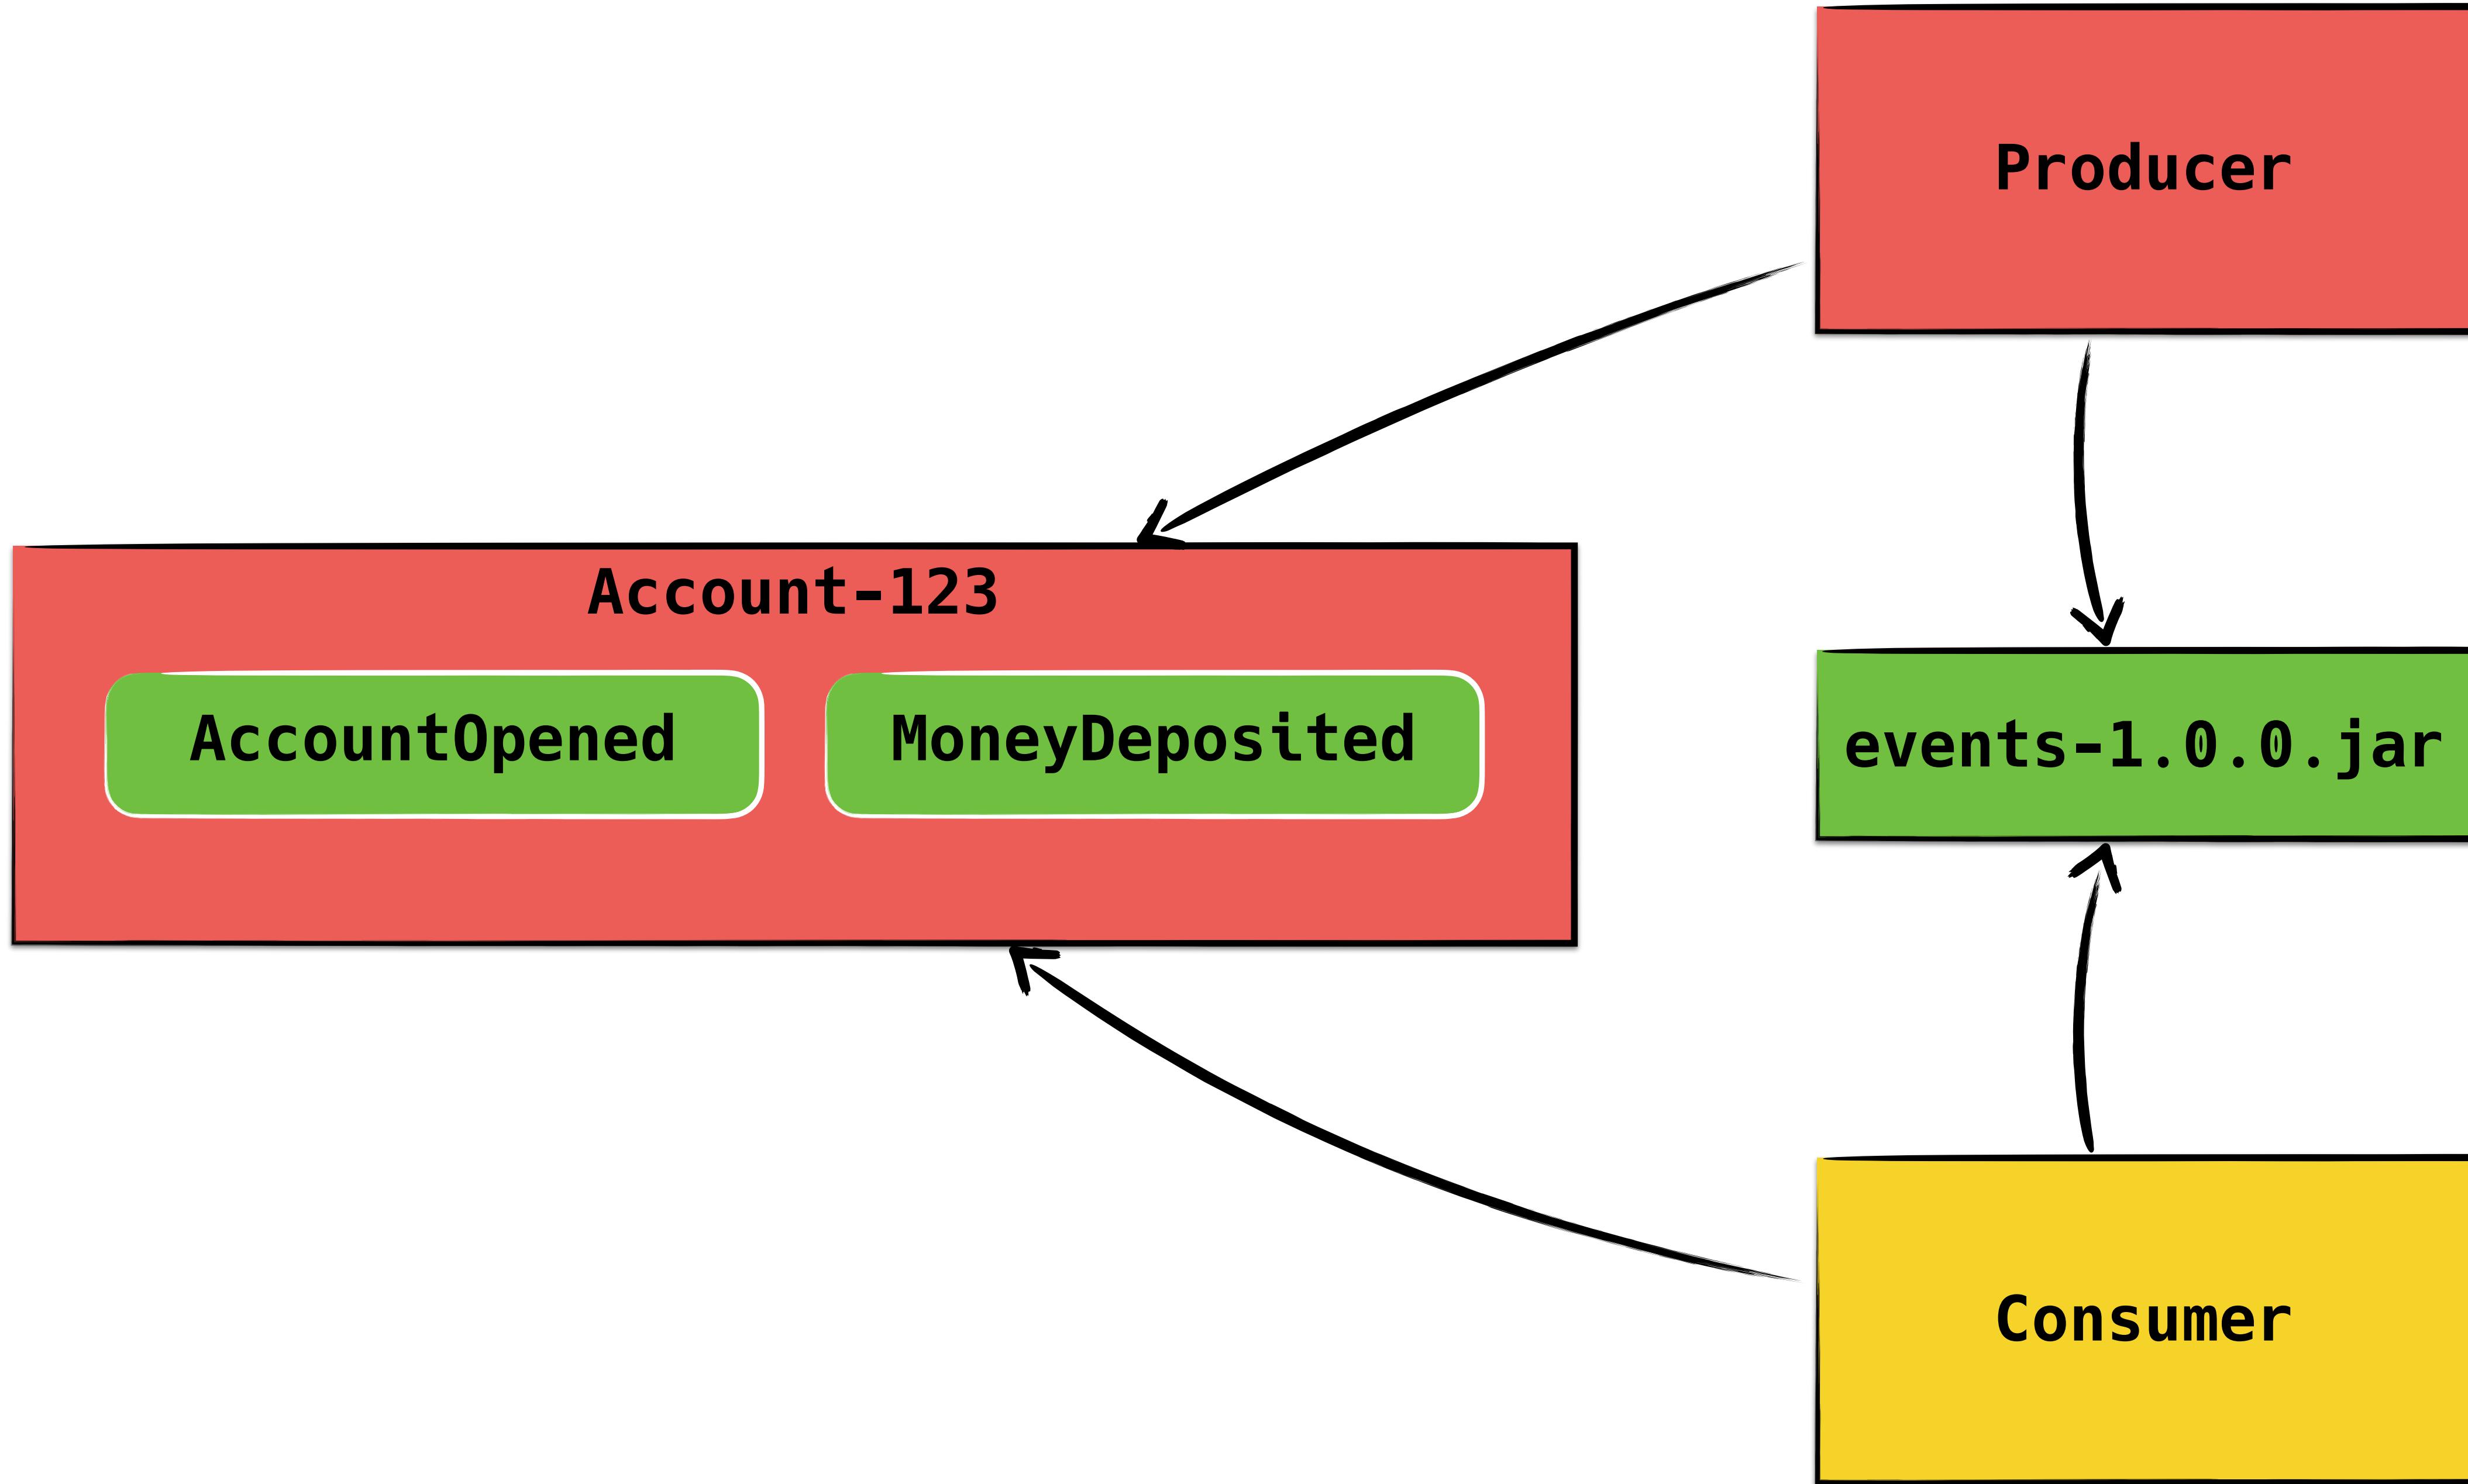
How can we deal with
versions without going crazy?

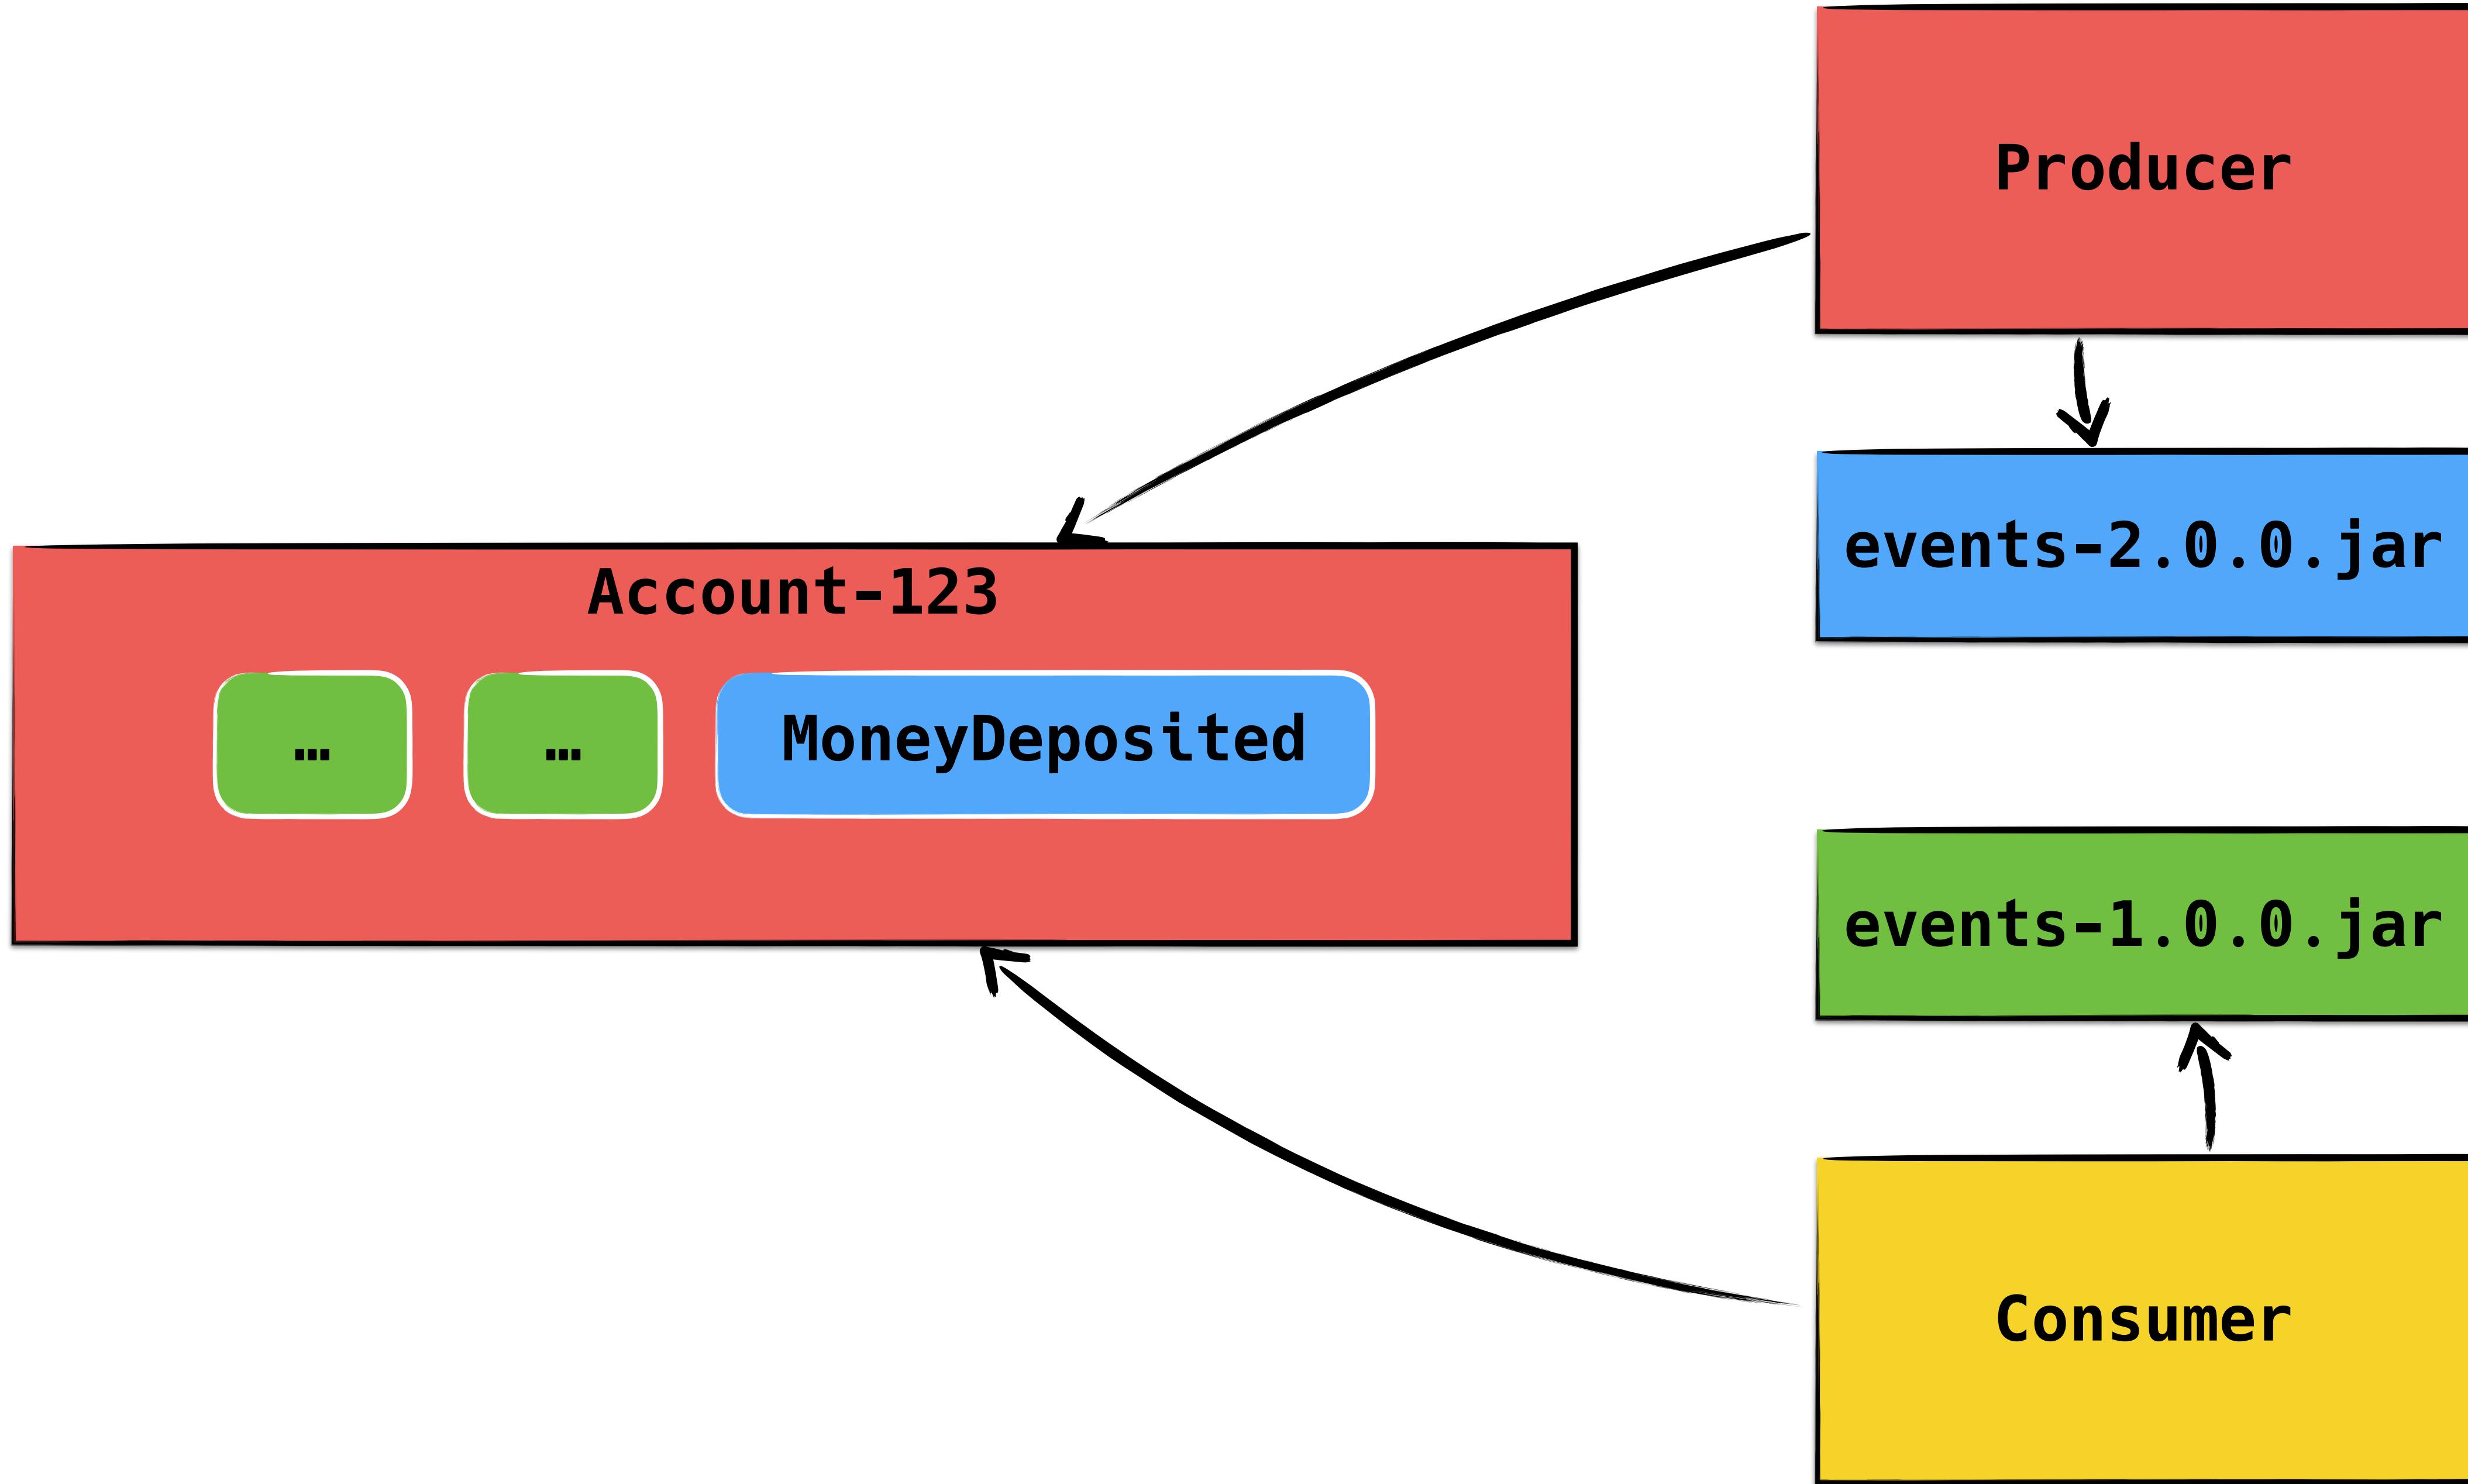
“Just” use semantic
versioning and types





...then change some event types







```
public class InvalidClassException extends ObjectOutputStream {  
    ...  
}
```

Gosh... “just” apply
Double Write



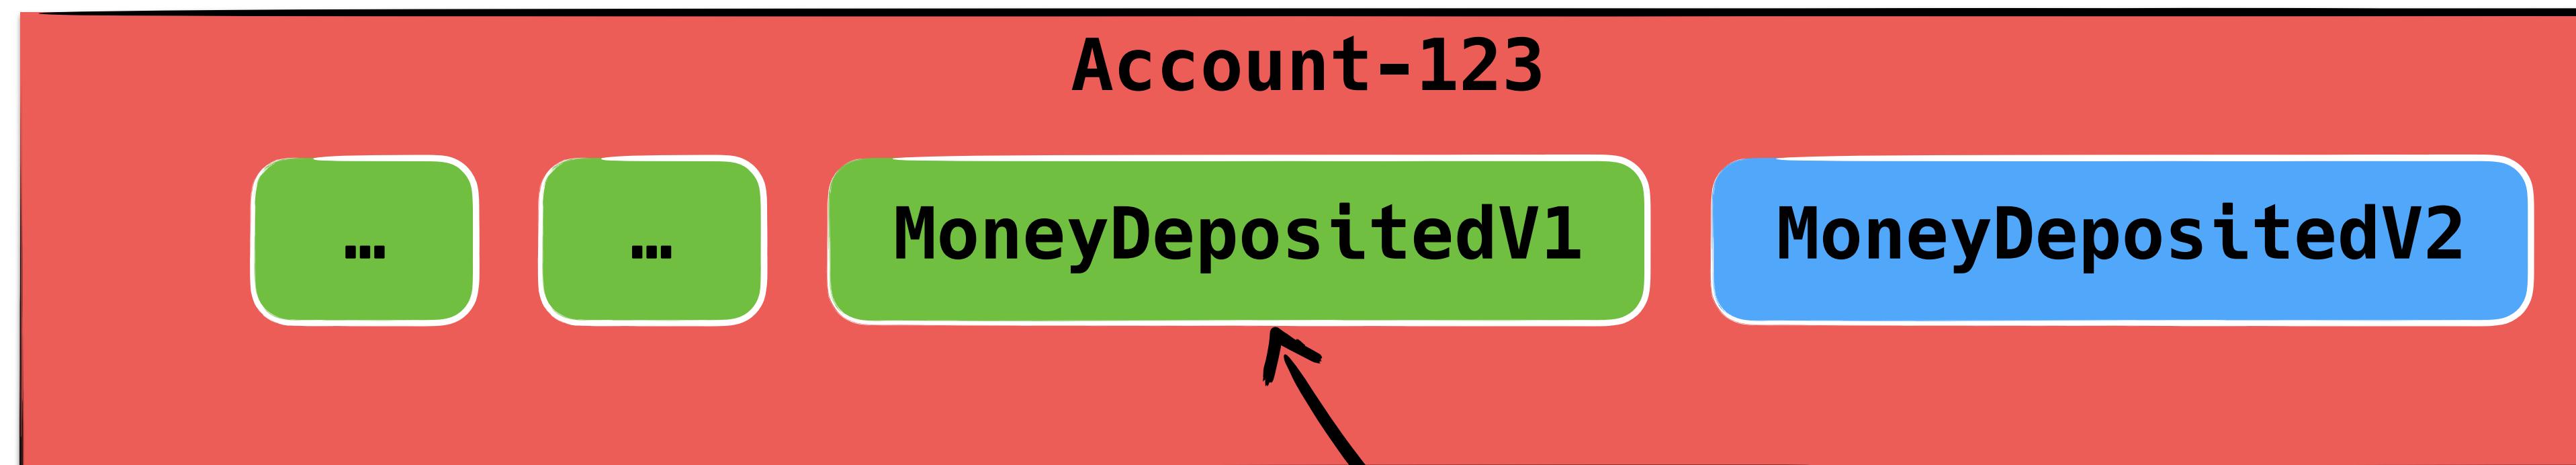
Account-123

...

...

MoneyDepositedV1

MoneyDepositedV2



Should I process V1?





Or wait for a V2...which
might never arrive?



MoneyDeposited_v1

MoneyDeposited_v2

...

MoneyDeposited_v100

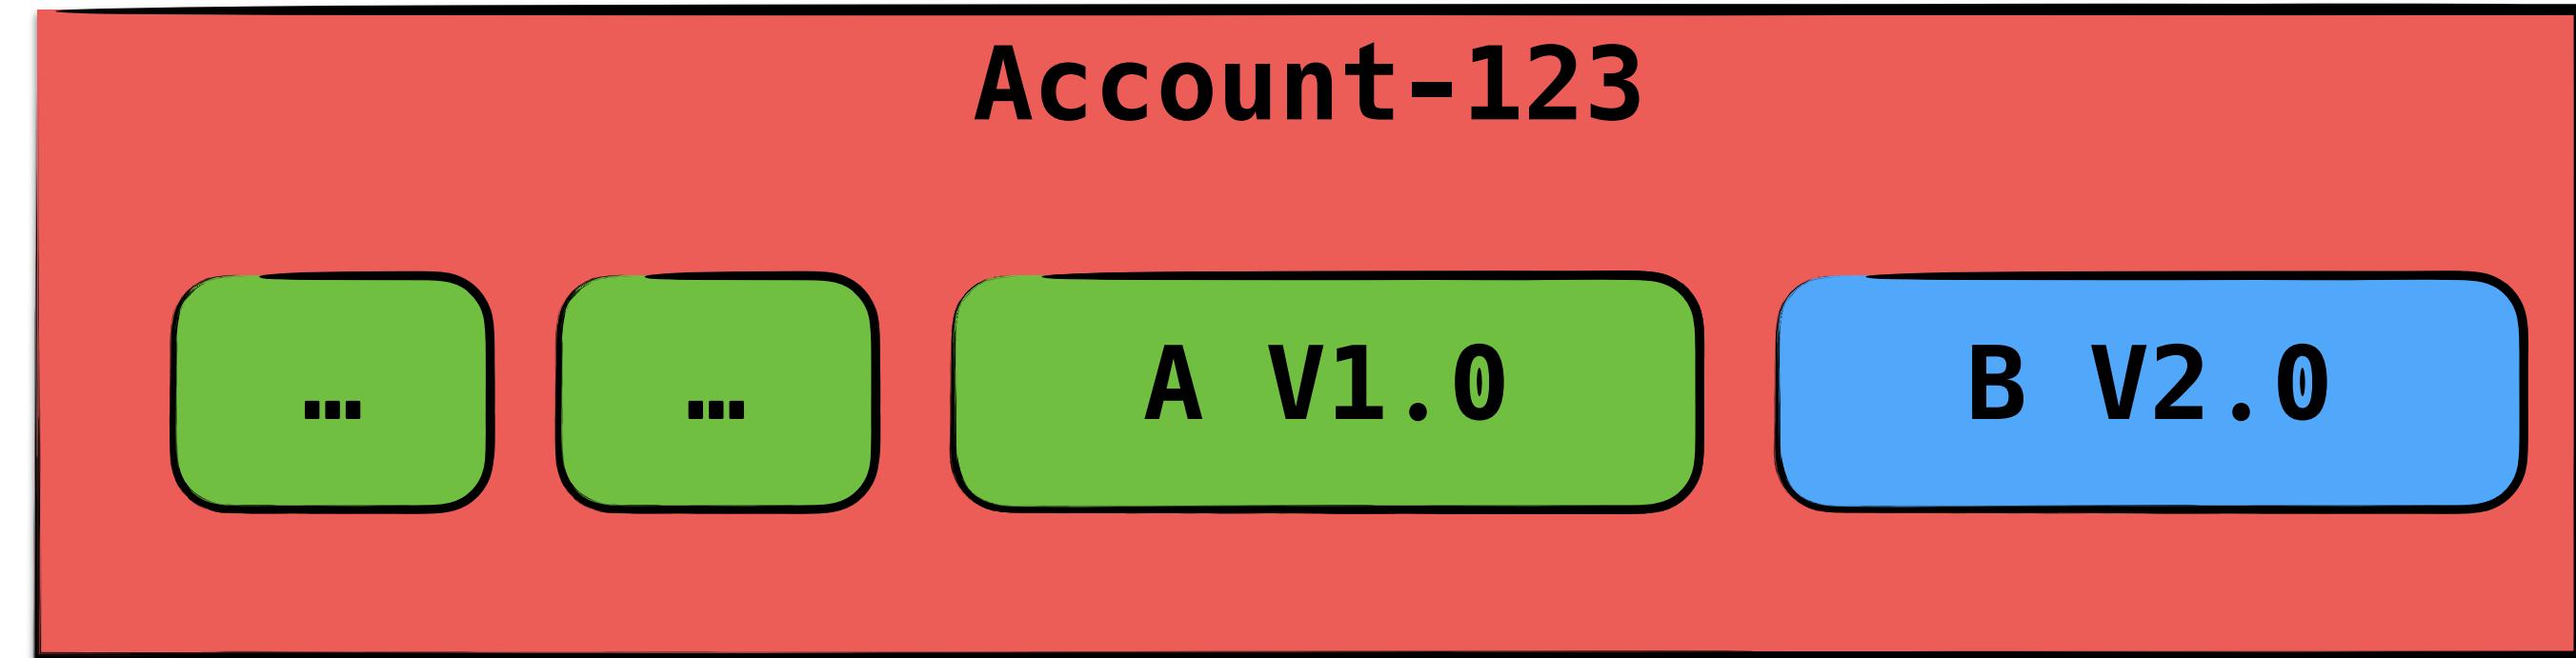
MoneyDeposited_v1Handler

MoneyDeposited_v2Handler

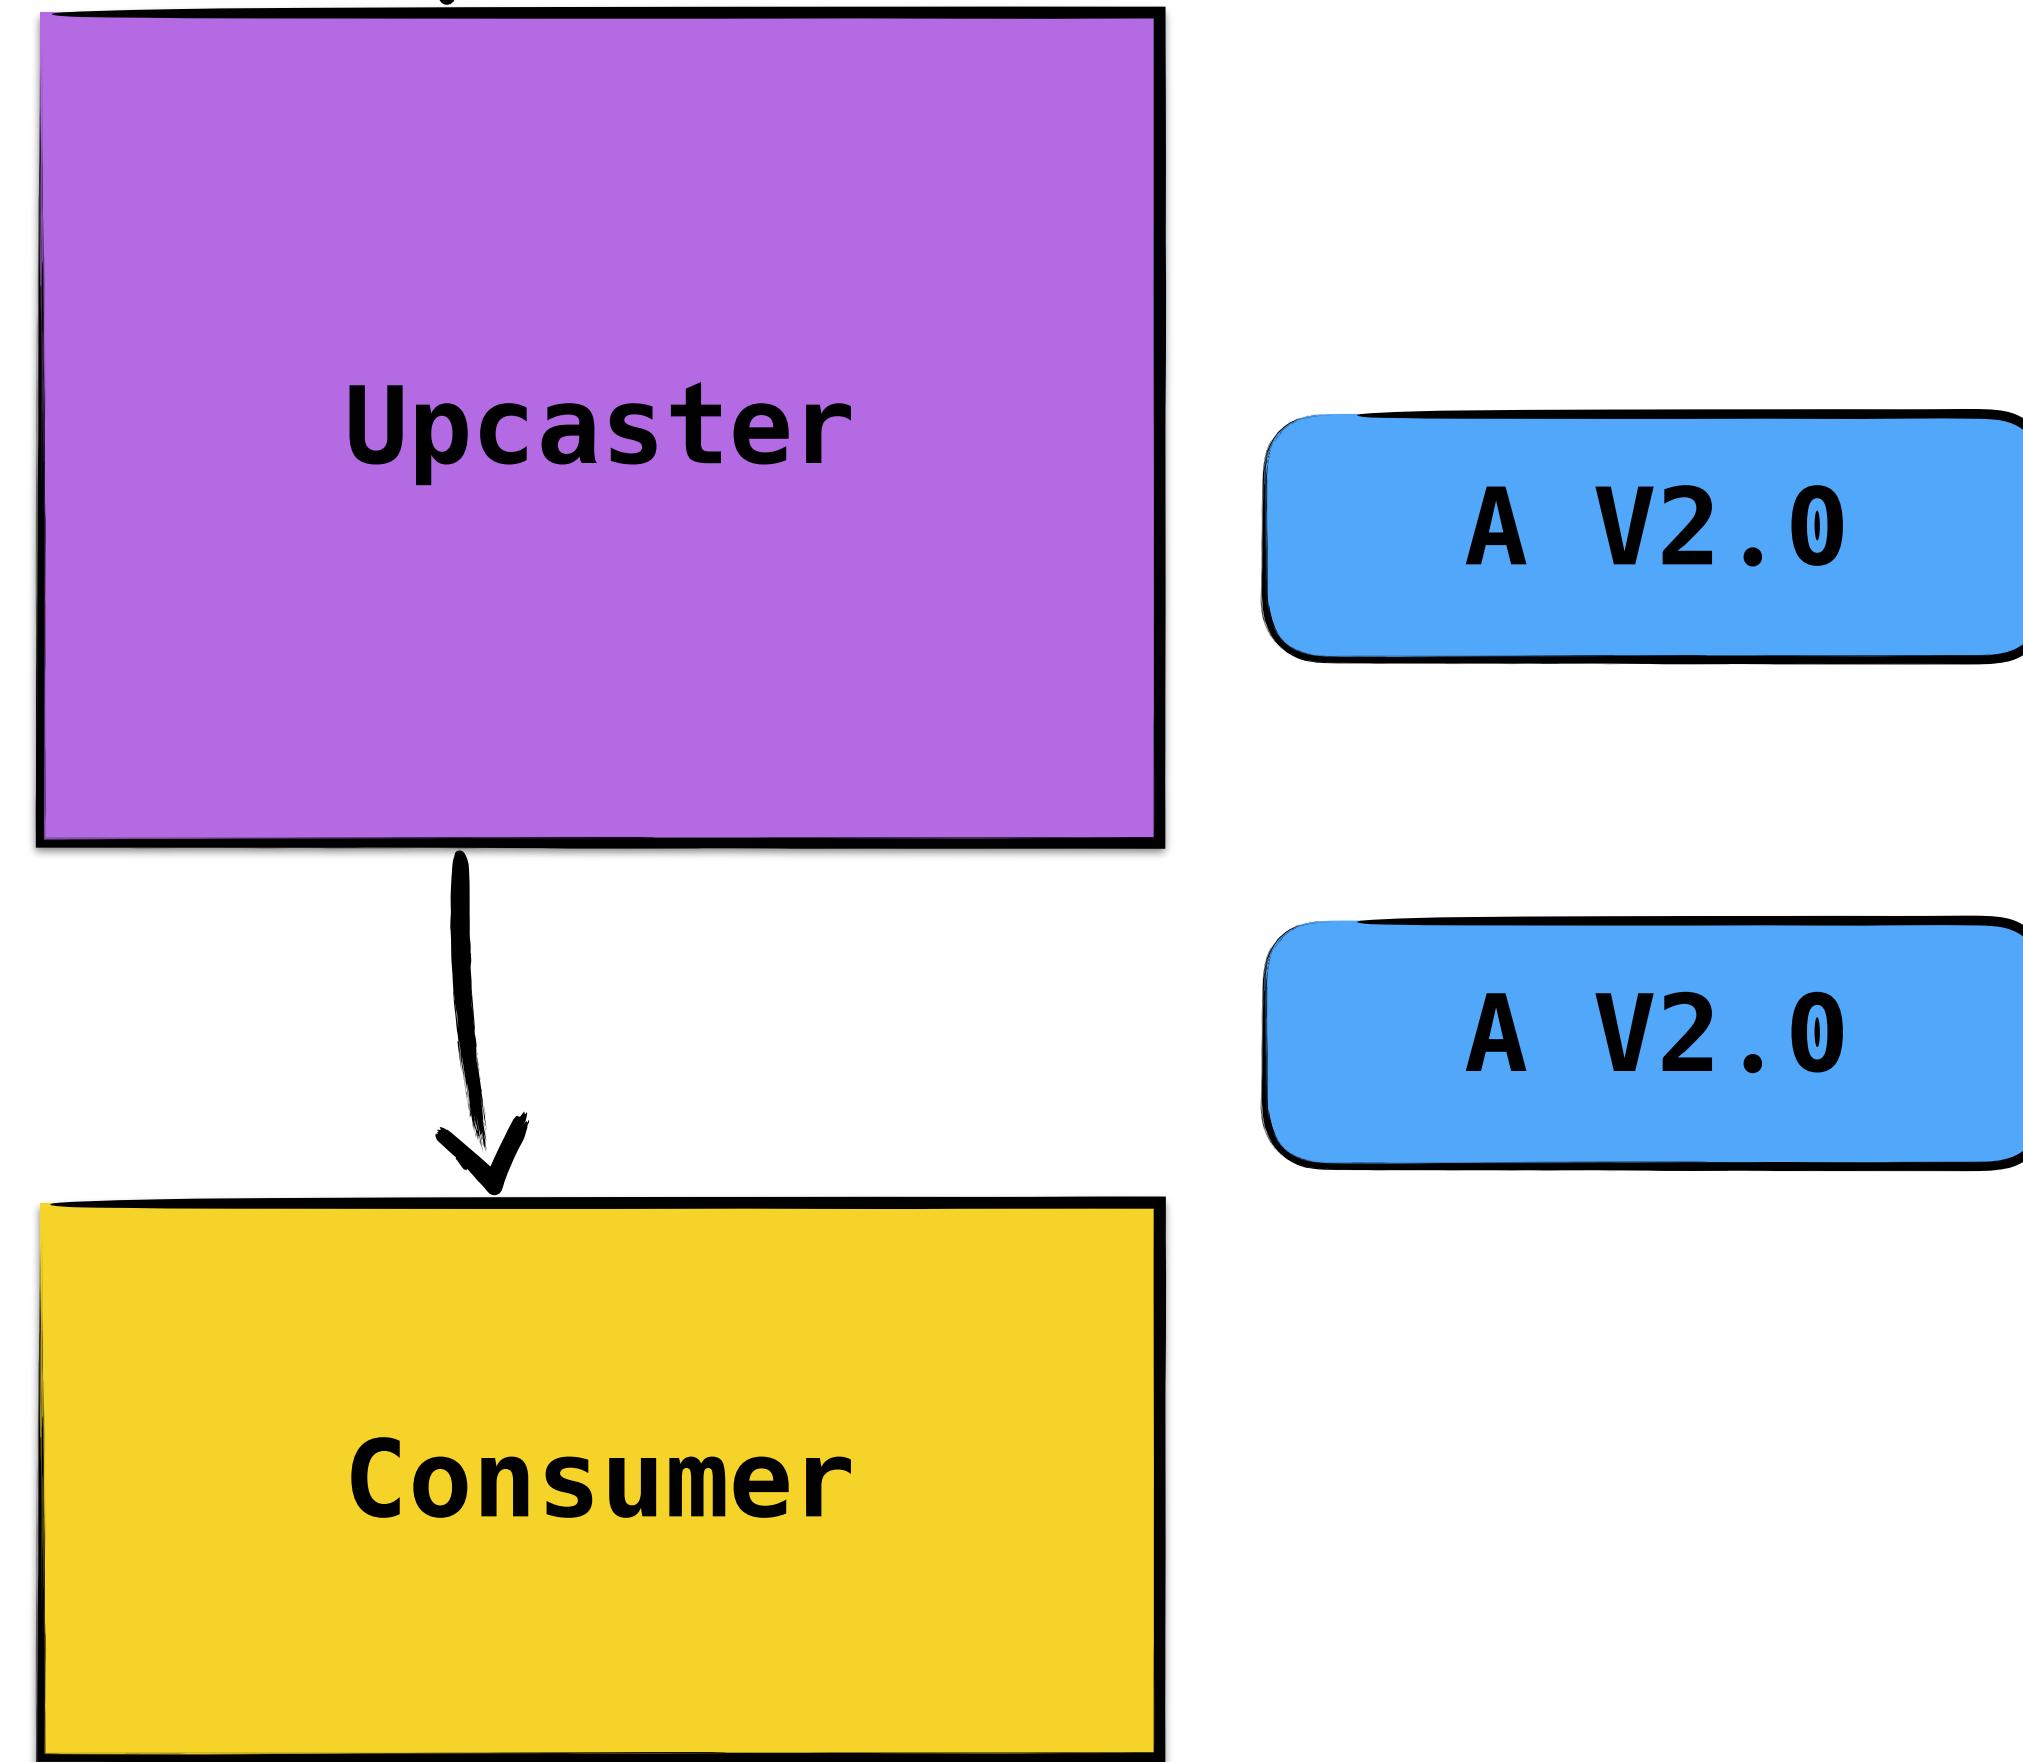
...

MoneyDeposited_v100Handler

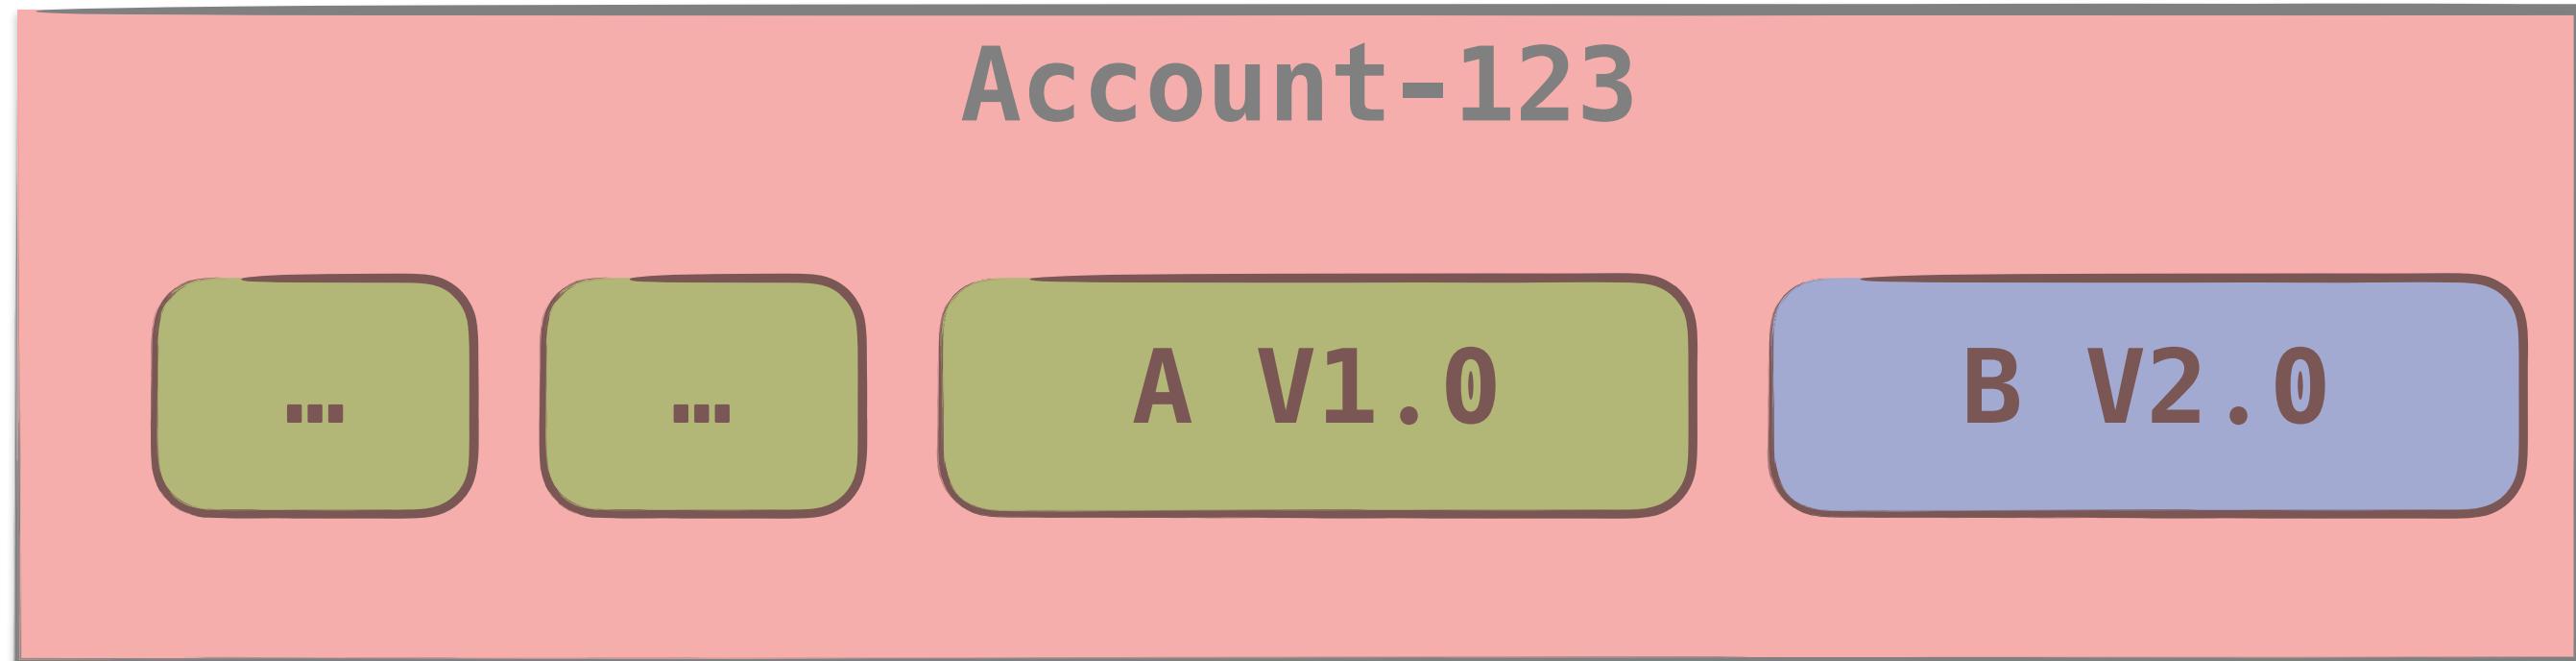
Upcaster



$f: v1 \rightarrow v2$



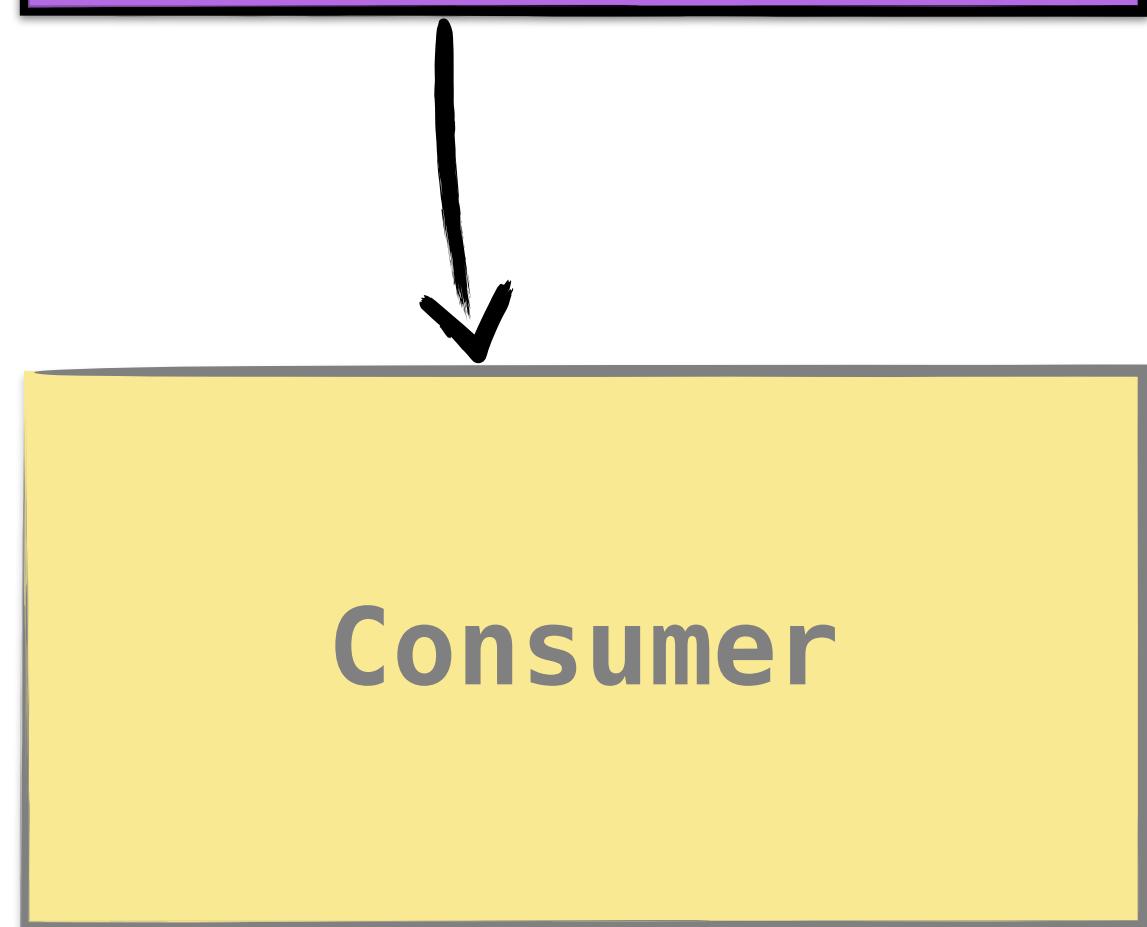
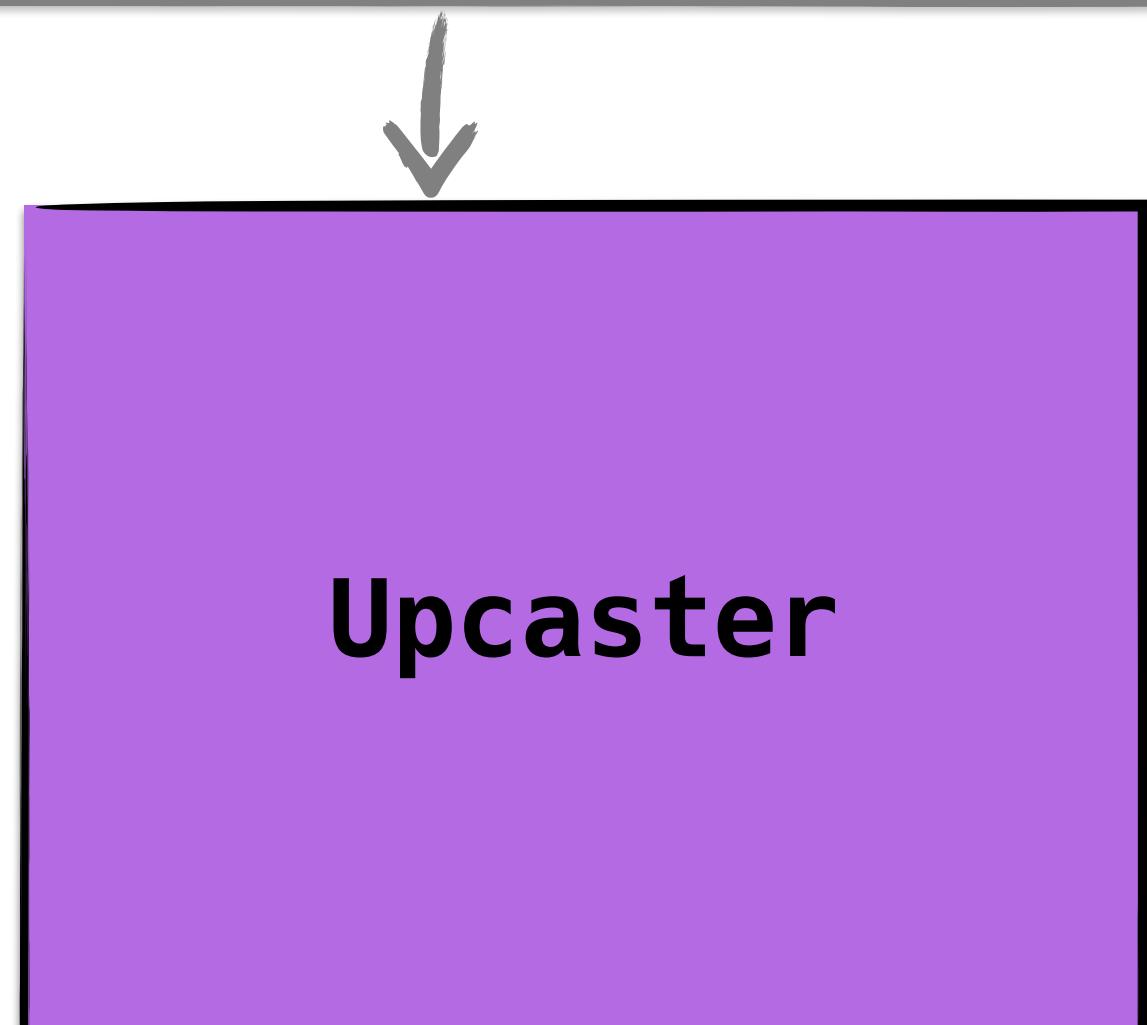
5 months later...



f: v1->v2
f: v2->v3
f: v3->v4

...

f: v97->v98
f: v98->v99
f: v99->v100



Good luck maintaining that monster

Prefer simple, text-based,
human readable events

Fancy speak for JSON



```
{  
  "eventType": "MoneyTransferred",  
  "aggregateId": "1234",  
  "iban": "DE12",  
  "accountNumber": "12312312",  
  "amount": 10,  
  "currency": "EUR"  
}
```

And correctness?

**“Just” generate
classes for JSON
mapping!**





```
public class UserCreatedEvent {  
    private UUID requestId;  
  
    ...  
}
```

NGINX \$request_id

unique request identifier
generated from 16 random bytes,
in hexadecimal (1.11.0)

“Oh, you changed the request id
from **uuid** to any **arbitrary string**”

String-ly typed events work really well

Weak schema to the rescue



```
{  
  "$schema": "http://json-schema.org/draft-07/schema",  
  "title": "UserCreated",  
  "description": "Creates a user",  
  "type": "object",  
  "properties": {  
    "userId": {  
      "description": "The new user's ID",  
      "type": "string",  
      "format": "uuid"  
    }  
  },  
  "additionalProperties": false,  
  "required": [  
    "userId"  
  ]  
}
```

```
{  
  "$schema": "http://json-schema.org/draft-07/schema",  
  "title": "UserCreated",  
  "description": "Creates a user",  
  "type": "object",  
  "properties": {  
    "userId": {  
      "description": "The new user's ID",  
      "type": "string",  
      "format": "uuid"  
    }  
  },  
  "additionalProperties": false,  
  "required": [  
    "userId"  
  ]  
}
```



```
{  
  "$schema": "http://json-schema.org/draft-07/schema",  
  "title": "UserCreated",  
  "description": "Creates a user",  
  "type": "object",  
  "properties": {  
    "userId": {  
      "description": "The new user's ID",  
      "type": "string",  
      "format": "uuid"  
    }  
  },  
  "additionalProperties": false,  
  "required": [  
    "userId"  
  ]  
}
```

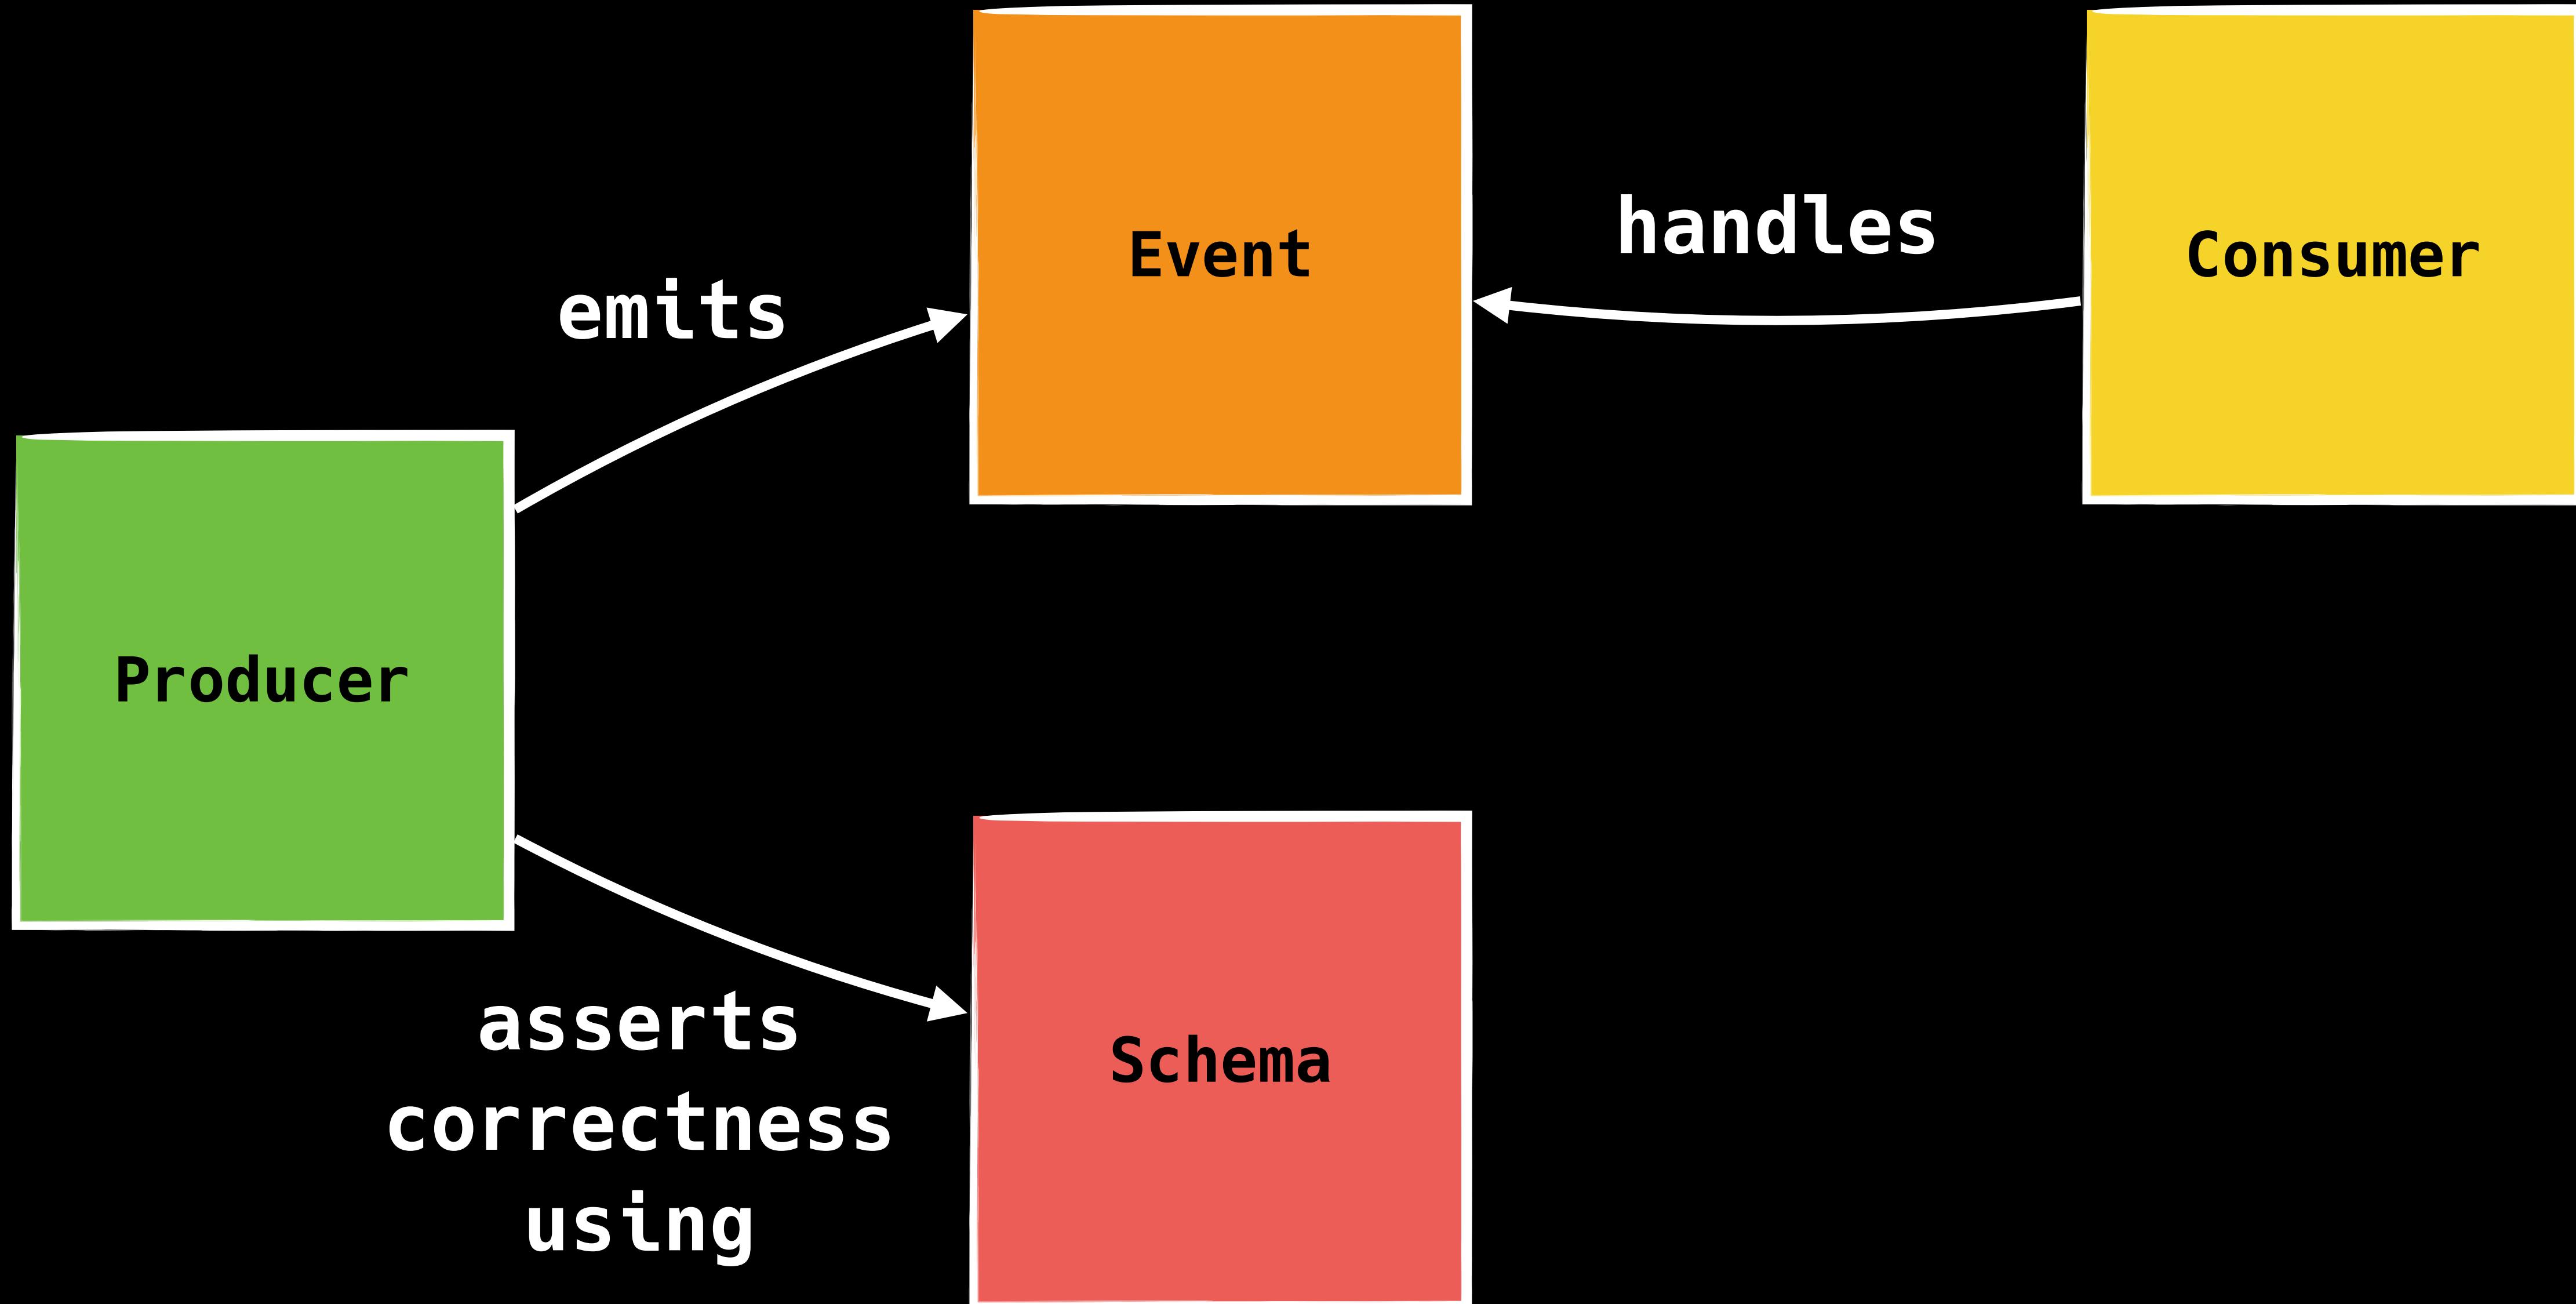


```
assertIsValid(eventData, ajv.compile(schema))
event = newEvent(
{
  aggregateId,
  aggregateType,
  eventData
}
)
```



```
assertIsValid(eventData, ajv.compile(schema))  
event = newEvent(  
{  
  aggregateId,  
  aggregateType,  
  eventData  
}  
)
```

Schema as a **description**
not as a contract



What about putting
versioned logic in handlers?

MoneyTransferred

eventId: 5

amount: 97 USD



```
handleMoneyTransferred( { amount, currency } ) {
    if (currency !== 'EUR') {
        rate = fxCalculator.currentExchangeRate(currency, 'EUR')
        this.transferVolume = rate * amount
    }
    else {
        this.transferVolume = amount
    }
}
```



```
handleMoneyTransferred( { amount, currency } ) {  
    if (currency !== 'EUR') {  
        rate = fxCalculator.currentExchangeRate(currency, 'EUR')  
        this.transferVolume = rate * amount  
    }  
    else {  
        this.transferVolume = amount  
    }  
}
```



```
handleMoneyTransferred( { amount, currency } ) {
    if (currency !== 'EUR') {
        rate = fxCalculator.currentExchangeRate(currency, 'EUR')
        this.transferVolume = rate * amount
    }
    else {
        this.transferVolume = amount
    }
}
```



```
handleMoneyTransferred( { amount, currency } ) {
  if (currency !== 'EUR') {
    rate = fxCalculator.currentExchangeRate(currency, 'EUR')
    this.transferVolume = rate * amount
  }
  else {
    this.transferVolume = amount
  }
}
```

Bravo, now your expense report of 2017
depends on today's exchange rates

Creating an event must **encapsulate all data** that lead to the emitting of the event

MoneyTransferred

eventId: 5

```
amount: 97 USD
exchangeRate": {
  "base": "USD",
  "date": "2018-02-13",
  "rates": { "EUR": 0.806942 }
}
```

Side-effect
manifested as
event payload

MoneyTransferred
eventId: 5

amount: 97 USD

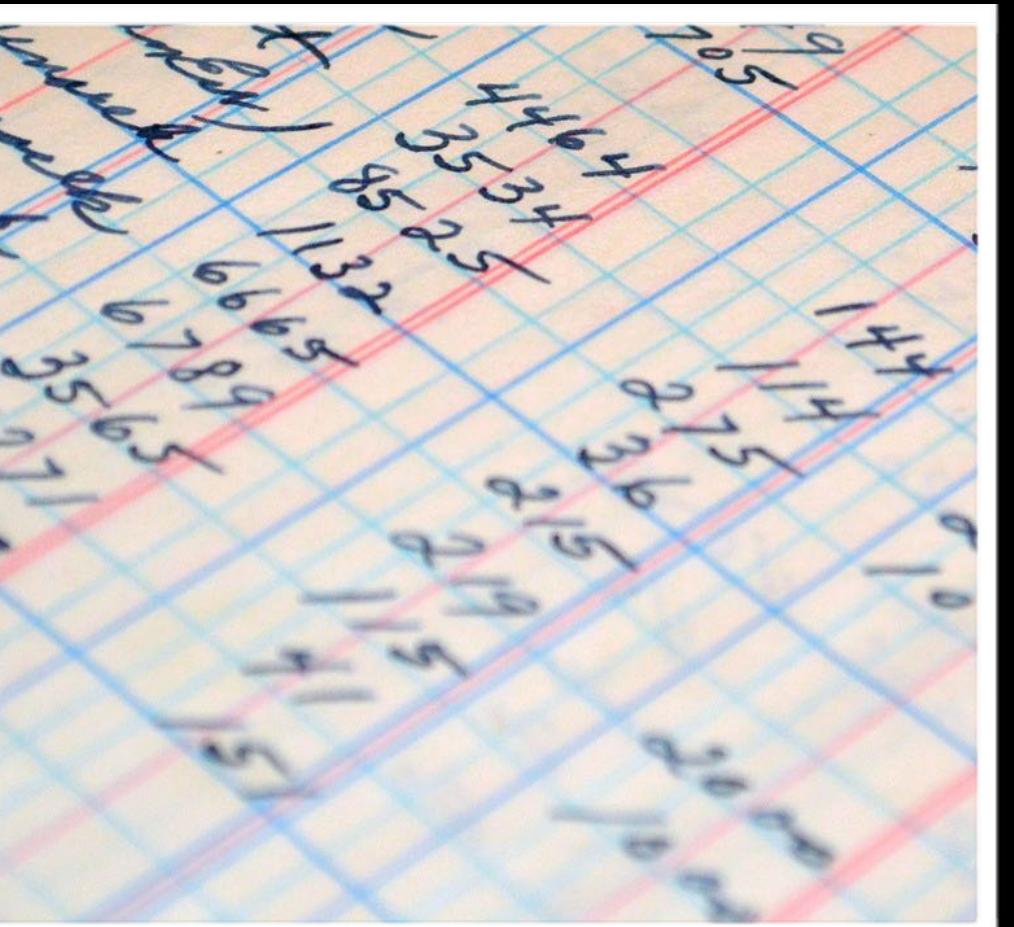
```
exchangeRate": {  
    "base": "USD",  
    "date": "2018-02-13",  
    "rates": { "EUR": 0.806942 }  
}
```



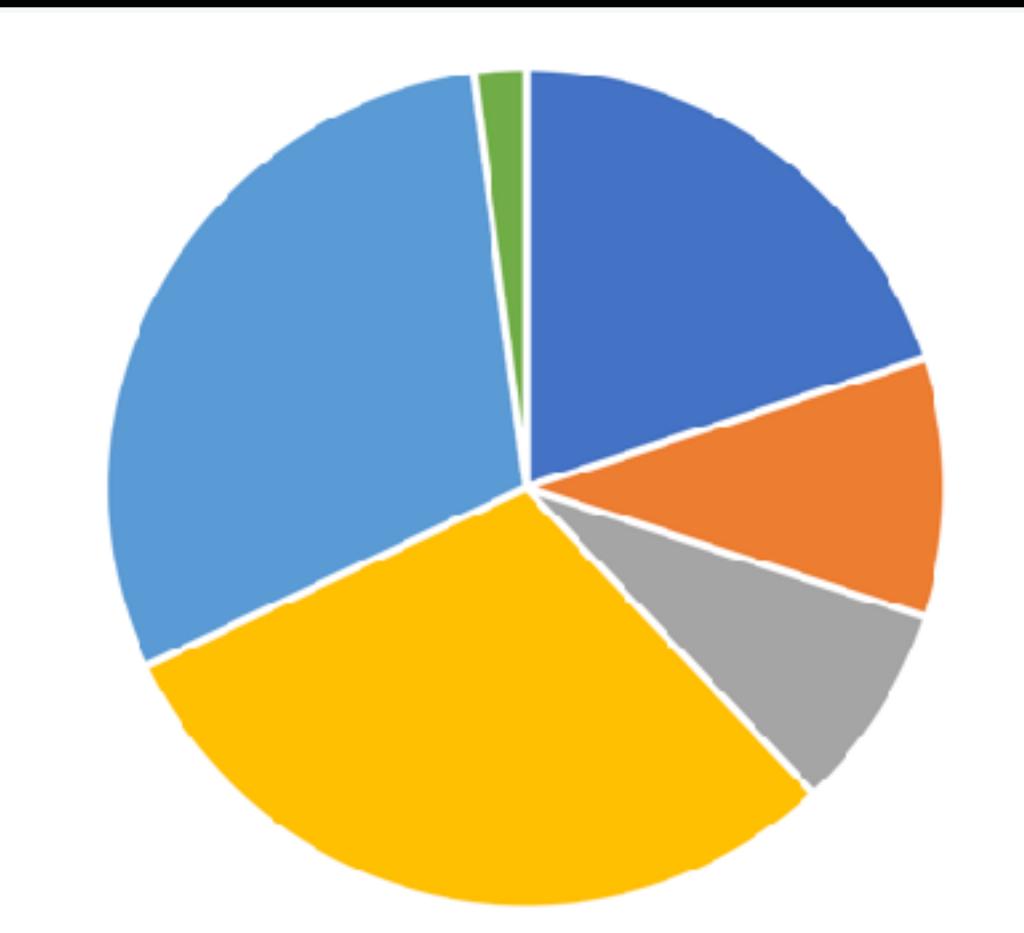
Reduce stream-replay headaches by
storing side-effects as event results

Reusing event data?

Transaction ledger Microservice



Budget Planer Microservice



**Transactionledger
Microservice**

**Budget Planer
Microservice**

TransactionBooked

TransactionCategorised

**“Just” copy data into
different events, “just”
so convenient**



TransactionBooked

transactionId: ...
accountNumber: ...
amount: ...
currency: ...
bookingTime: ...
purpose: ...

TransactionBooked

transactionId: ...
accountNumber: ...
amount: ...
currency: ...
bookingTime: ...
purpose: ...

TransactionCategorised

tagId: ...
categoryName: “...”
transactionId: ...
amount: ...
currency: ...

TransactionBooked

transactionId: ...
accountNumber: ...
amount: ...
currency: ...
bookingTime: ...
purpose: ...

TransactionCategorised

tagId: ...
categoryName: “...”
transactionId: ...
amount: ...
currency: ...

**But I need to
display the
transaction
purpose, too**

**Budget Planer
Microservice**

TransactionBooked

transactionId: ...
accountNumber: ...
amount: ...
currency: ...
bookingTime: ...
purpose: ...

TransactionCategorised

tagId: ...
categoryName: “...”
transactionId: ...
amount: ...
currency: ...
????

The lossy event

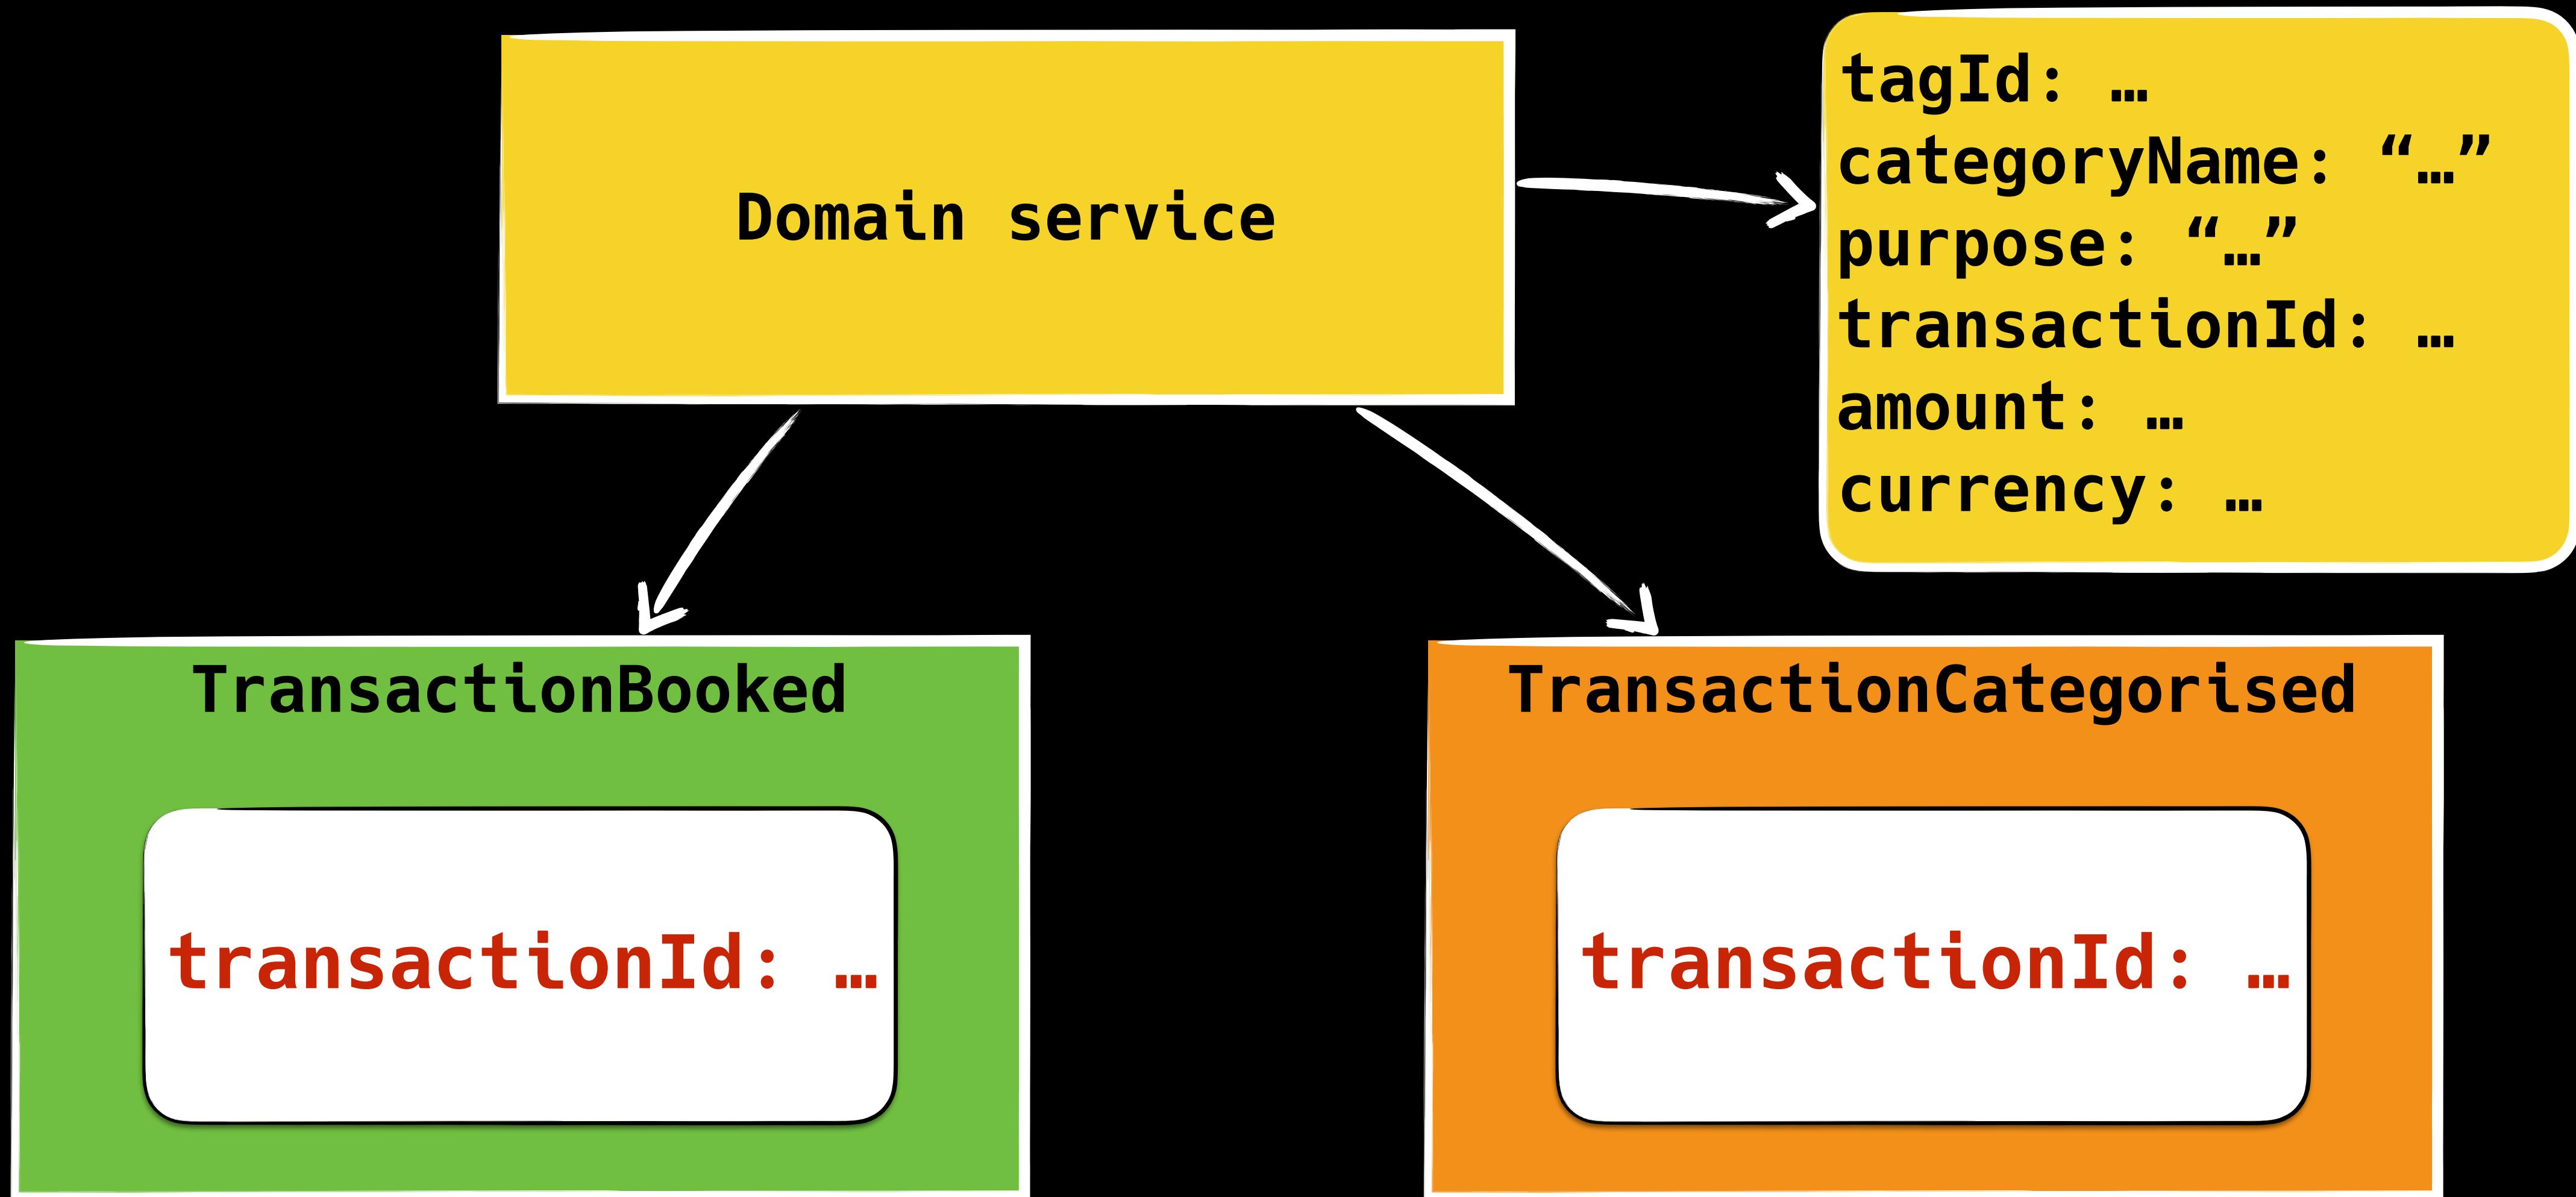
Only reference aggregates via their **root id**

TransactionBooked

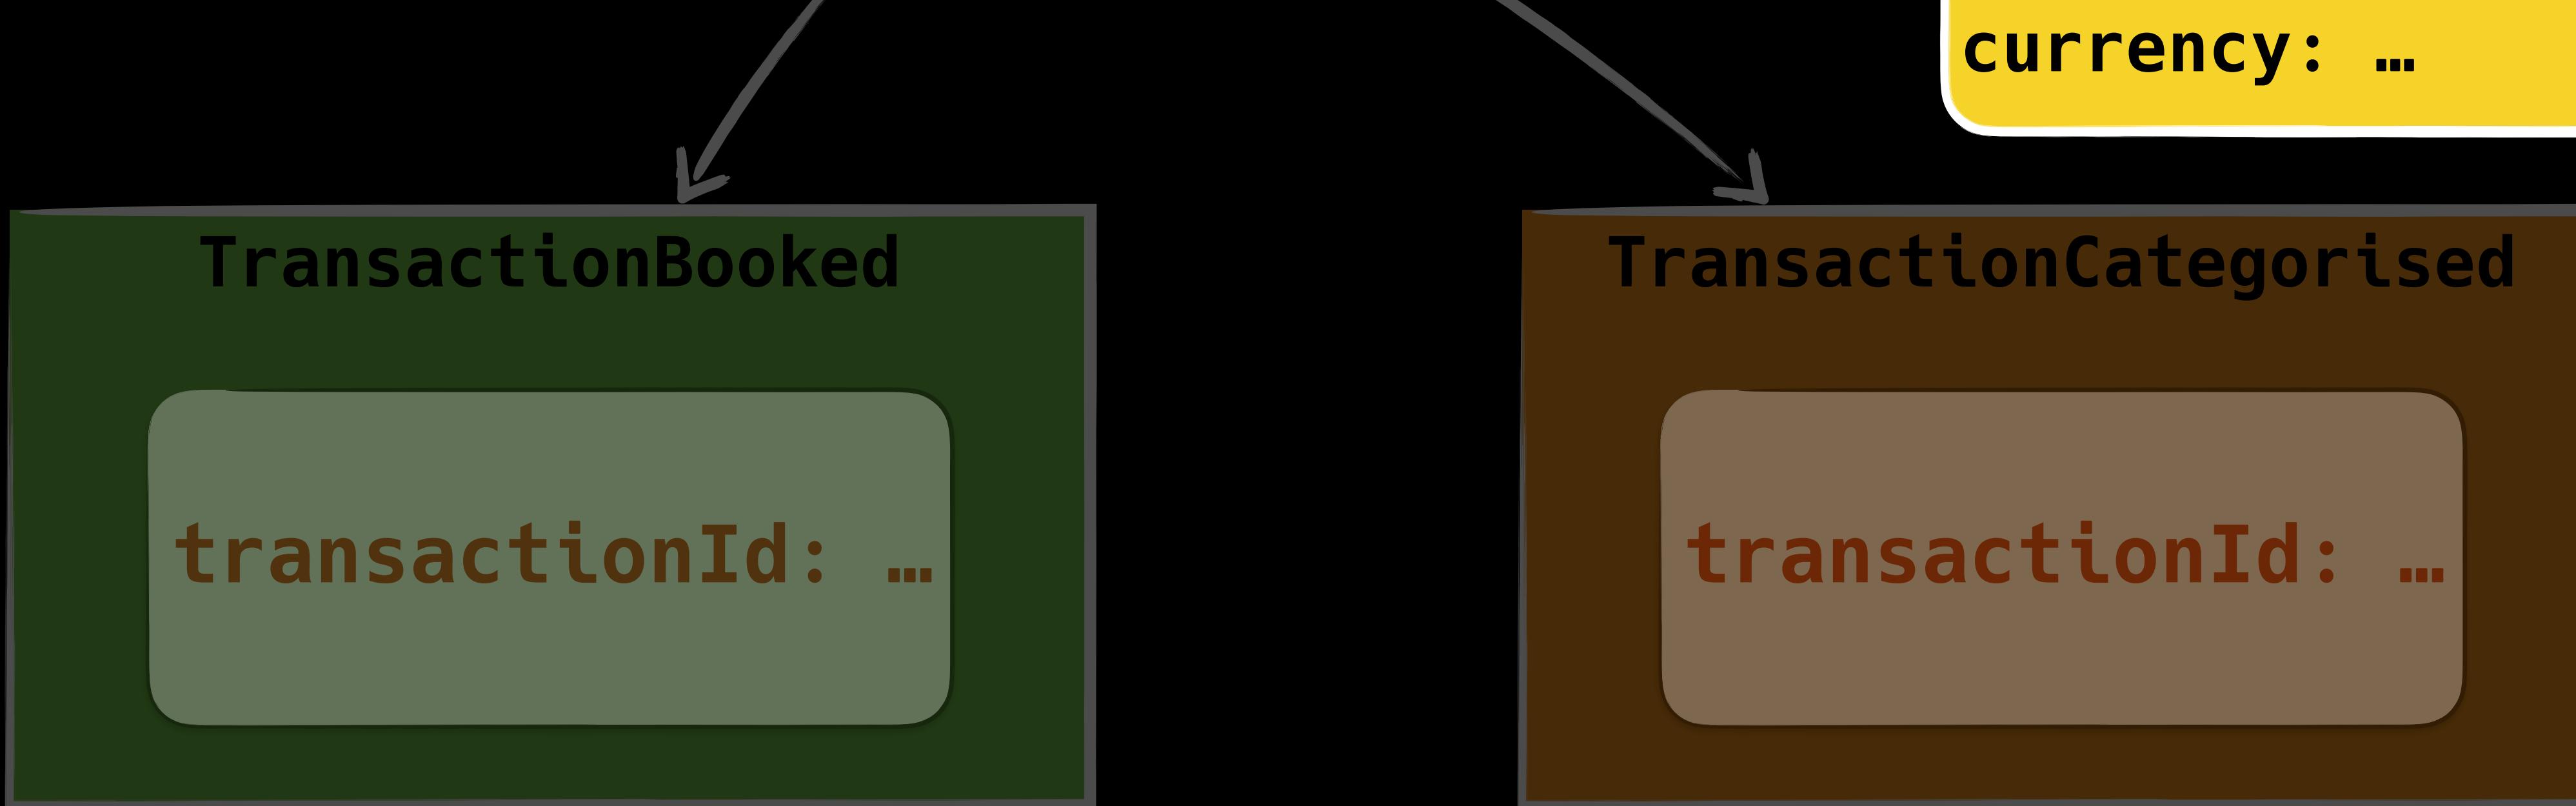
transactionId: ...
accountNumber: ...
amount: ...
currency: ...
bookingTime: ...
purpose: ...

TransactionCategorised

tagId: ...
categoryName: “...”
transactionId: ...



**Good candidate for
a read model resp.
projection btw.**



Don't copy parts of an event.
Prefer building use case
specific projections

How can you handle event data
over a long period of time?

You don't

**“Just” take a
snapshot of the
stream**



Year's end procedure

Year end – also known as an accounting reference date – is the completion of an accounting period. At this time, businesses need to carry out specific procedures to close their books.

<https://debitoor.com/dictionary/year-end>

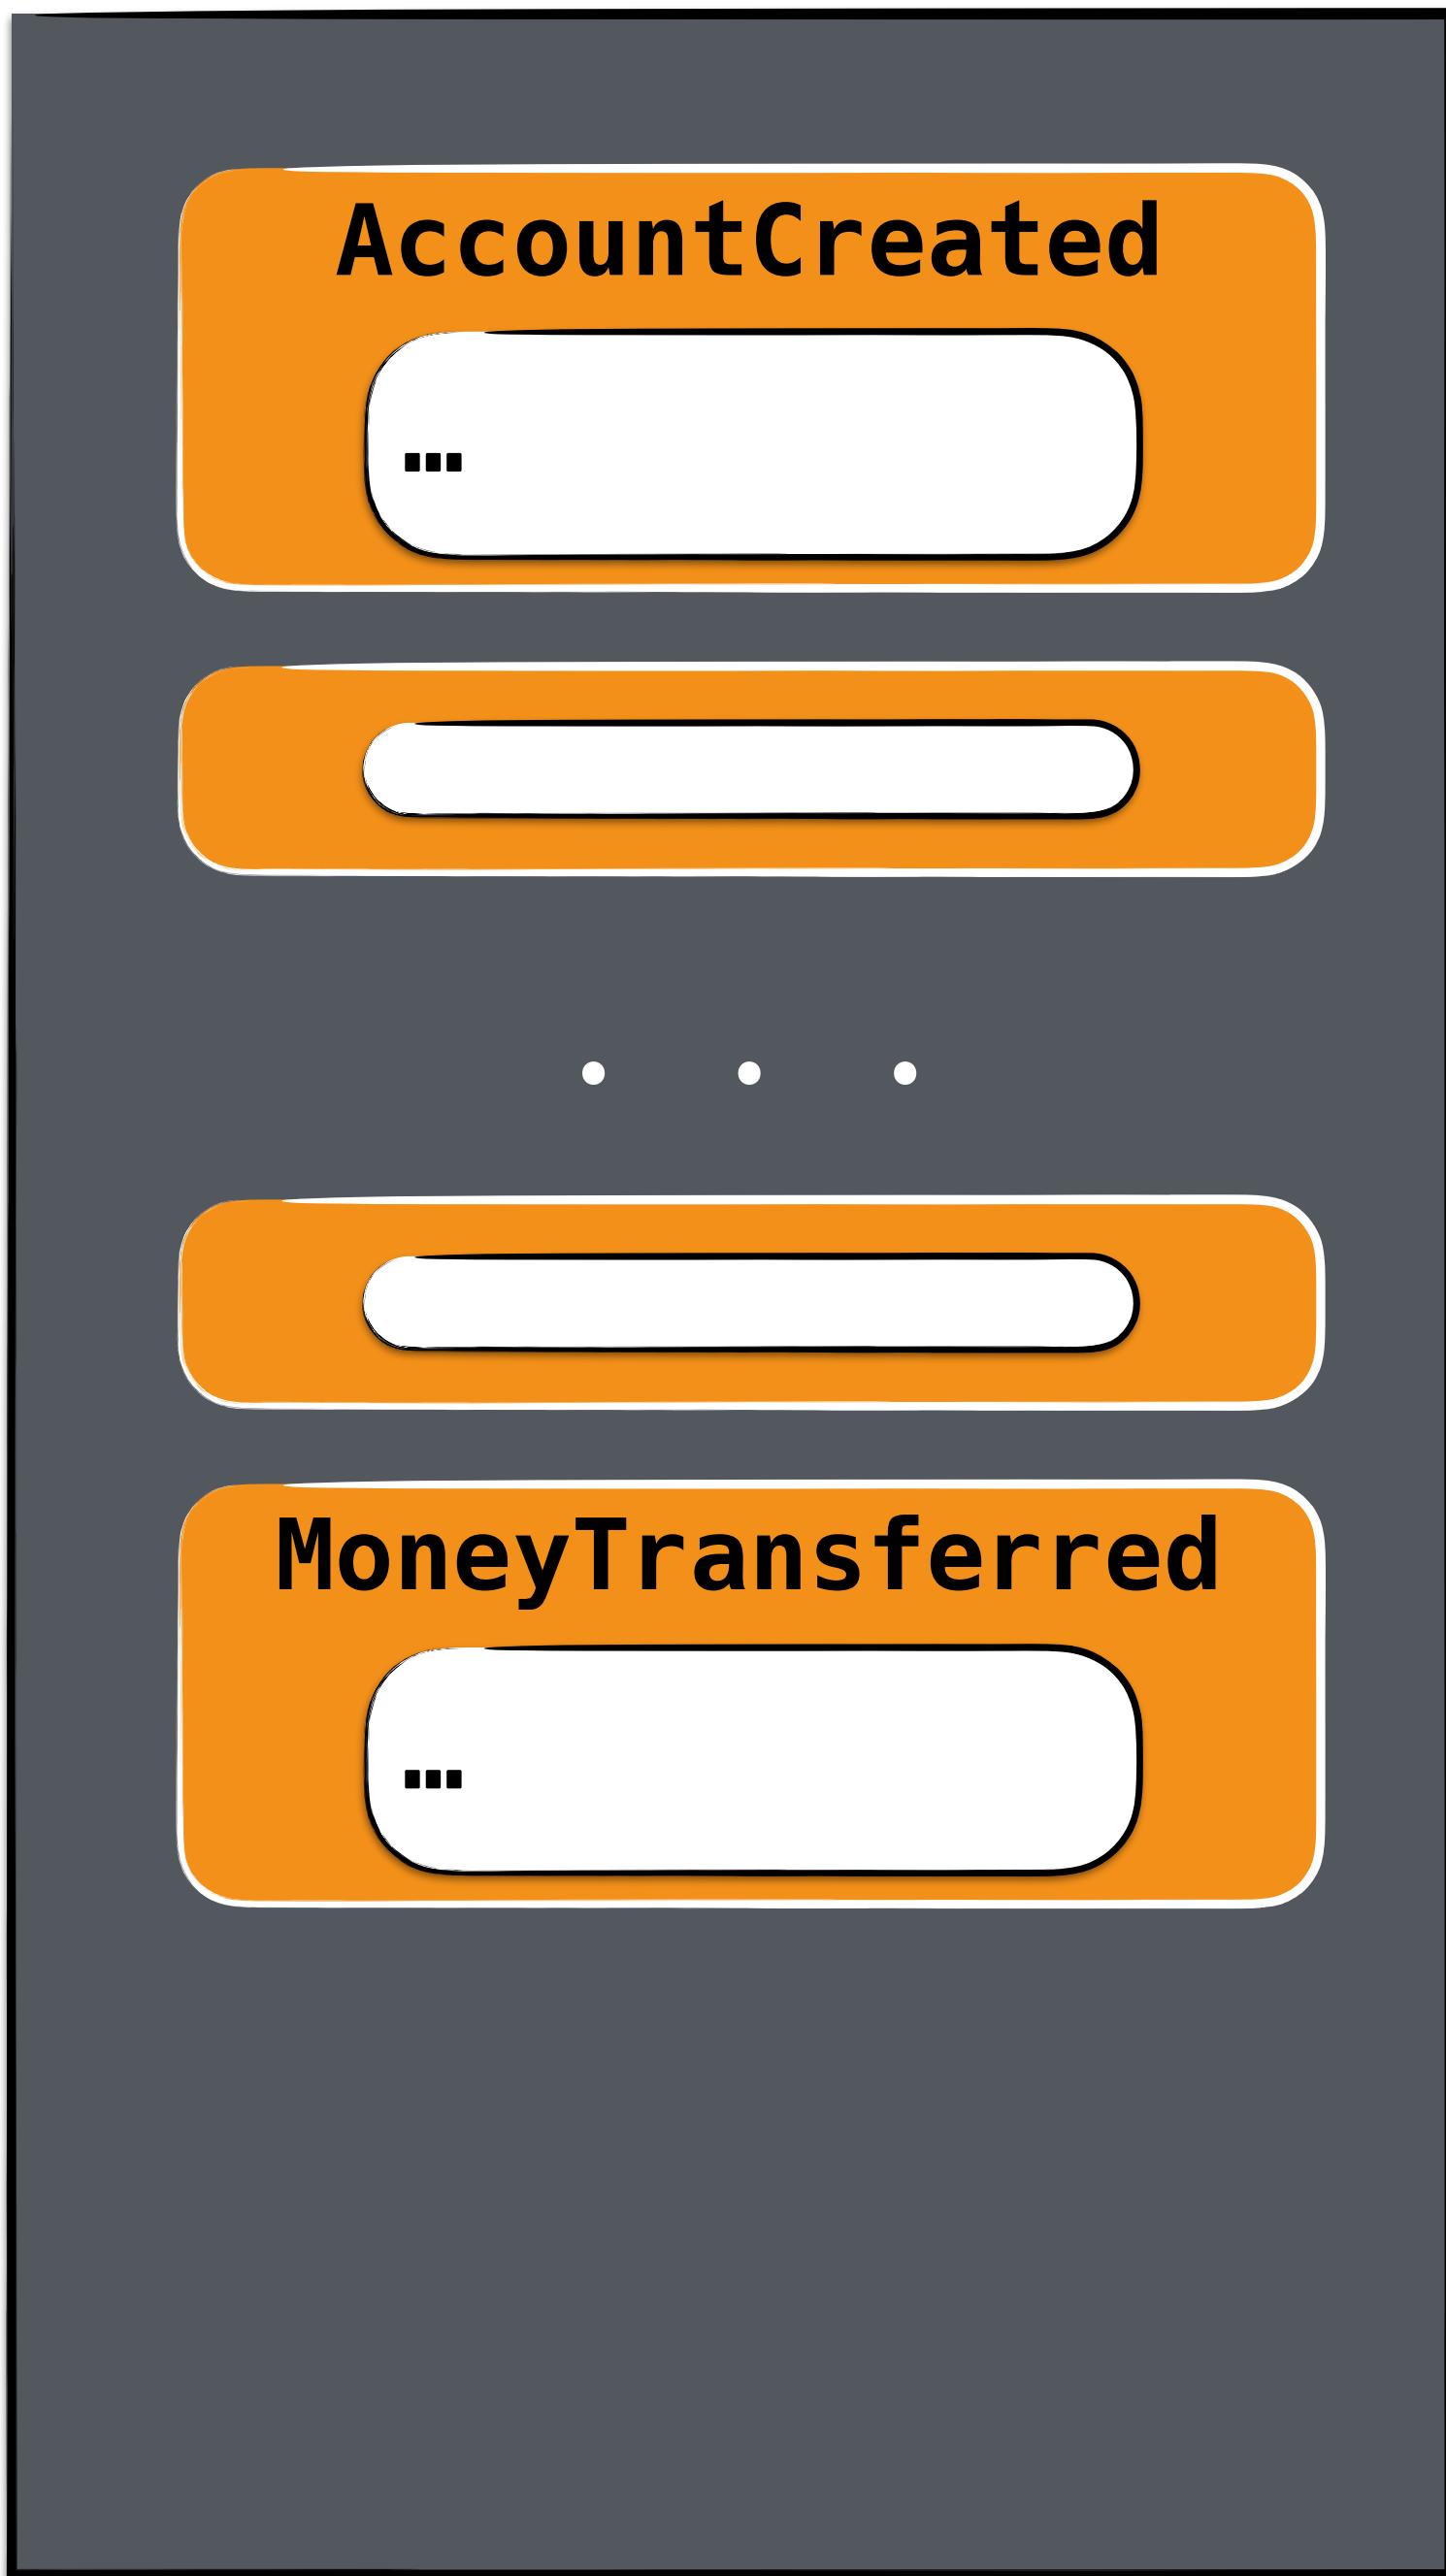
Year end – also known as an accounting reference date – is the completion of an accounting period. At this time, businesses need to carry out specific procedures to close their books.

<https://debitoor.com/dictionary/year-end>

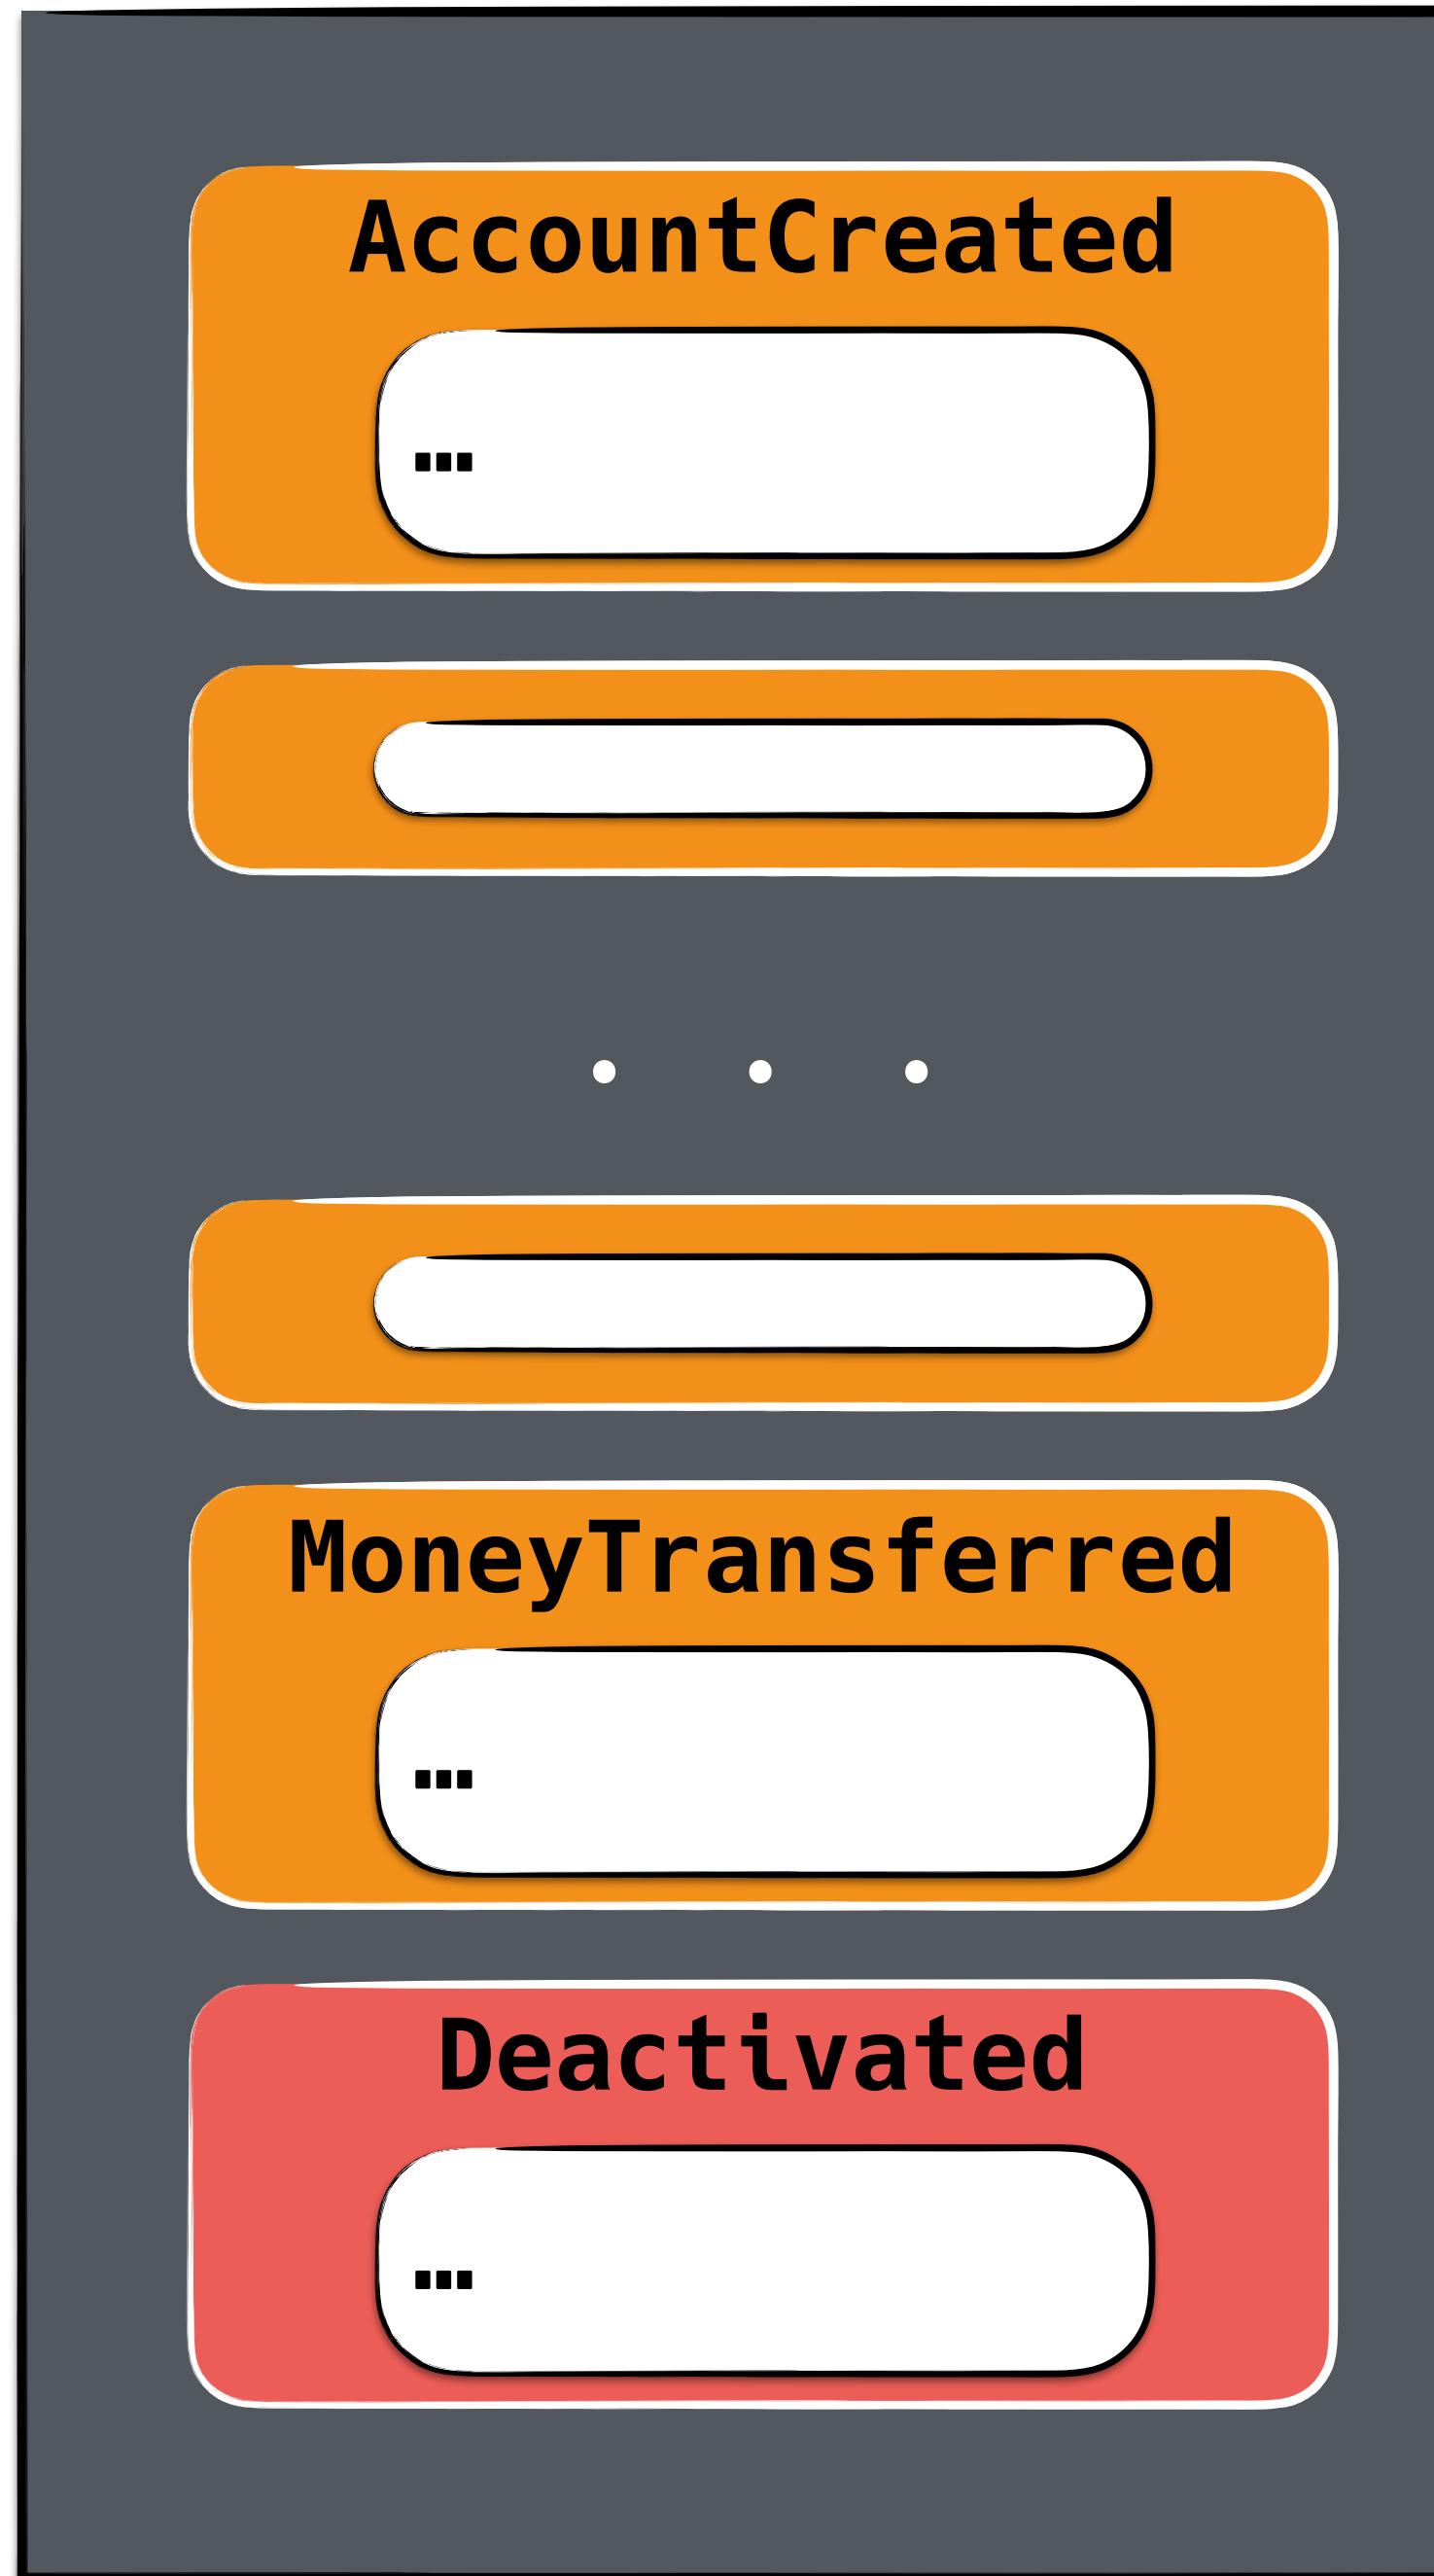
Copy-Transform

a.k.a. eventsourcing **refactoring** powertool

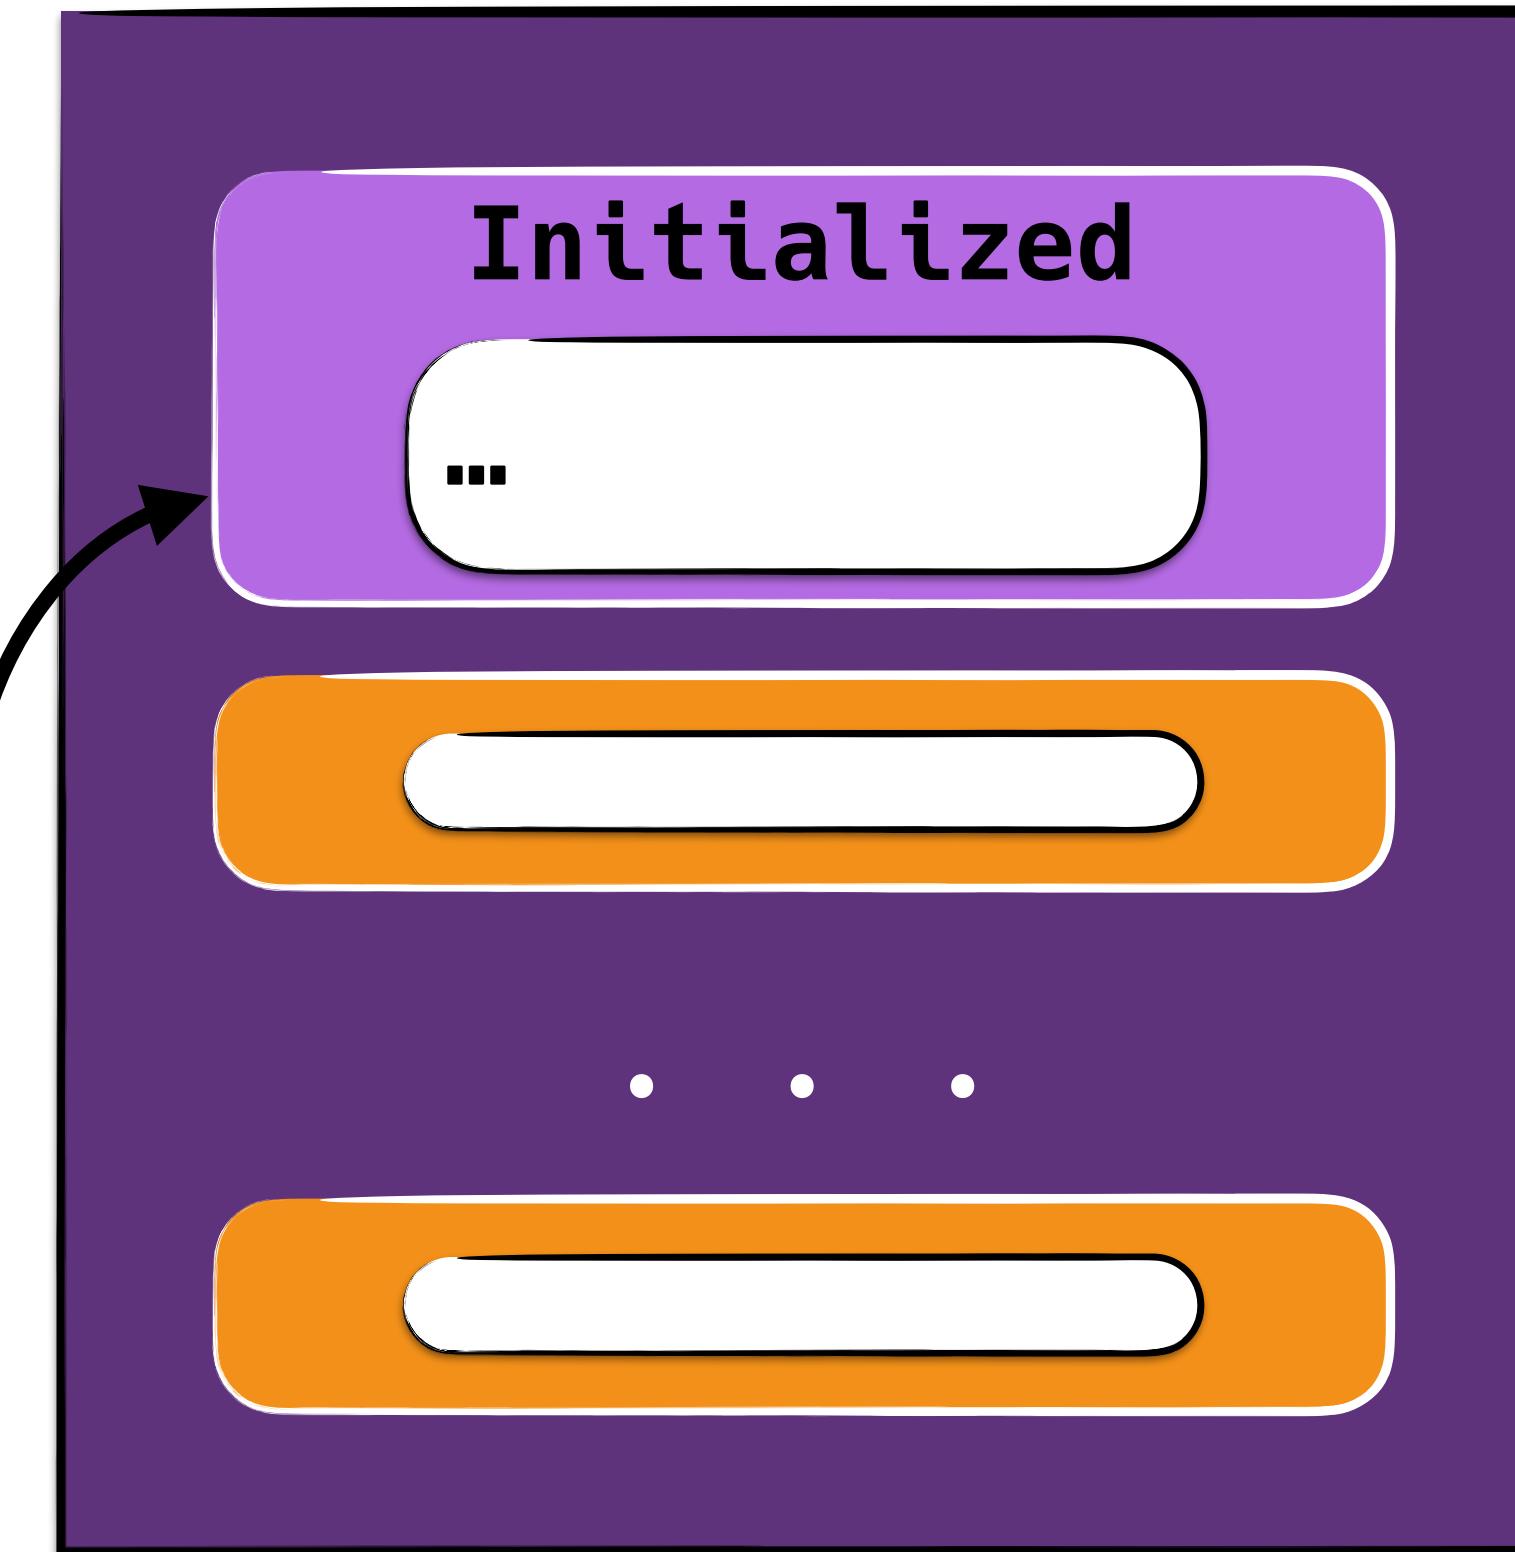
2017

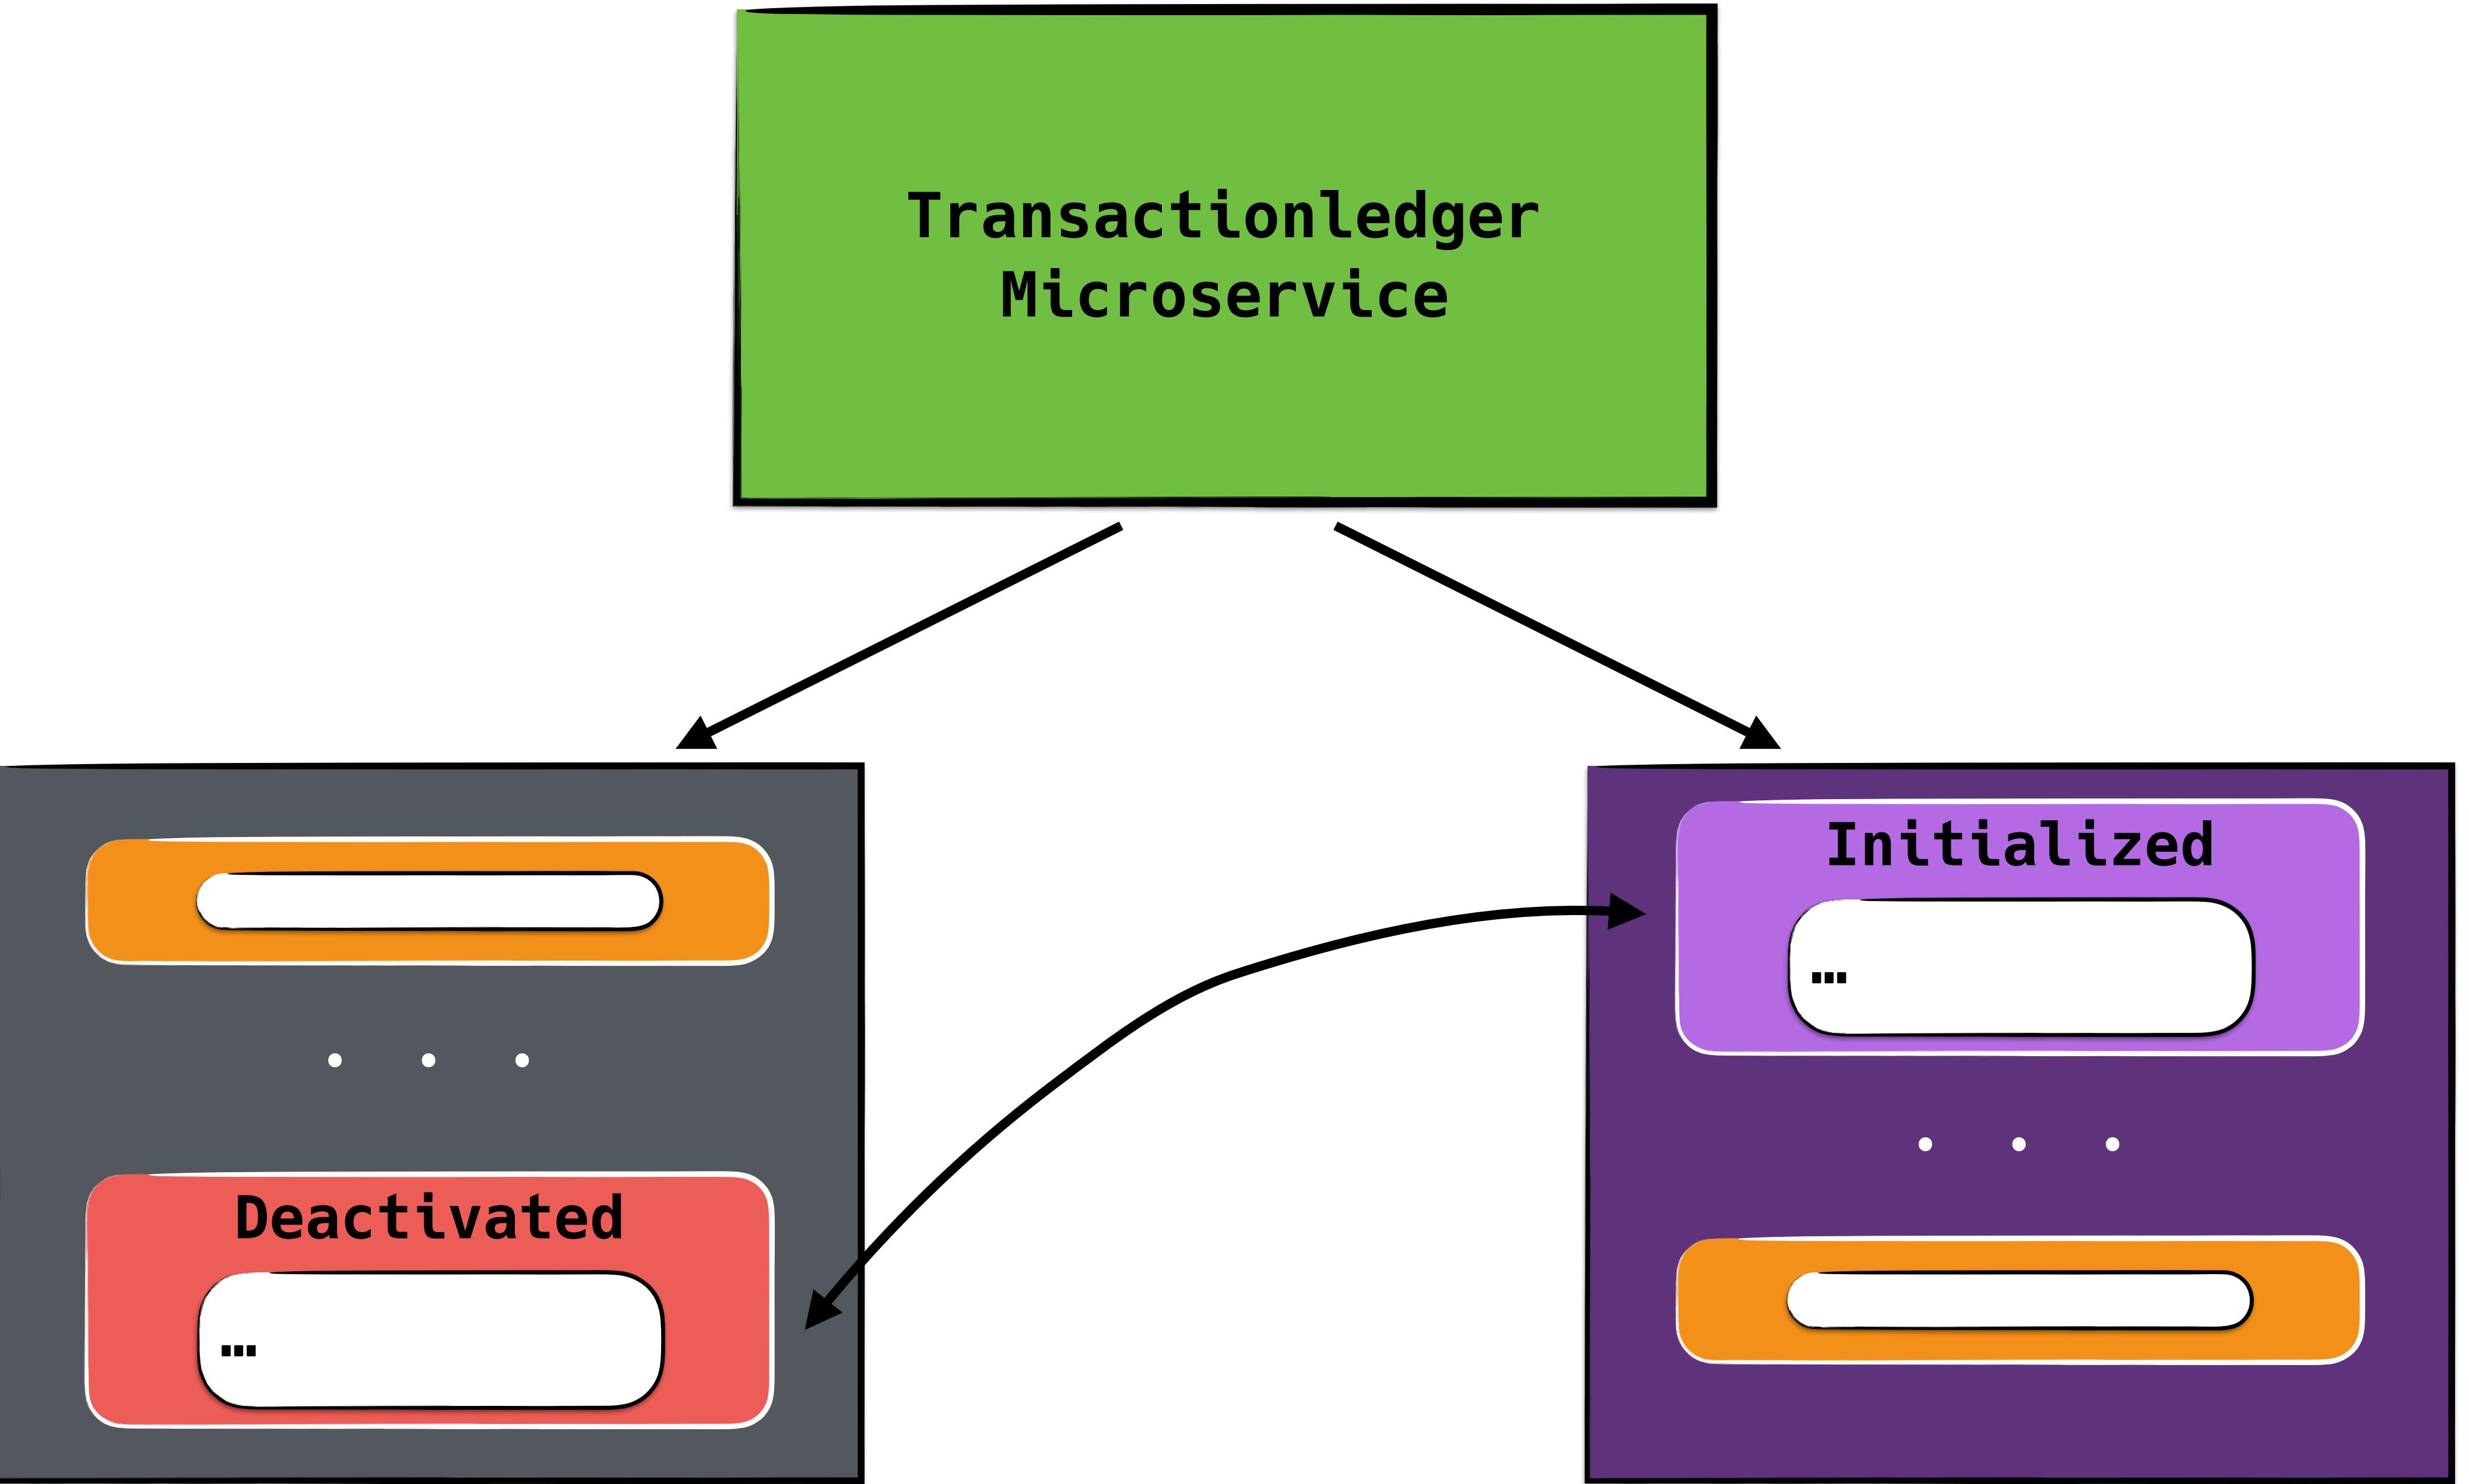


2017



2018





**Same idea if you need to
remodel your domain!**

The devil is in the detail

Dealing with errors

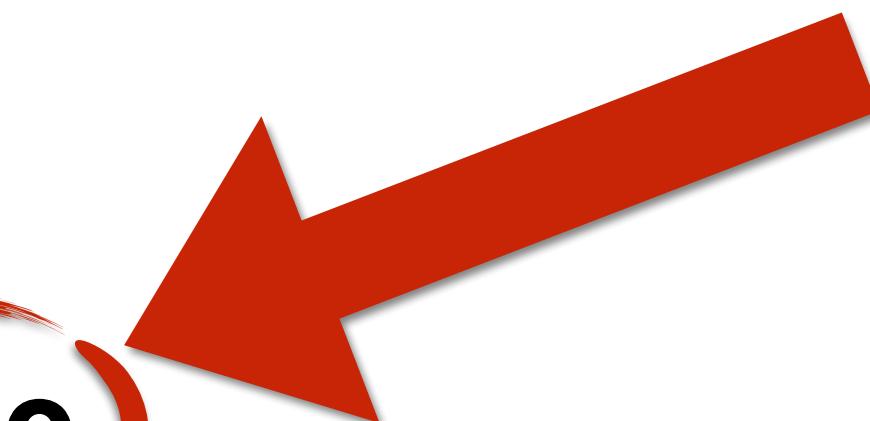


MoneyTransferred

eventId: 231233

amount: 97

Euro



withdrawnAt: 2018-08-30T08:58:26.624

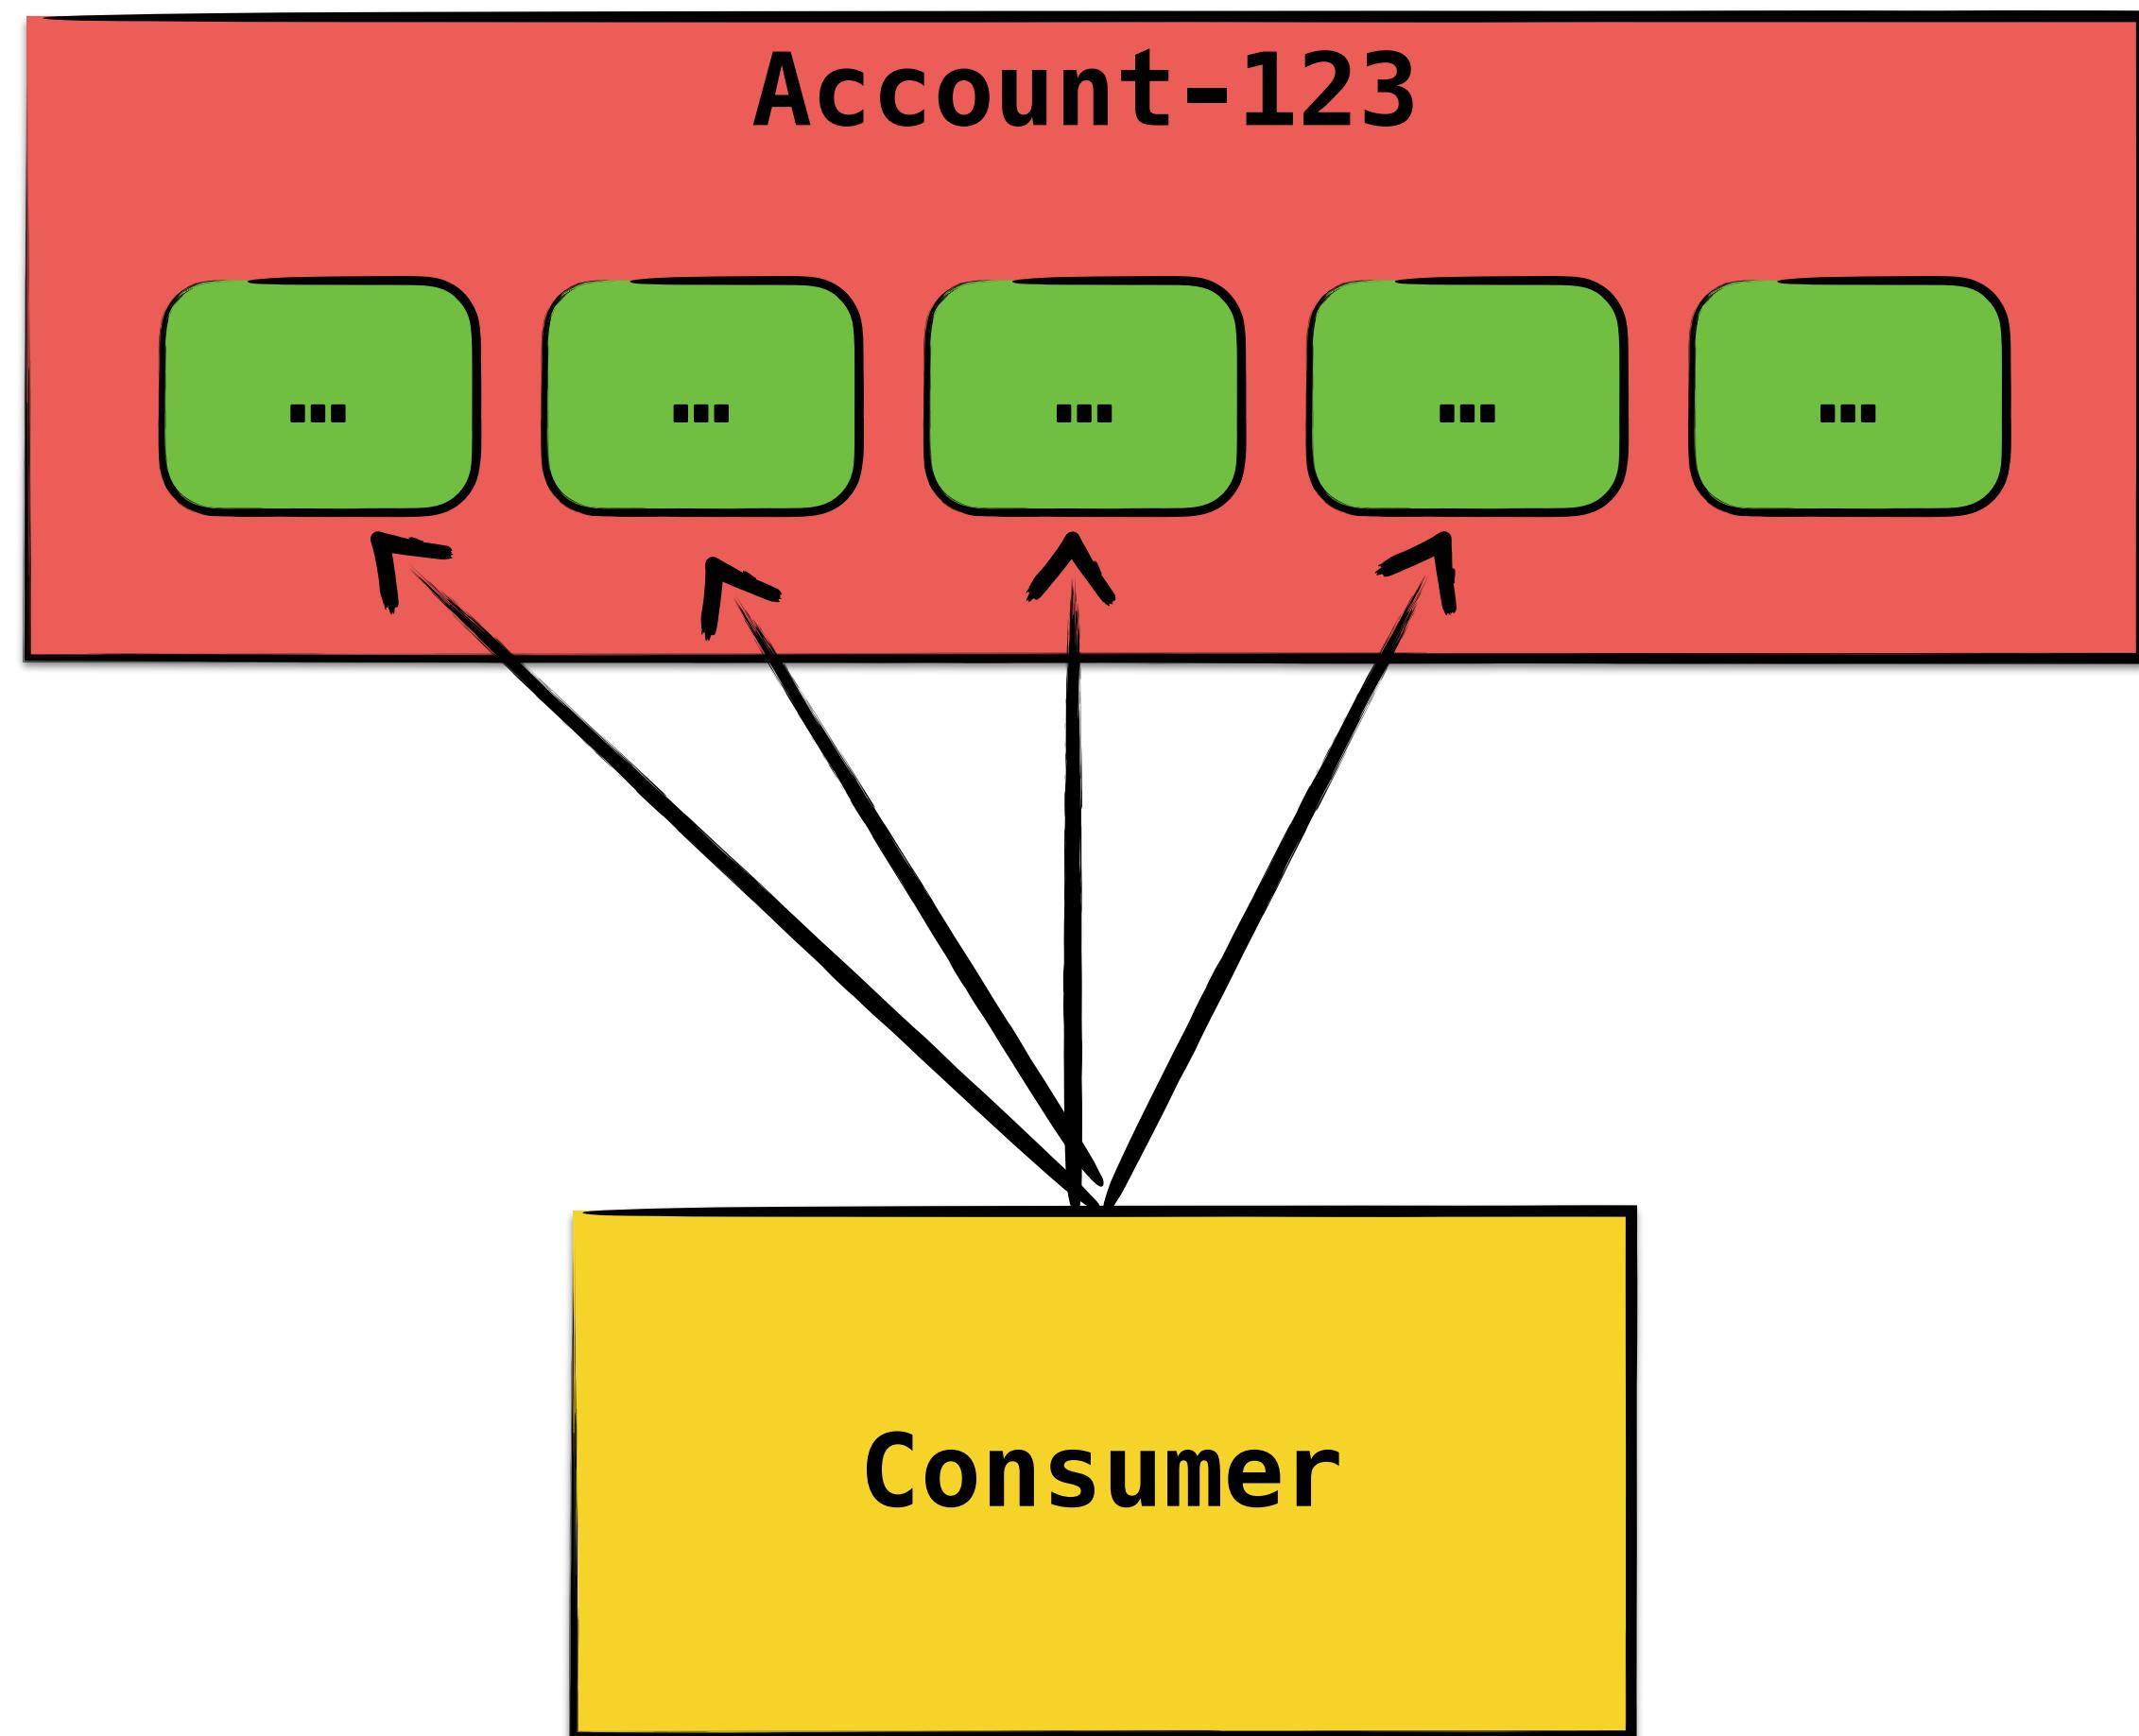
**“Just” update the
event in the
eventstore!**



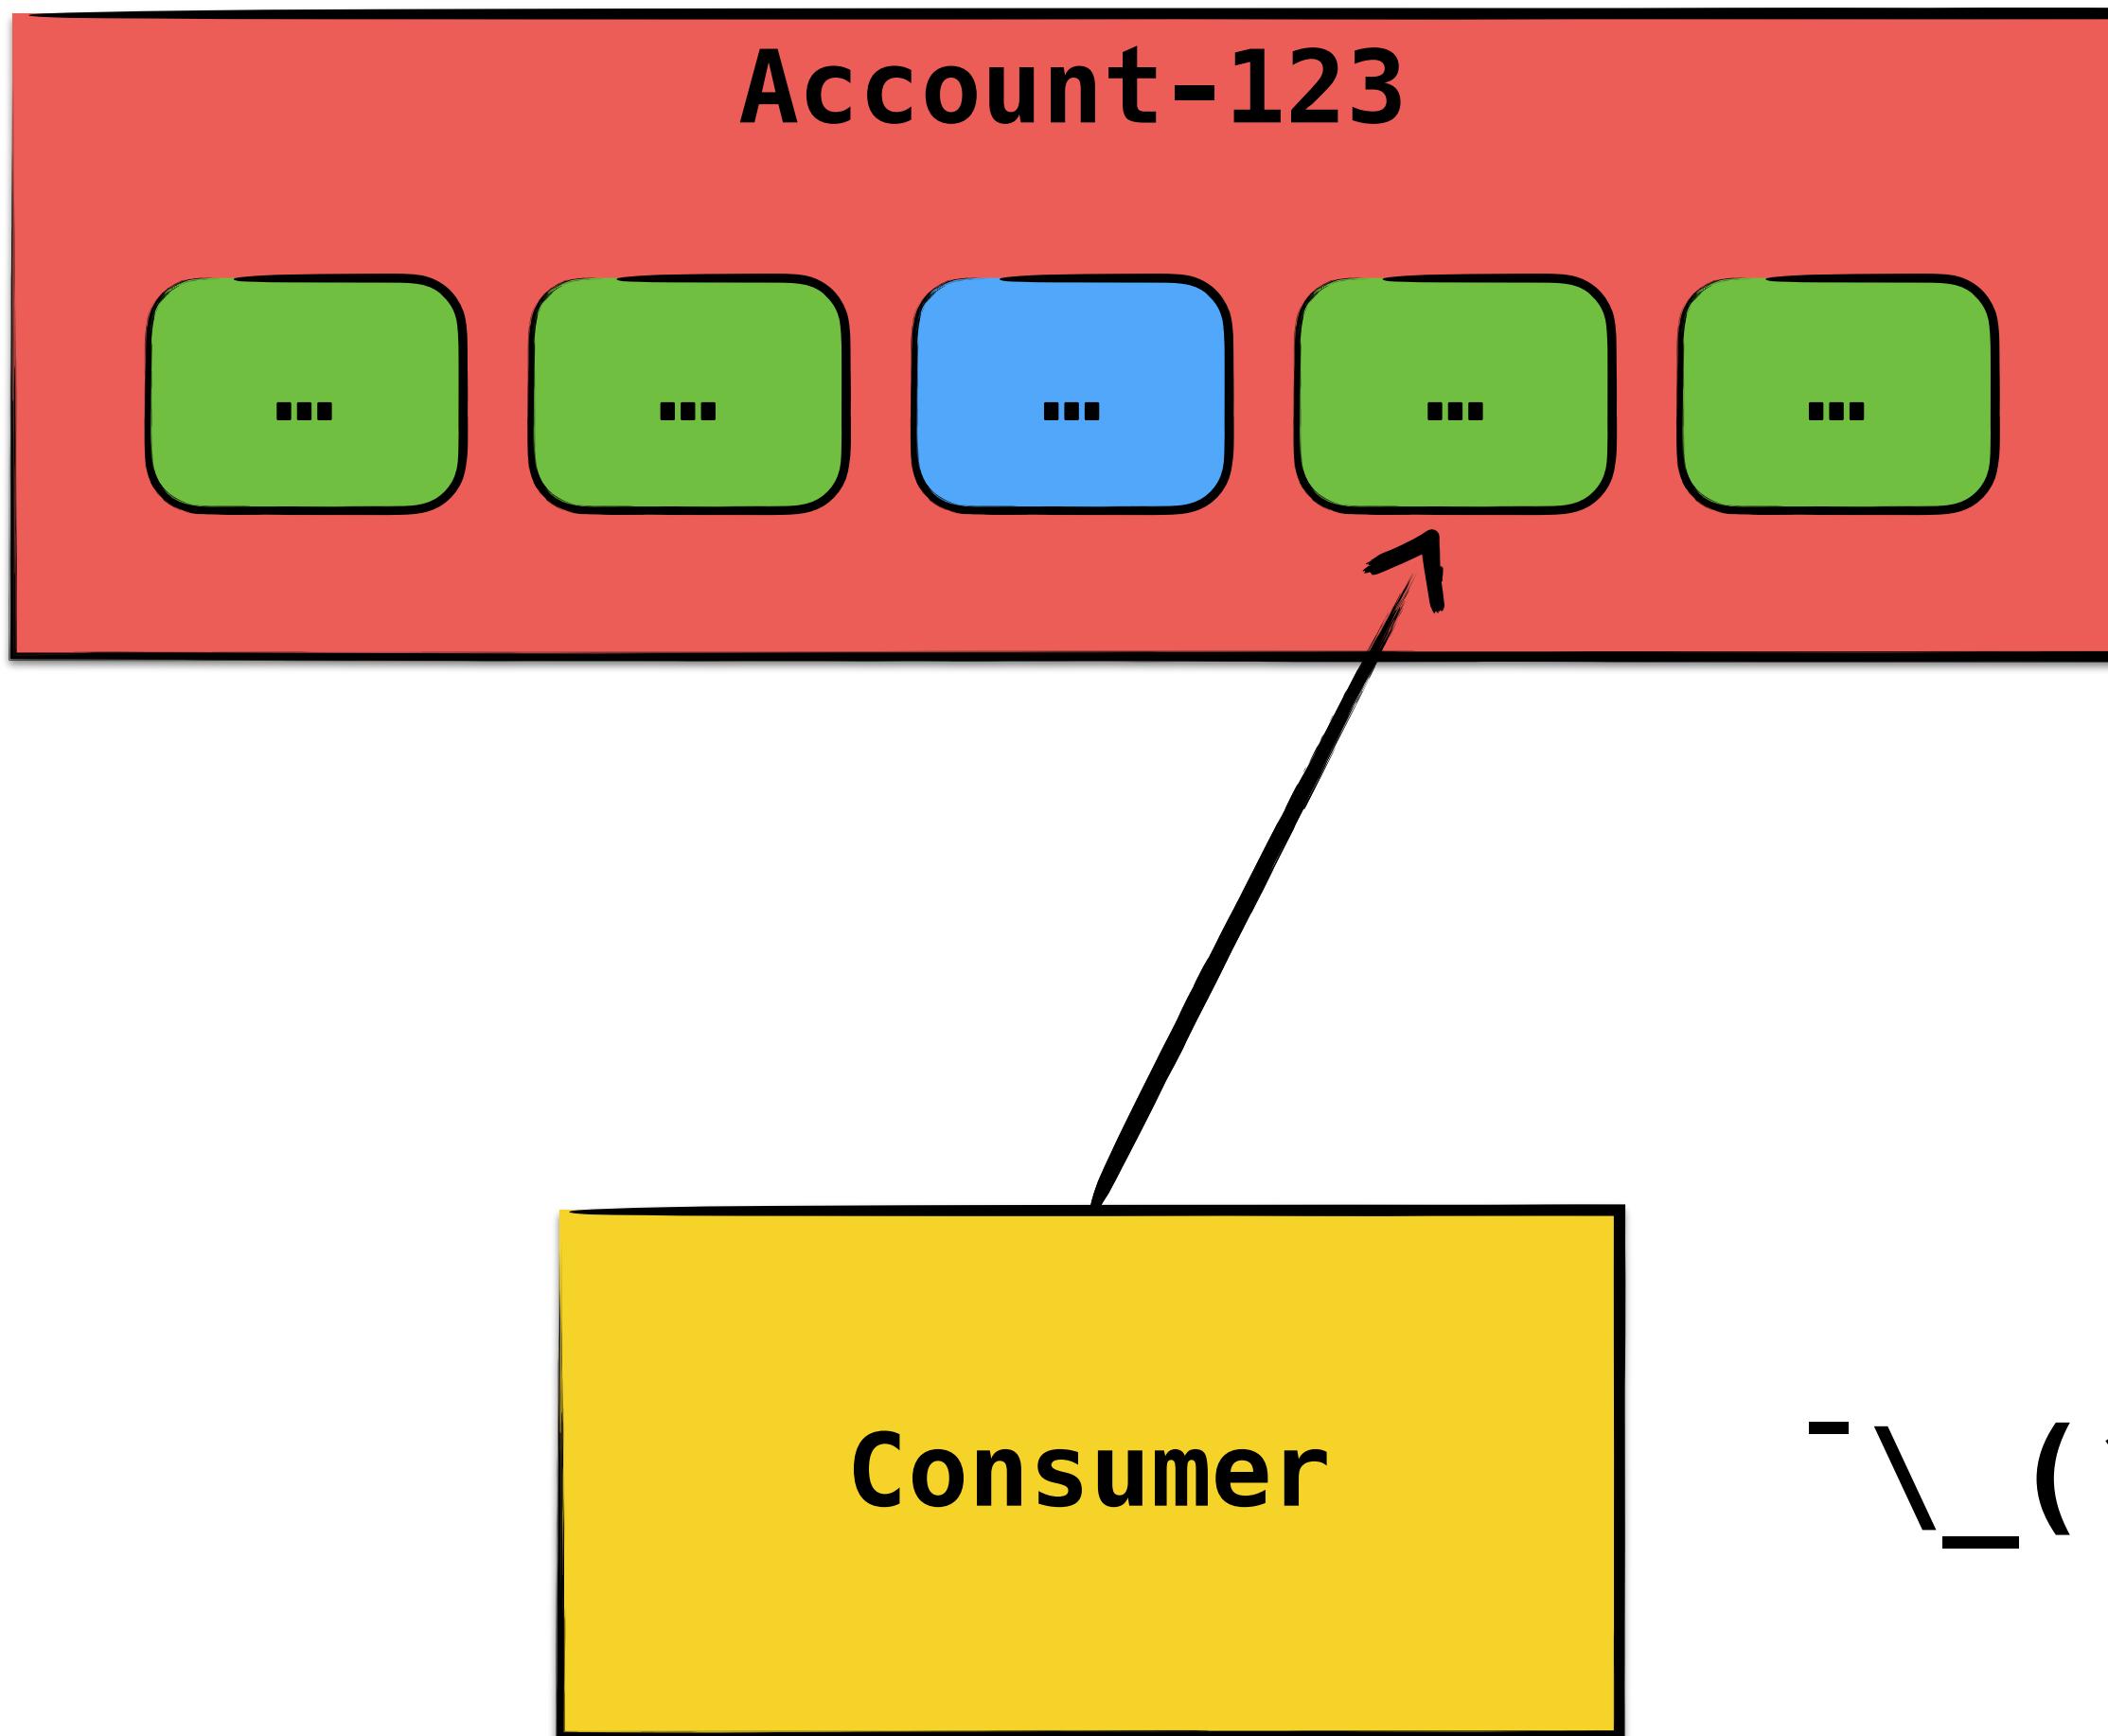


```
UPDATE transactions  
SET currency='EUR'  
WHERE eventId='231233'
```

No!



Update ...



I already
know that
event.

Why
should I
re-read?

-_(ツ)_/-

Ok. Then “**just**” use
compensation events



The cancelled or corrected event

Partial compensation?

MoneyTransferred
eventId: 1

amount: 97 Euro

...

MoneyTransferAmountCorrected
eventId: 2

amount: 97 EUR
eventId: 1

Full compensation

- do as accountants do

MoneyTransferred

eventId: 1

amount: 97 Euro

...

MoneyTransferCancelled

eventId: 2

reasonEventId: 1

reason: ...

MoneyTransferred

eventId: 3

amount: 97 EUR

...

The full compensation makes the
reason for compensation explicit

Consumers must mostly be forward and
backward compatible

Beware lossy events

Prefer projections to event data copying

Refer across aggregates using root ids

GDPR, compliance and eventsourcing





**REGULATION (EU) 2016/679 OF THE EUROPEAN PARLIAMENT AND OF THE COUNCIL
of 27 April 2016**

on the protection of natural persons with regard to the processing of personal data and on the free movement of such data, and repealing Directive 95/46/EC (General Data Protection Regulation)





Article 17

Right to erasure ('right to be forgotten')

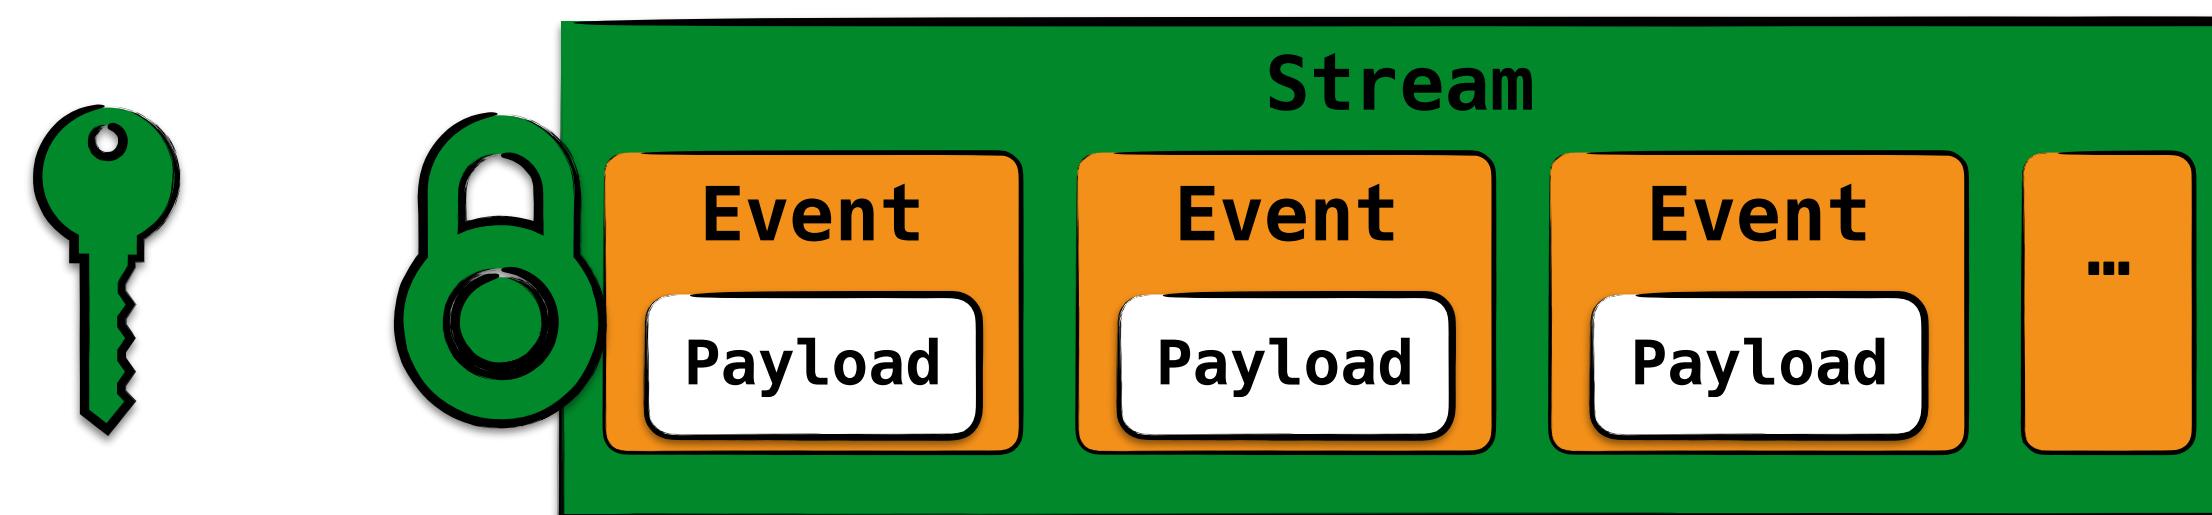
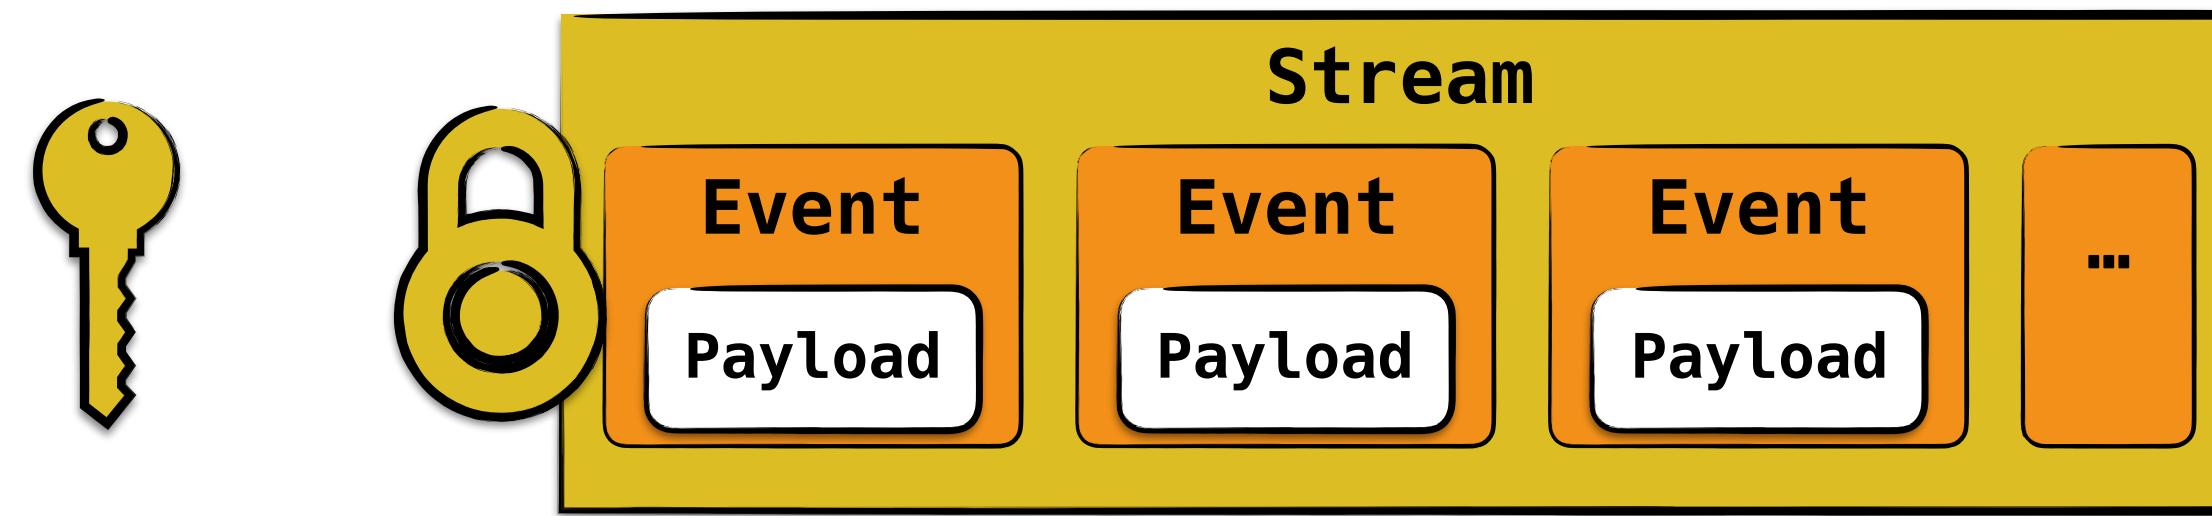
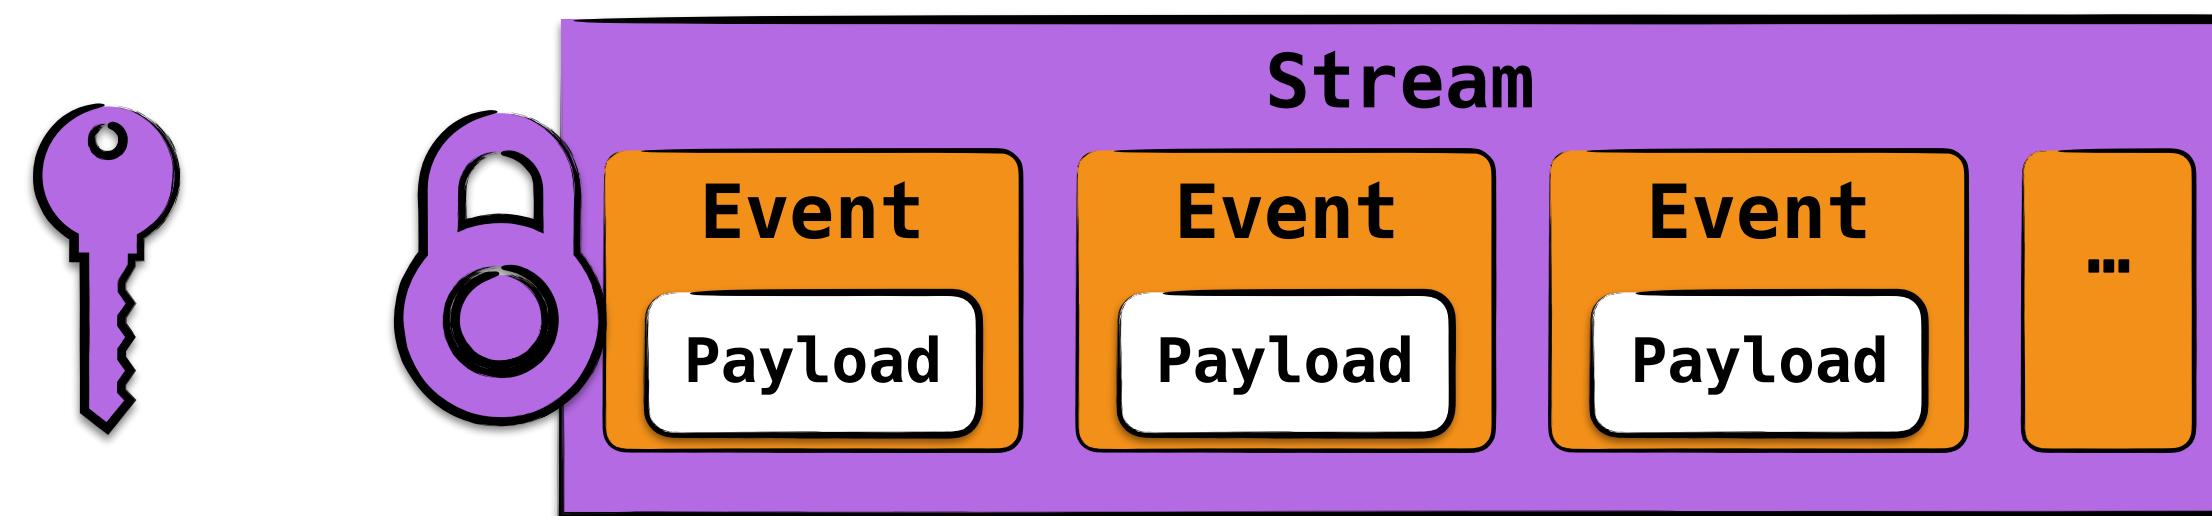
1. The data subject shall have the right to obtain from the controller the erasure of personal data concerning him or her without undue delay and the controller shall have the obligation to erase personal data without undue delay where one of the following grounds applies:



“Just” encrypt and
throw the key away

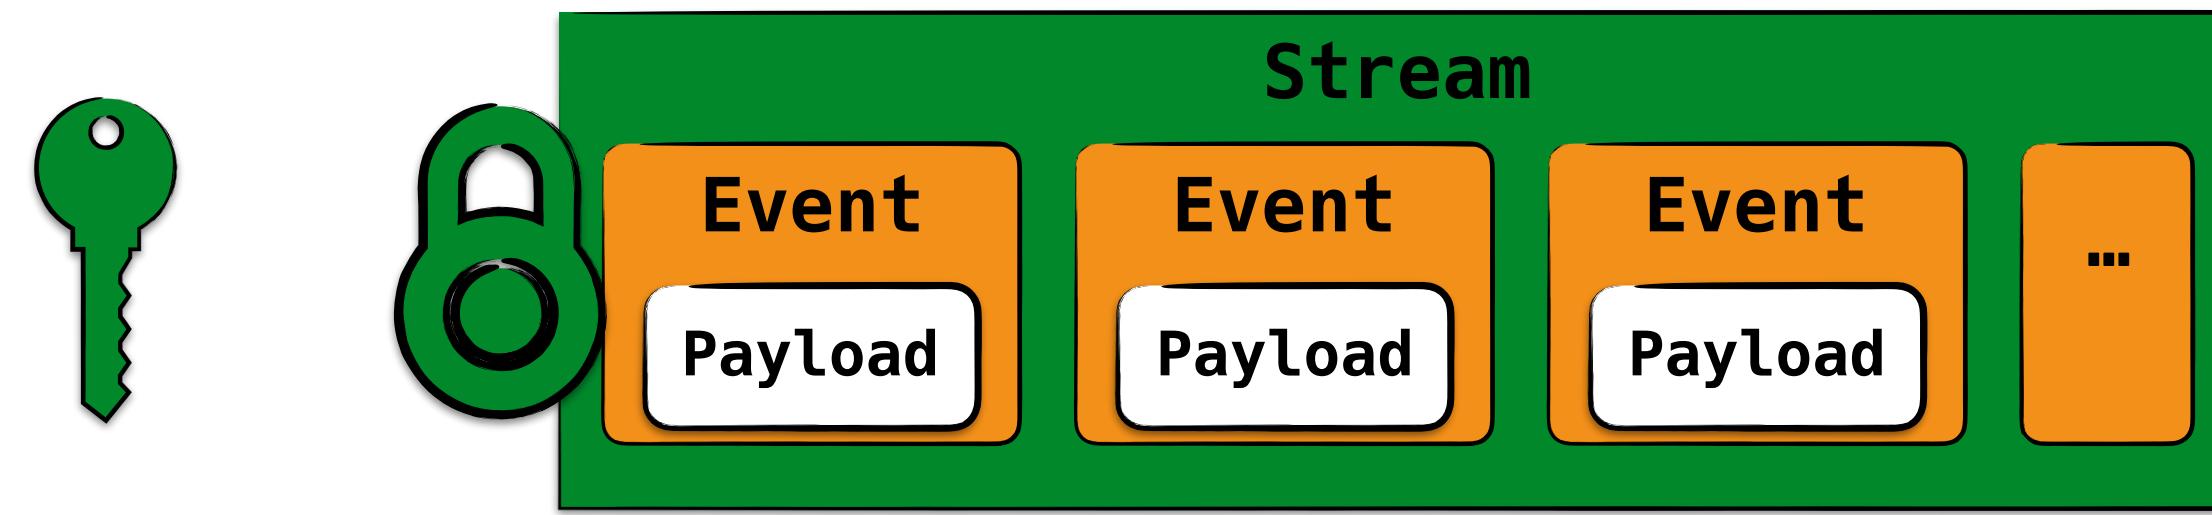
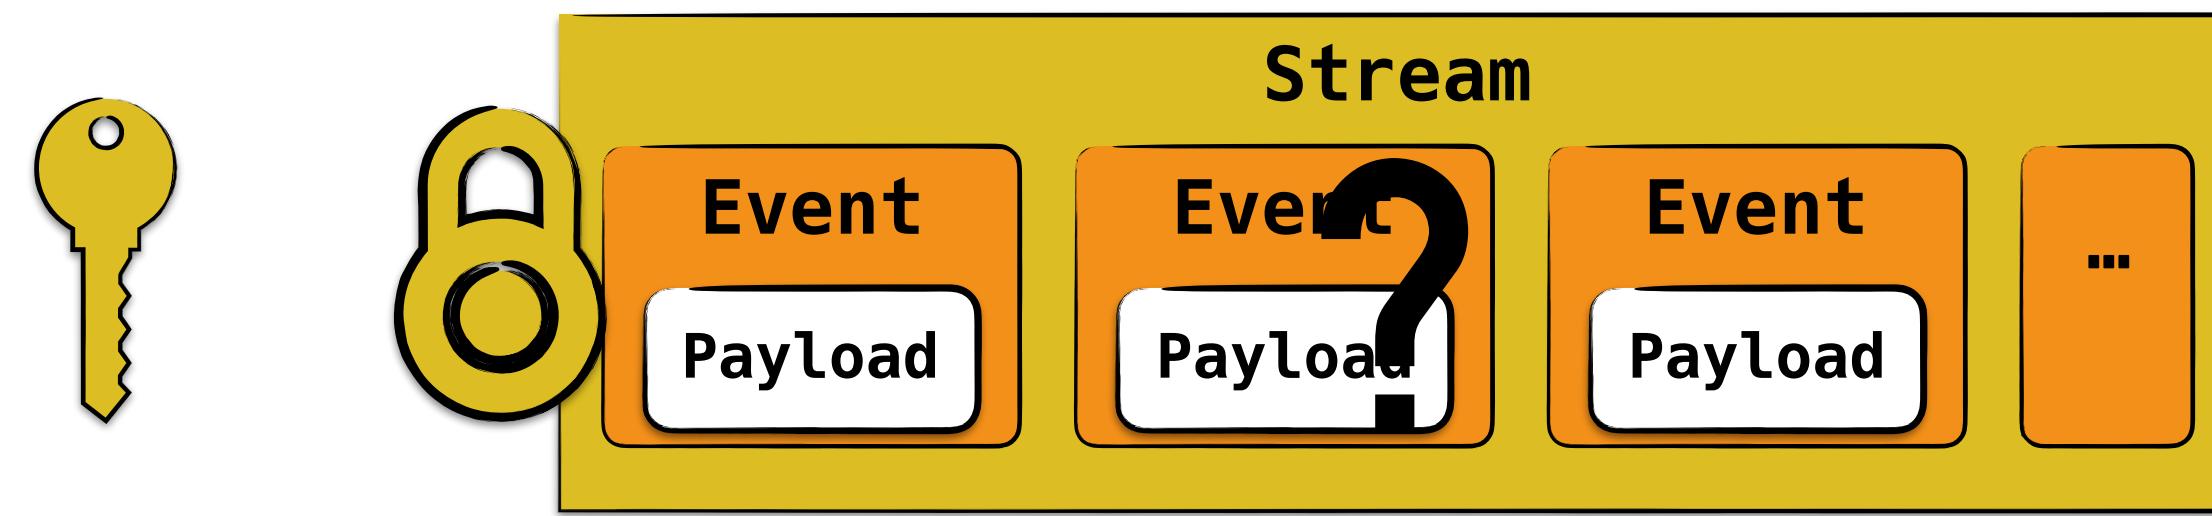
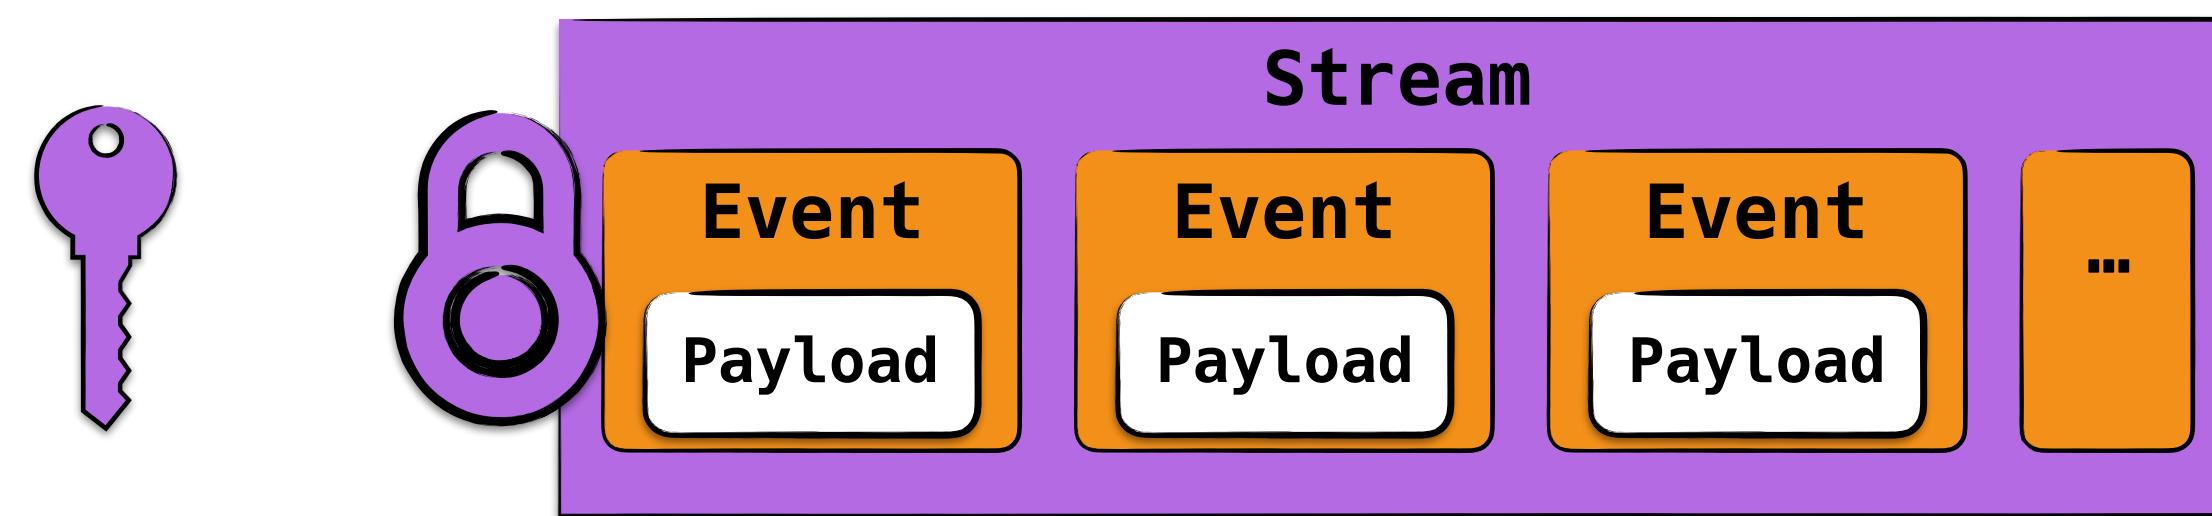


**Every stream is encrypted
using a stream-specific key**



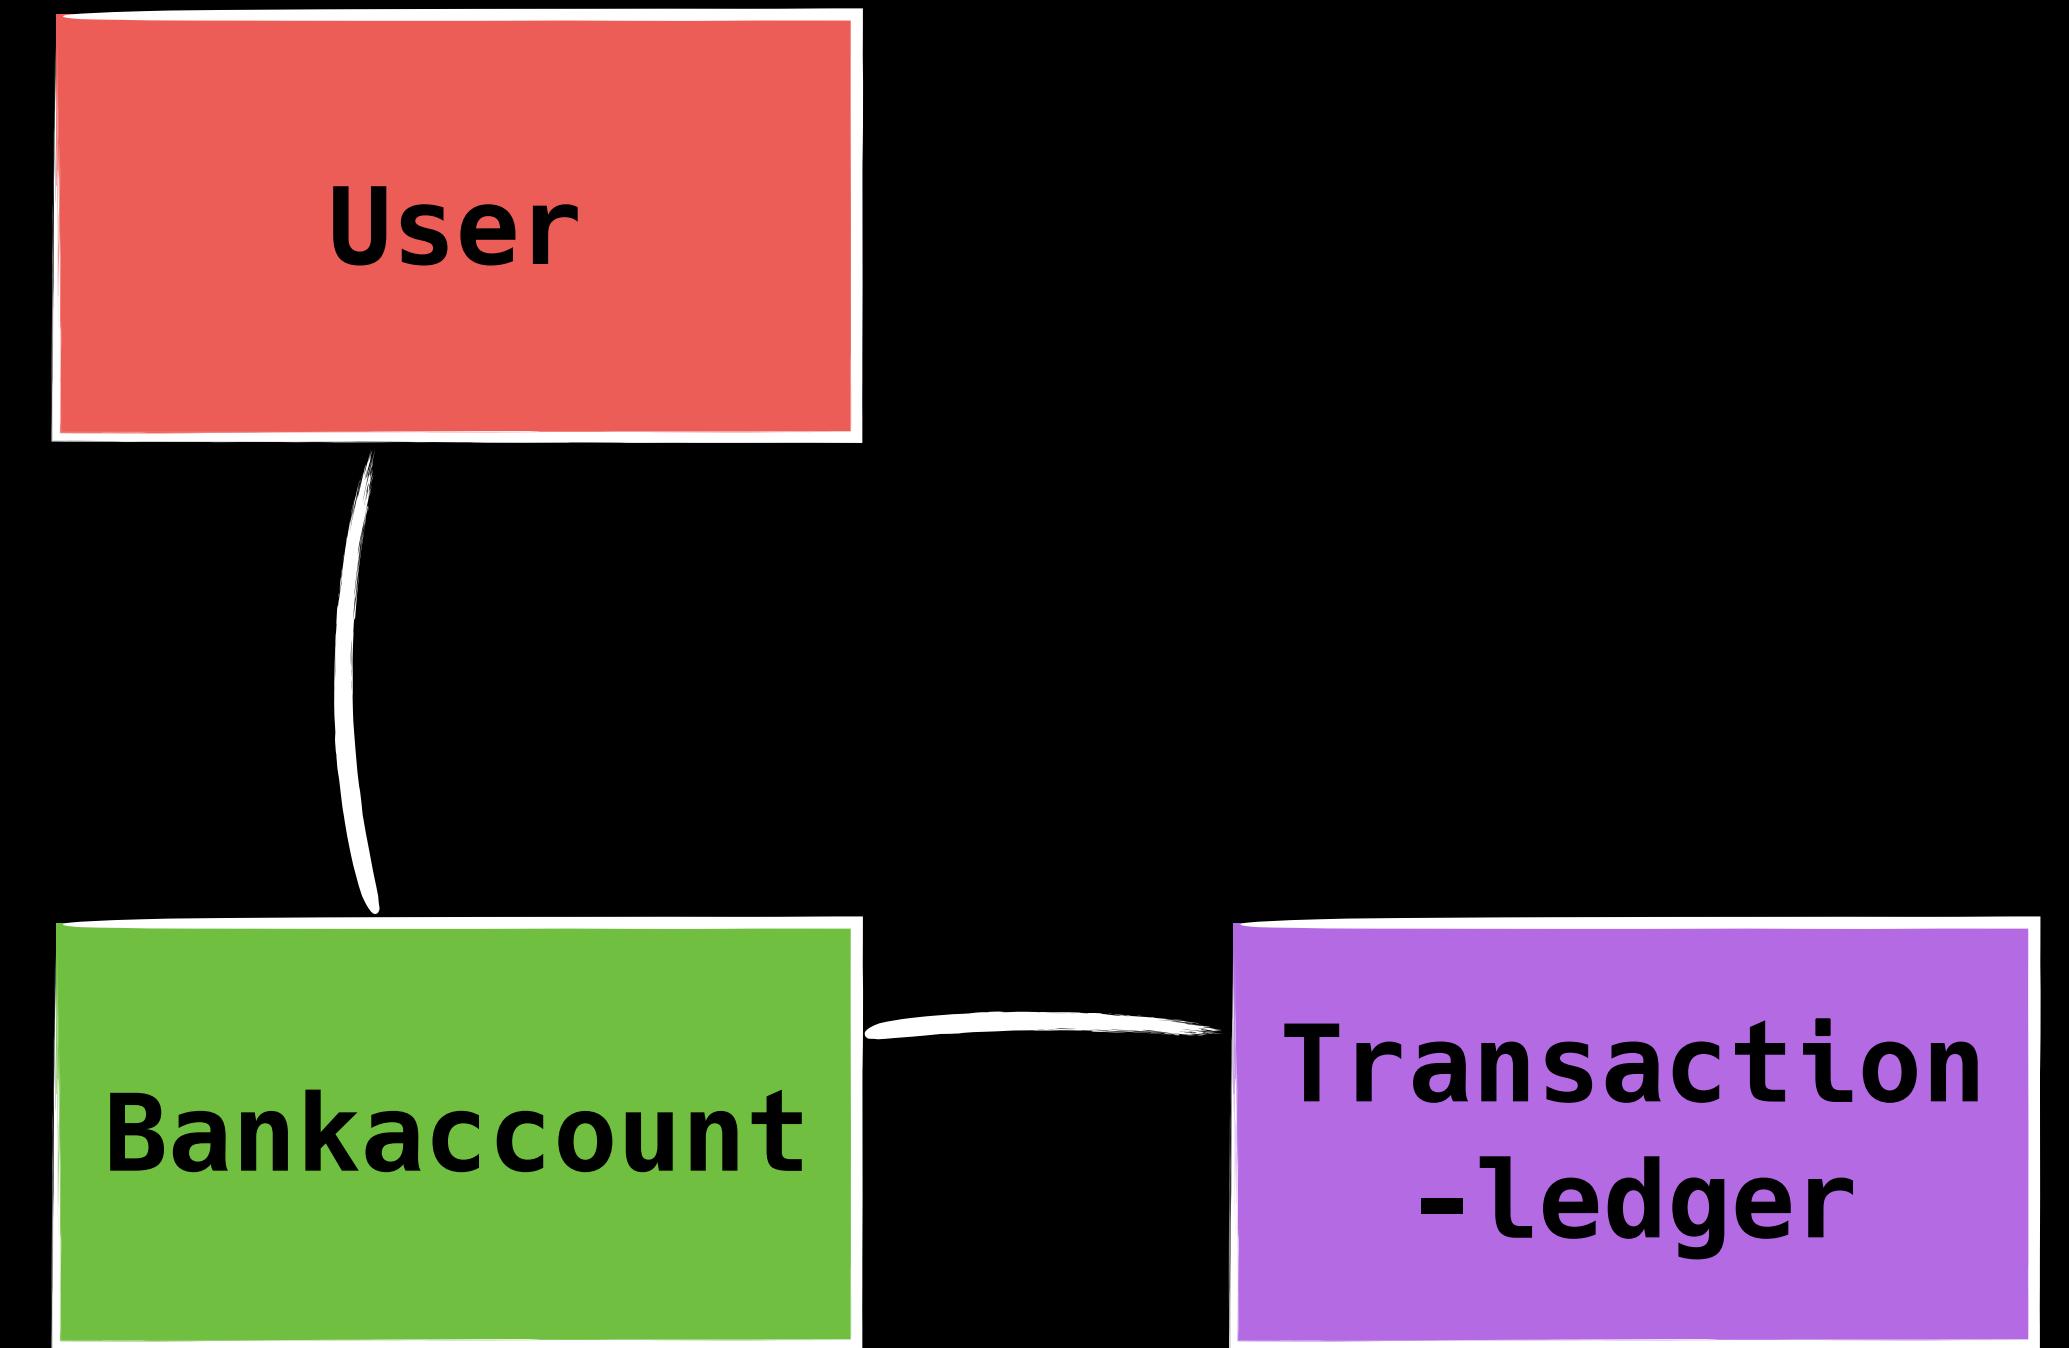
“Please delete all my data”

**Deletion is effectively deleting
the stream-specific key**



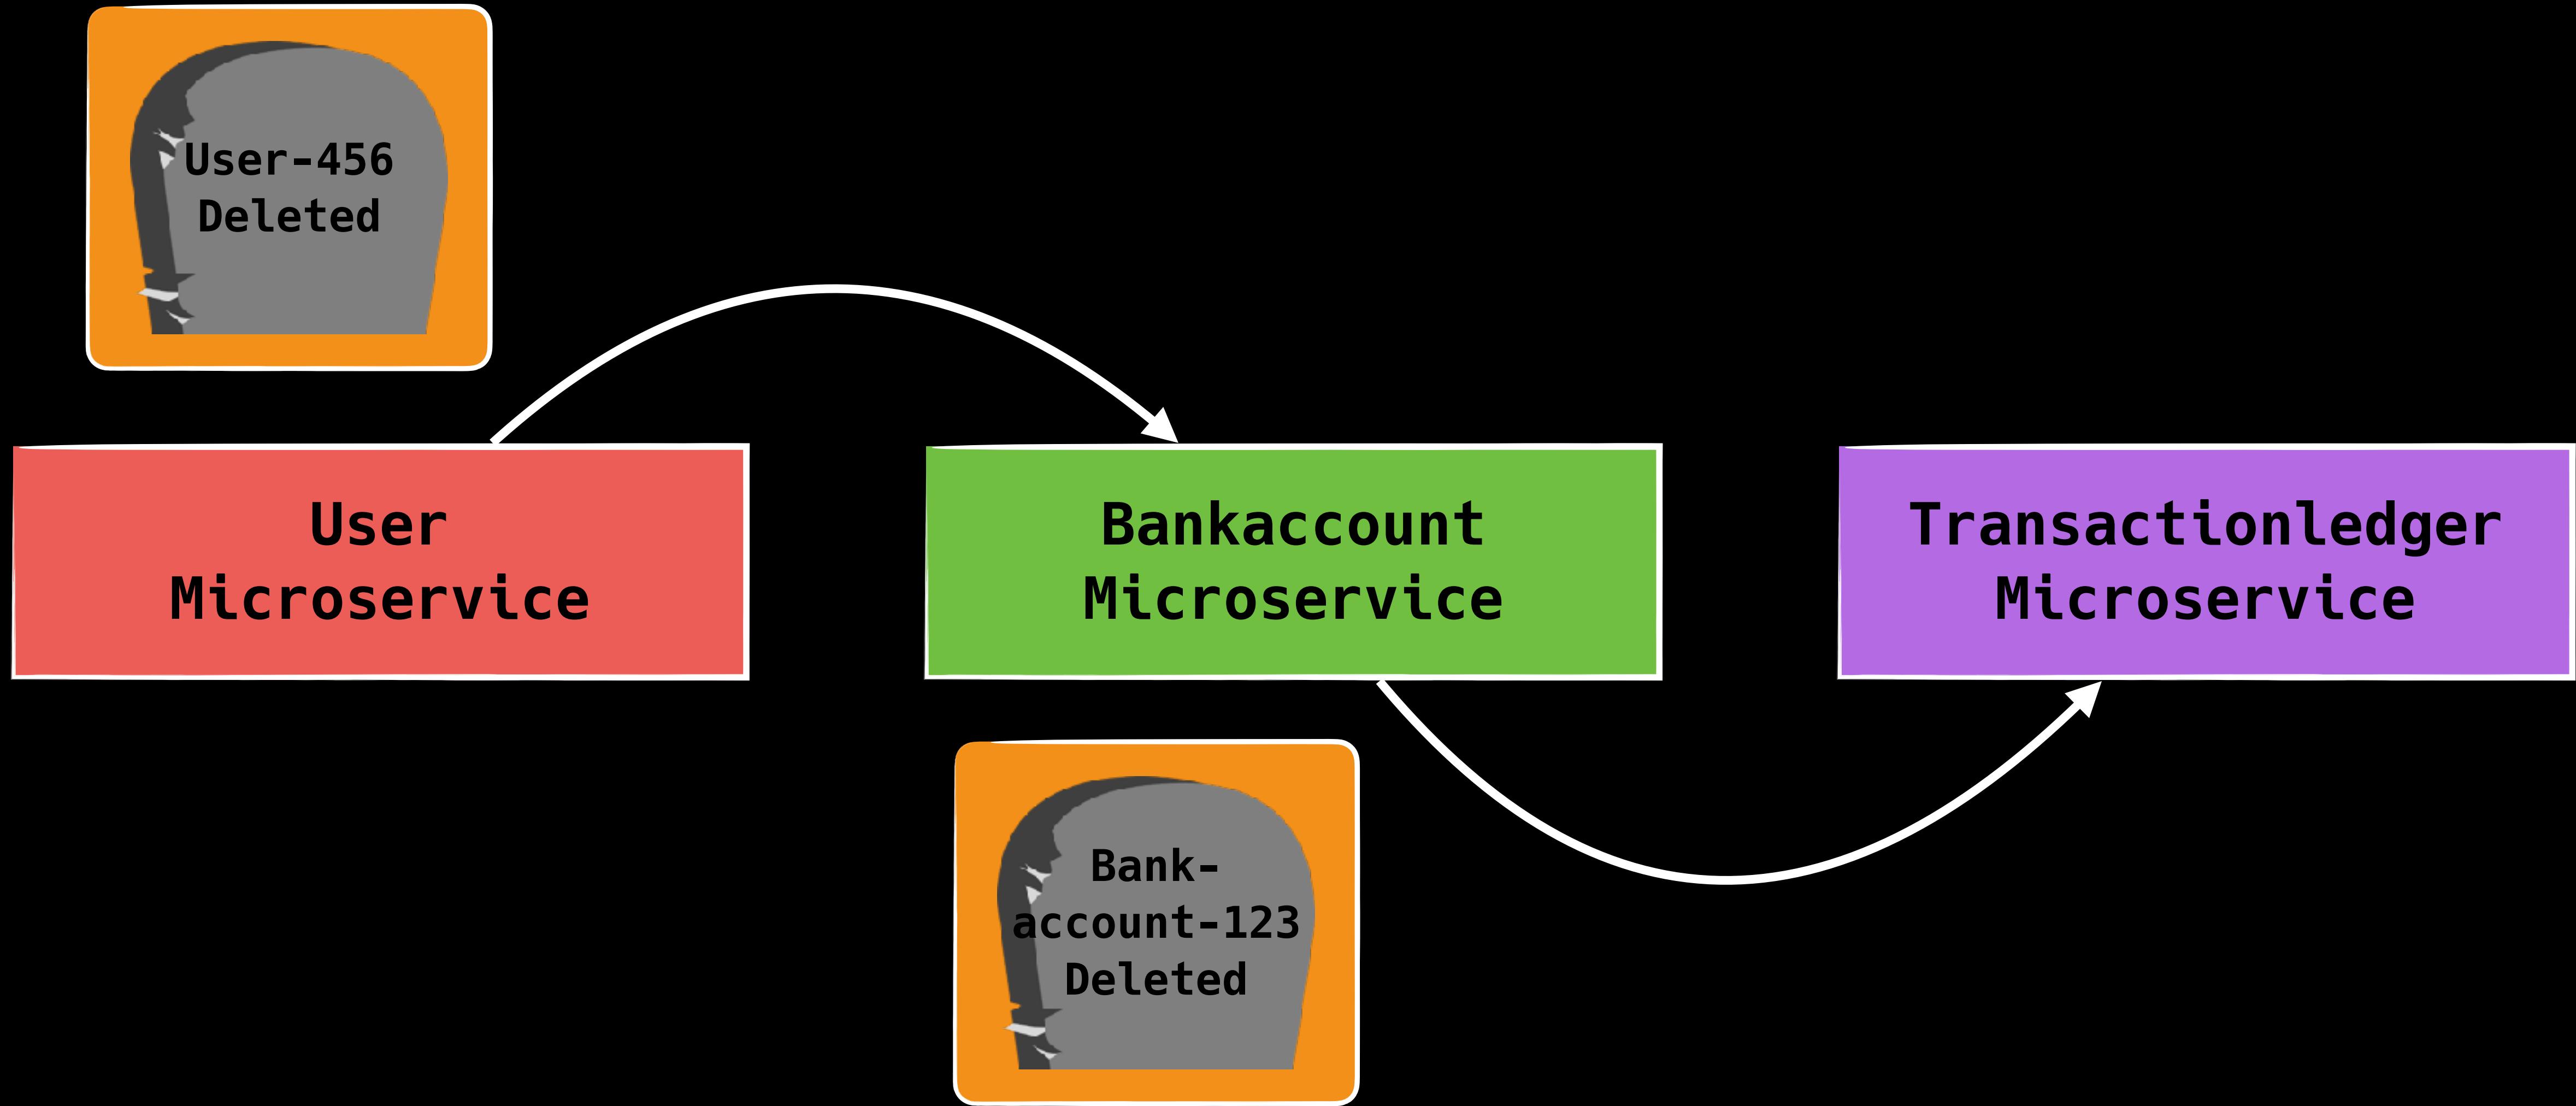
Key administration
Finding what needs to be deleted
Storage implications
Coding complexity
Dashboards, Monitoring

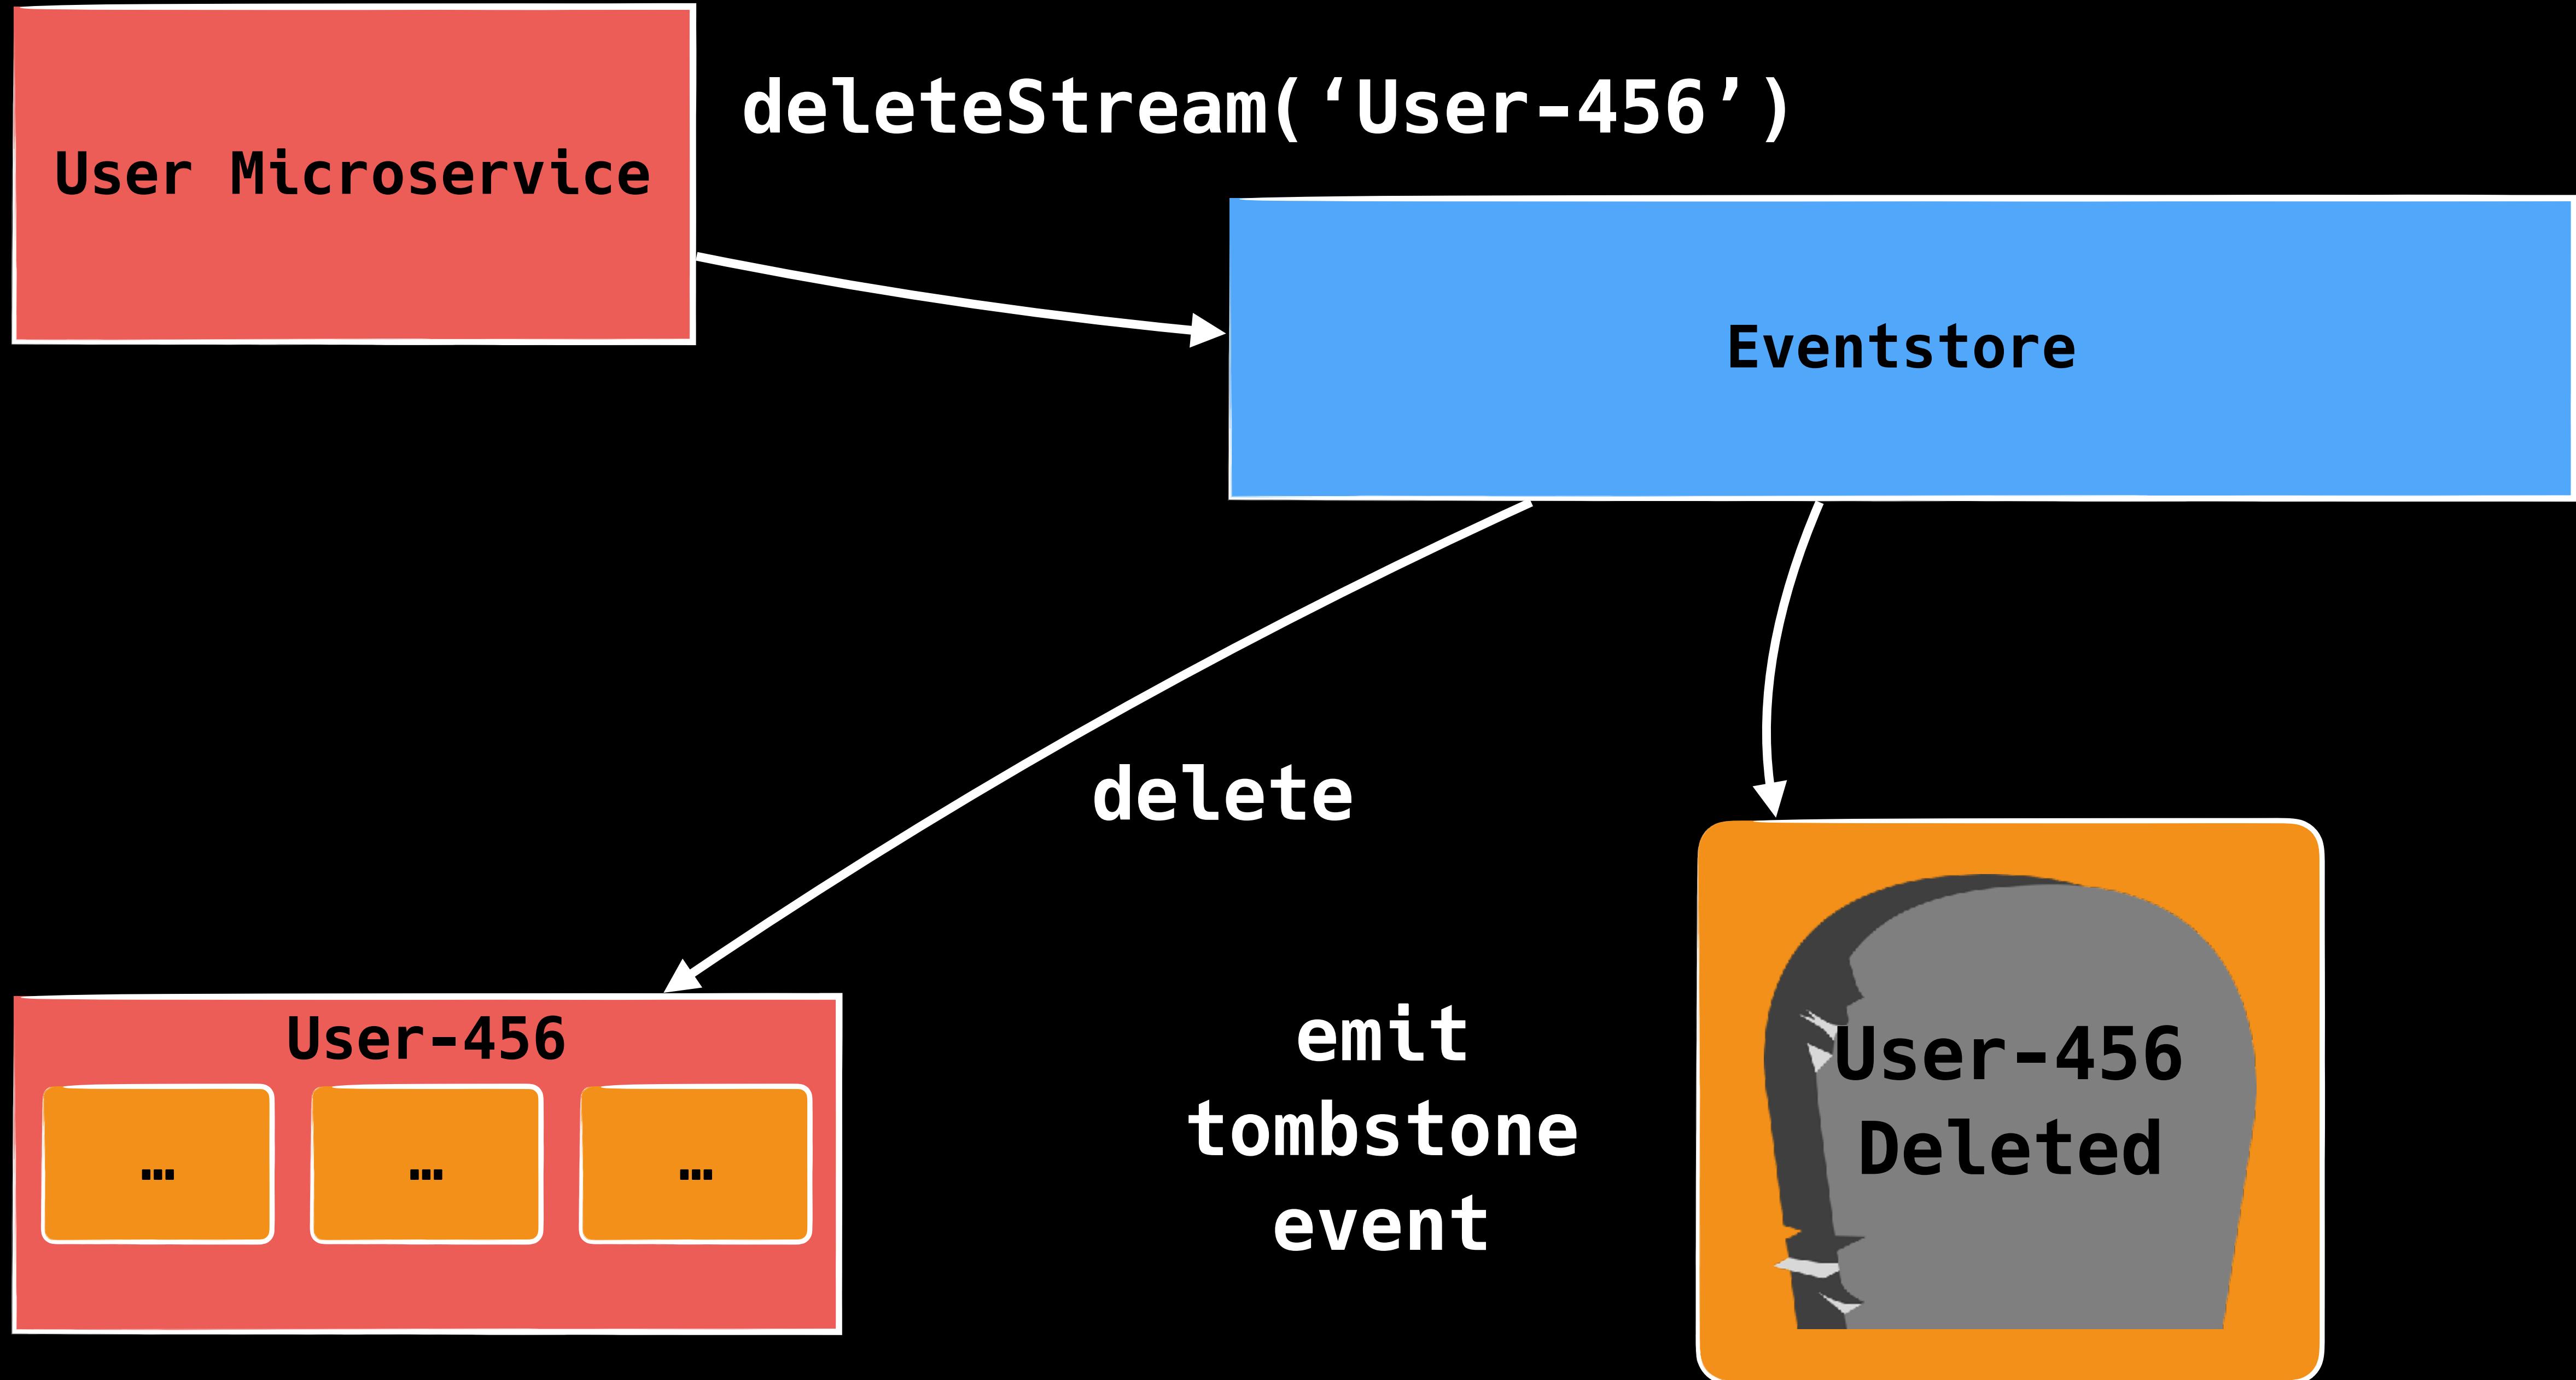
Being able to delete is awesome

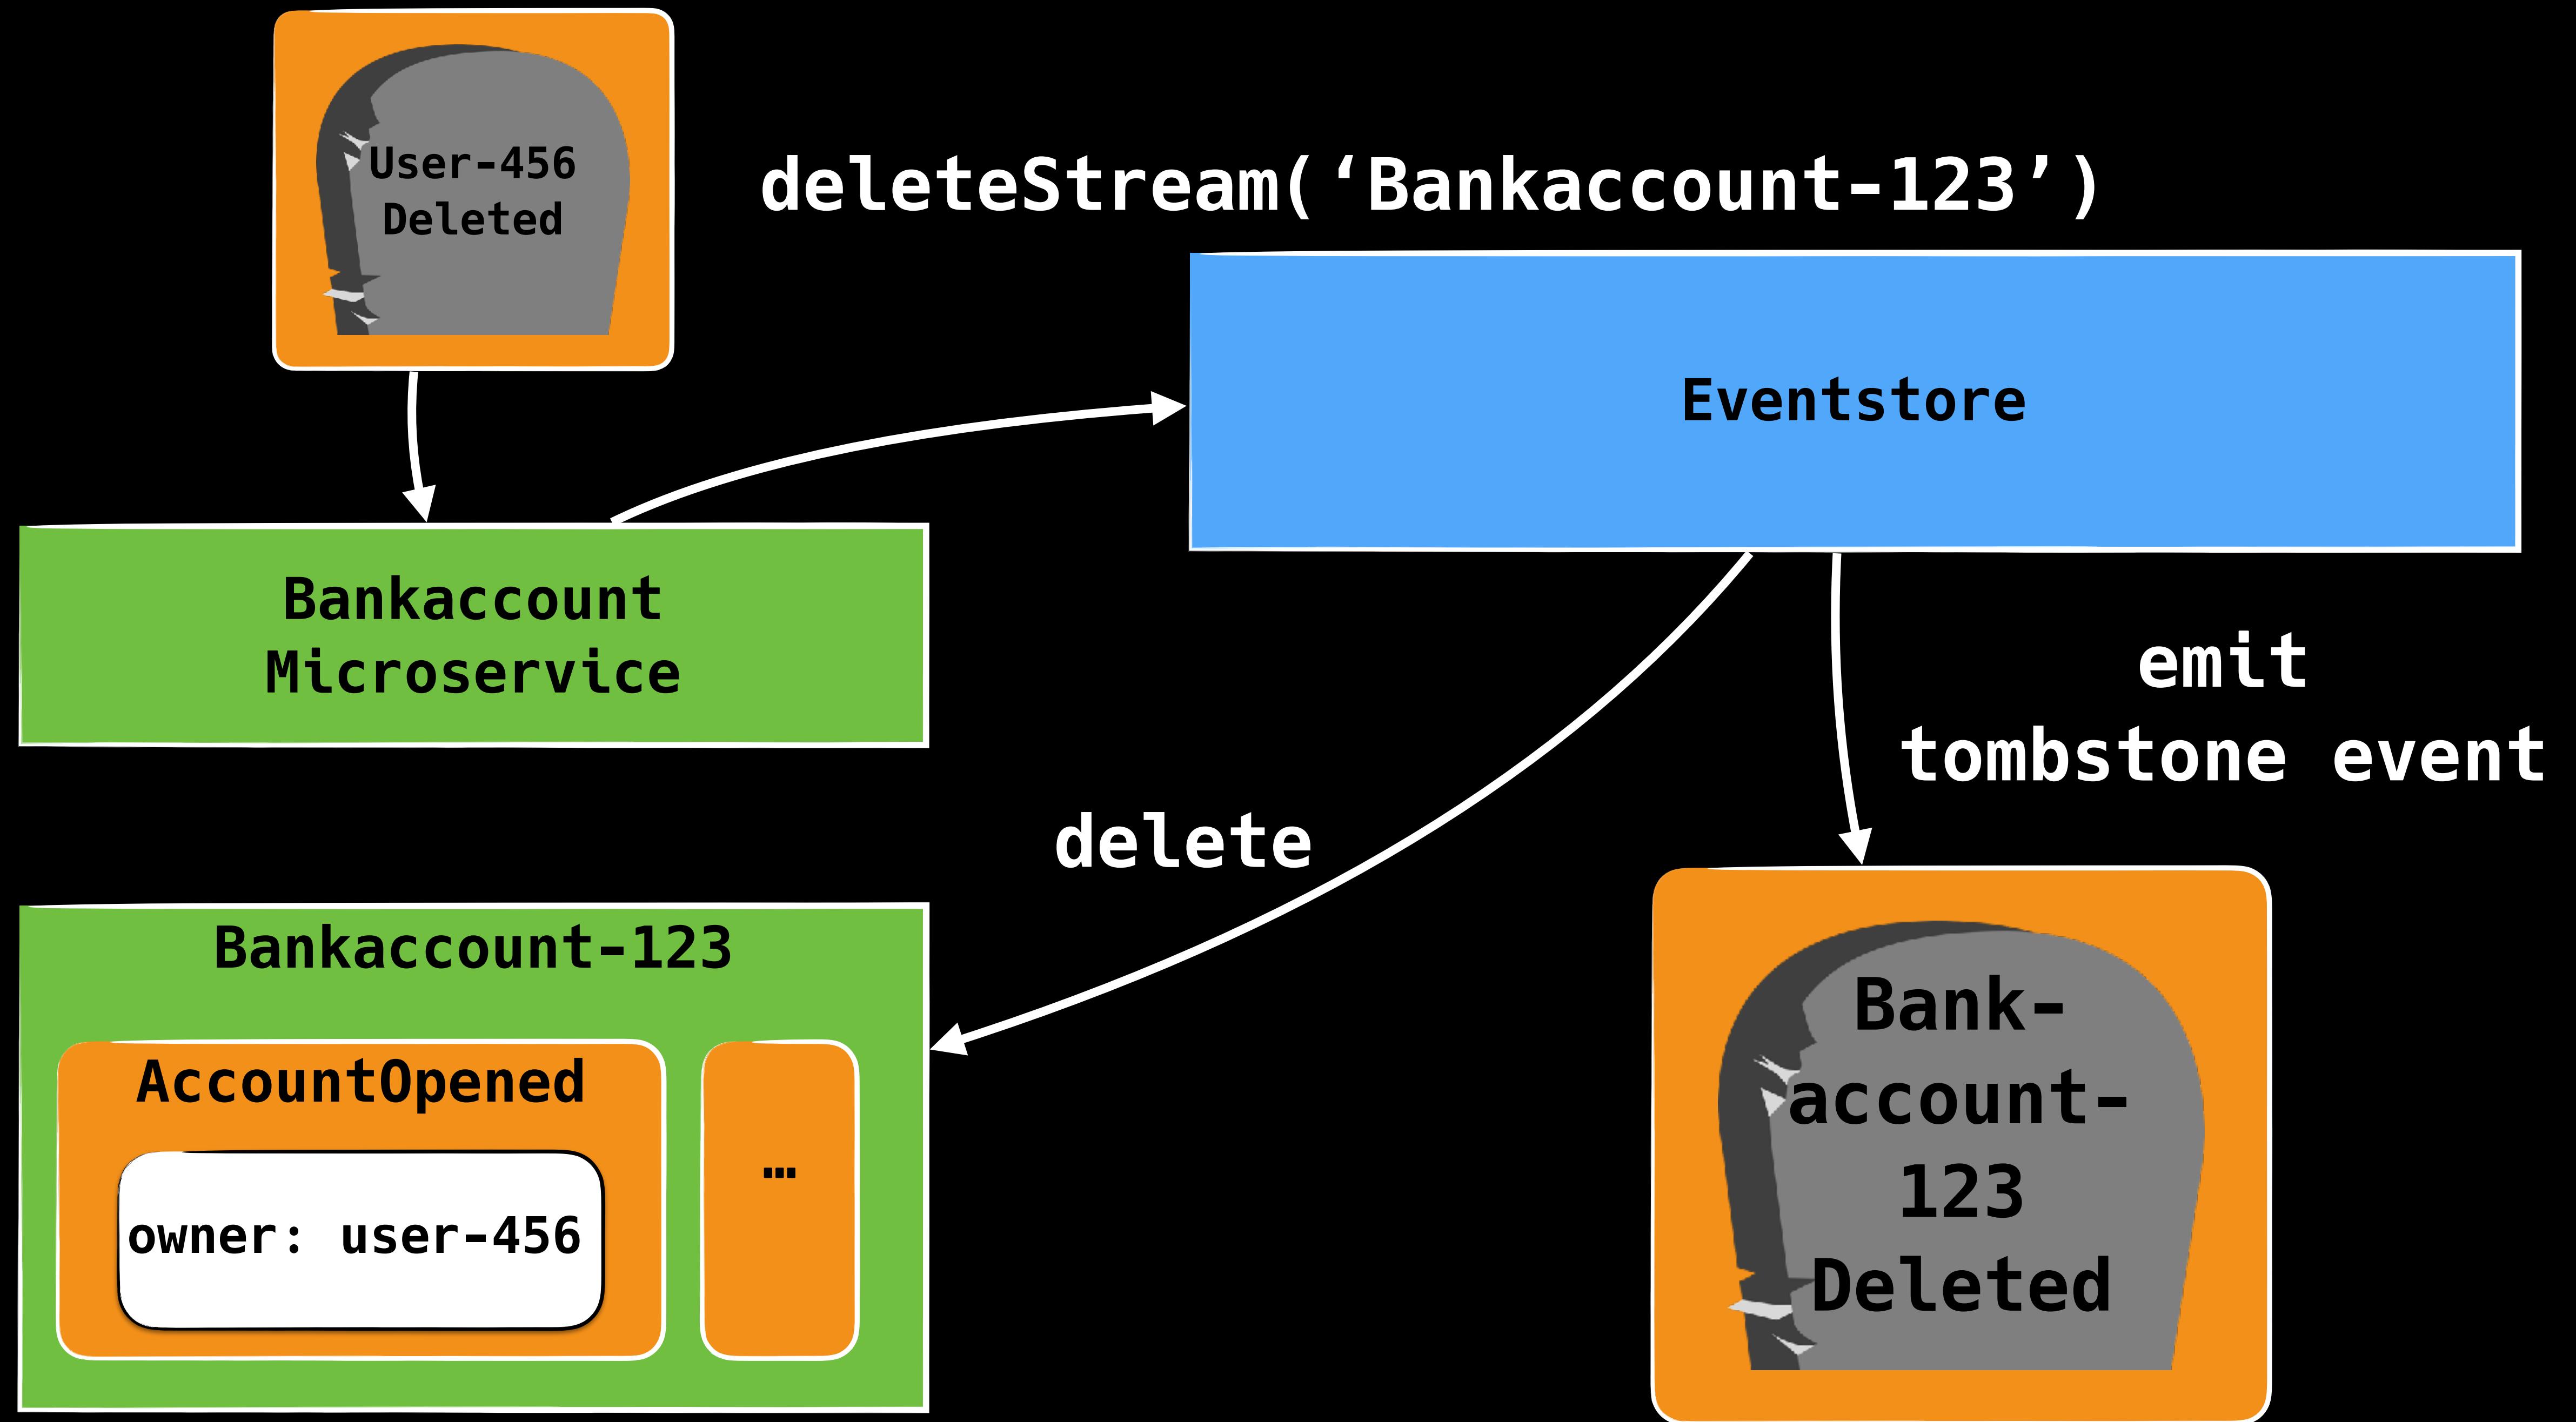


“Please delete all my data”

Cascading deletes with tombstones









Dealing with dependent events

Public/private data

User-Public-456

Id

Blue

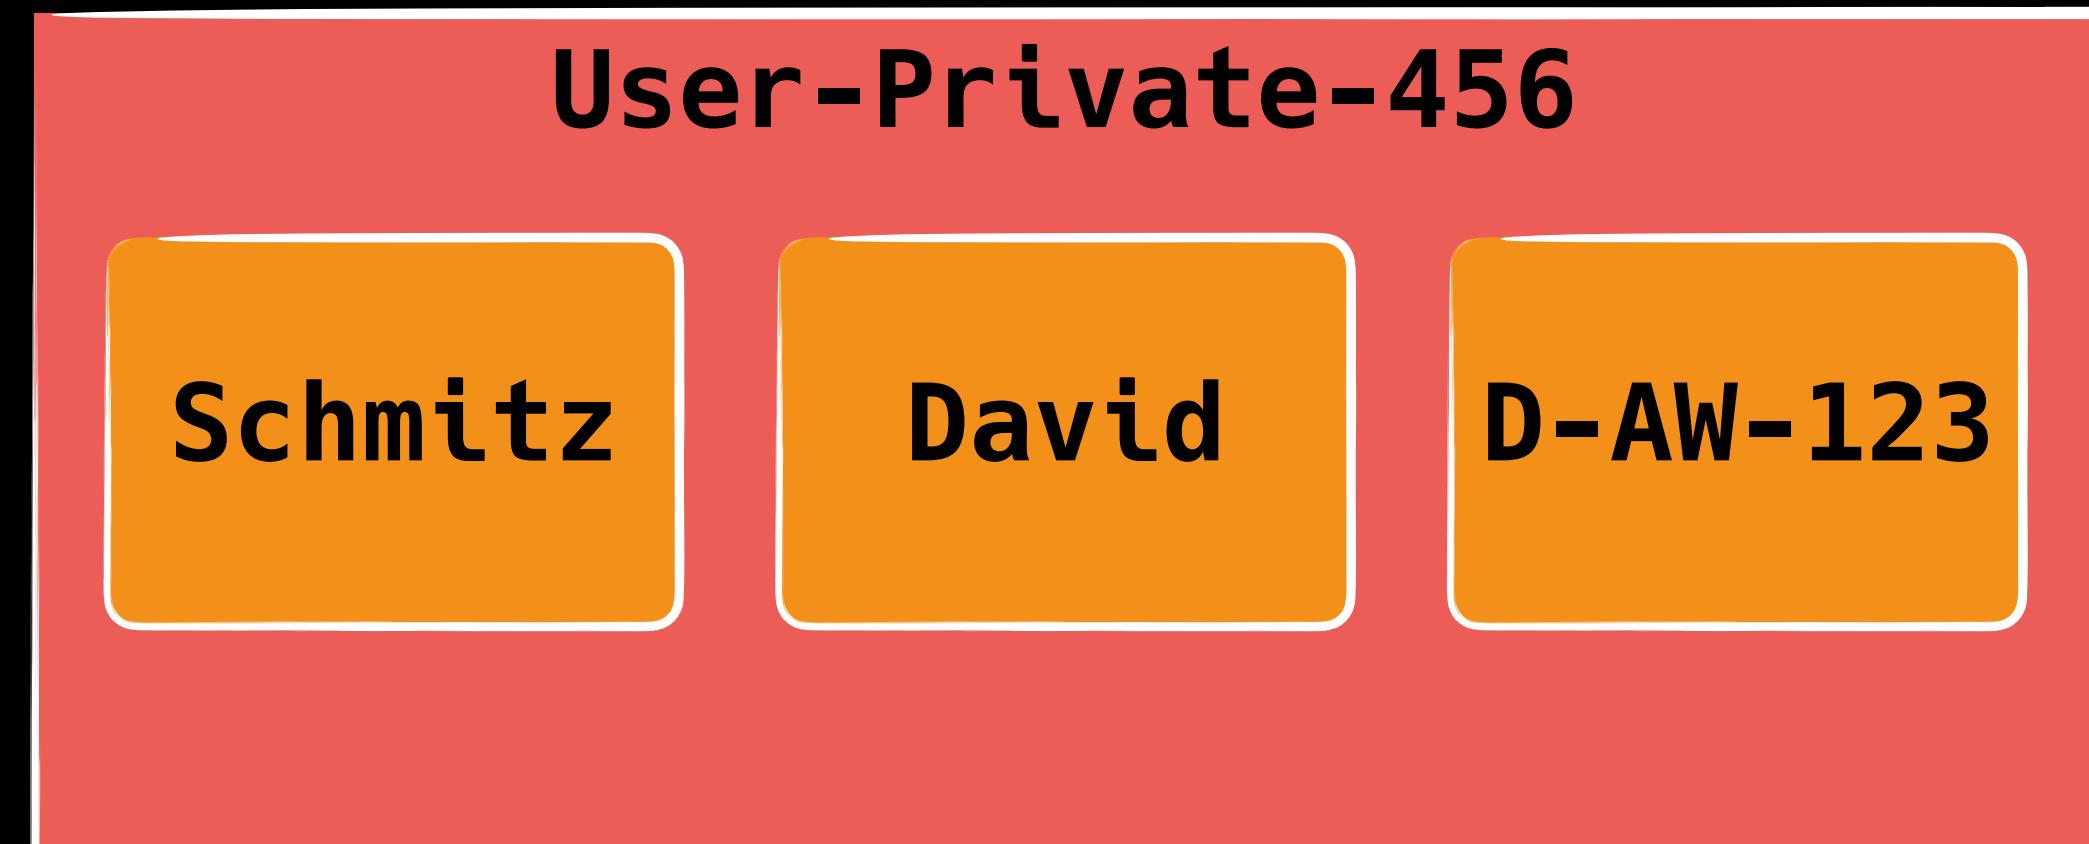
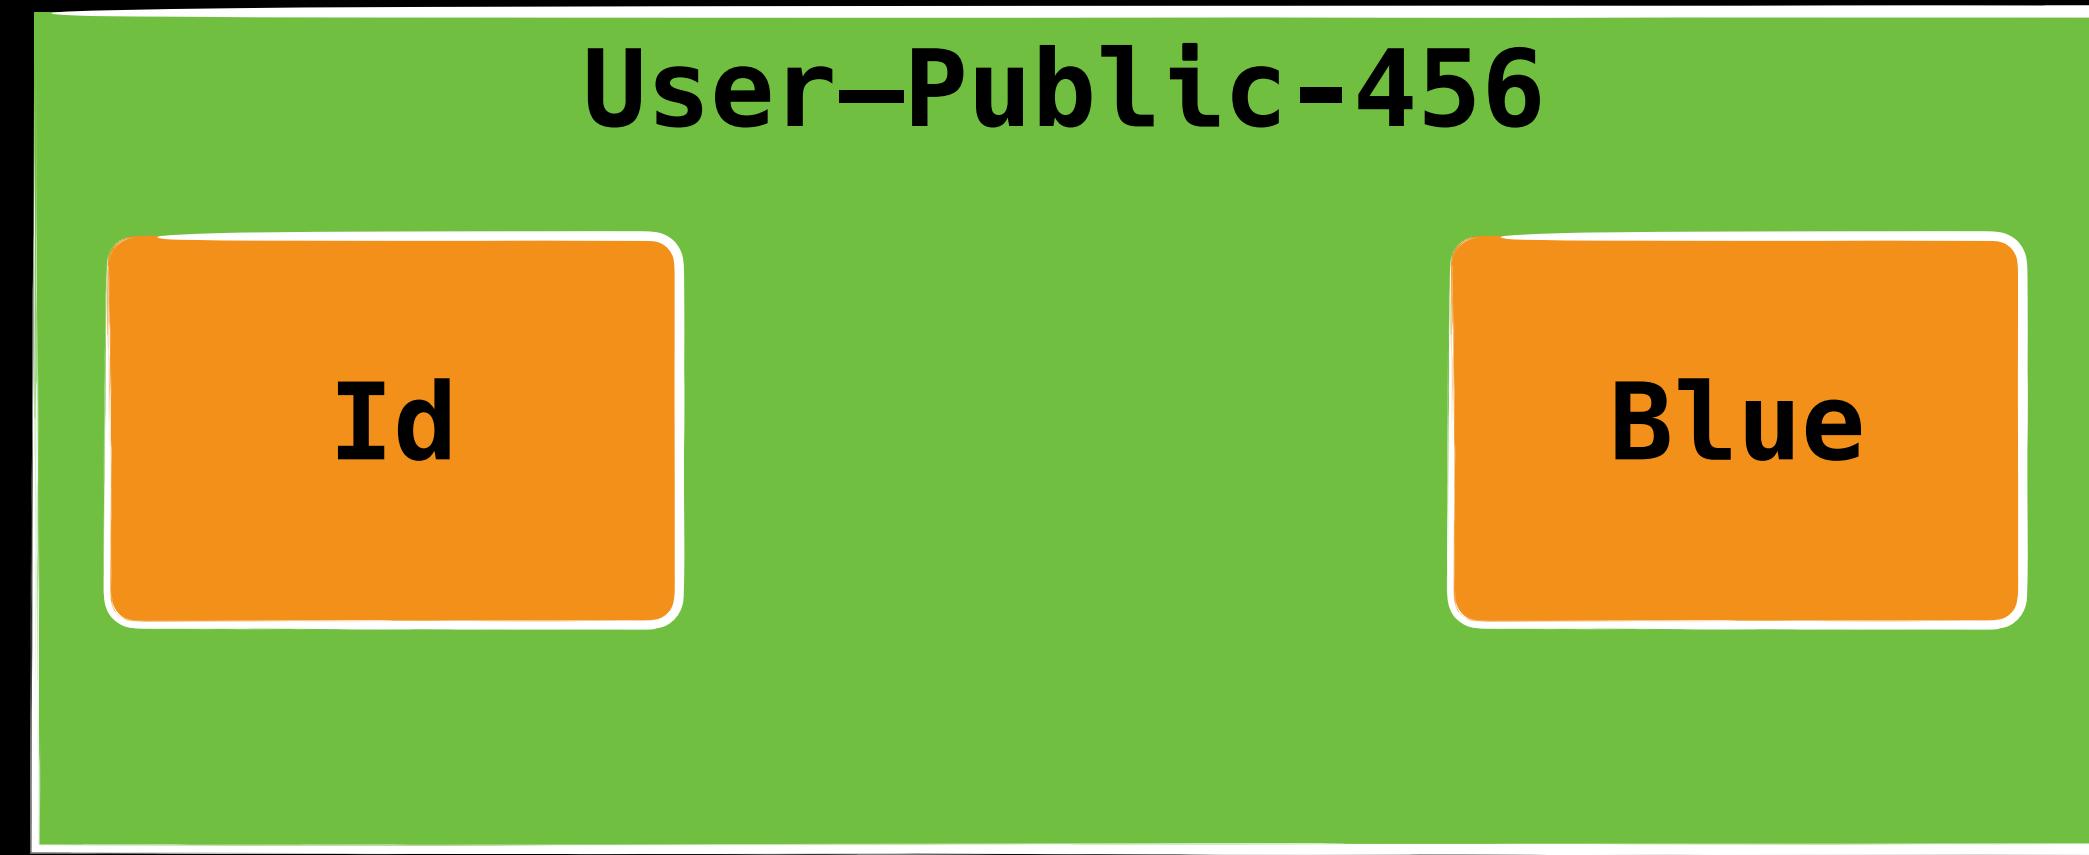
User-Private-456

Schmitz

David

D-AW-123

Keep this →



User-Public-456

Id

Blue

Delete this ➔

User-Private-456

Schmitz

David

D-AW-123

You **may** be able to keep
references to the public data

**“Just” anonymise the
data**





Recital 26 EU GDPR

(26) The principles of data protection should apply to any information concerning an identified or identifiable natural person.

Personal data which have undergone pseudonymisation, which could be attributed to a natural person by the use of additional information should be considered to be information on an identifiable natural person.





Recital 26 EU GDPR

(26) The principles of data protection should apply to any information concerning an identified or identifiable natural person.

Personal data which have undergone pseudonymisation, which could be attributed to a natural person by the use of additional information should be considered to be information on an identifiable natural person.



Surprise: No **easy** answers

Ask your lawyer or CISO

That's it?



ES + DDD = ❤

Needs more up-front design

You can refactor, you can clean up

Not enough in-depth books

Avoid frameworks

Beware: “just...” or “...made easy”

Forget this talk... read these:

The Dark Side of Event Sourcing: Managing Data Conversion

Michiel Overeem¹, Marten Spoor¹, and Slinger Jansen²

Effective Aggregate Design Part I: Modeling a Single Aggregate

Vaughn Vernon: vvernon@shiftmethod.com

Versioning in an Event Sourced System

Gregory Young



EXPLORING CQRS AND EVENT SOURCING

A journey into high scalability, availability,
and maintainability with Windows Azure

Dominic Betts
Julián Domínguez
Grigori Melnik
Fernando Simonazzi
Mani Subramanian

Foreword by Greg Young



patterns & practices

Choose the right tool?



[Community](#) [Support](#) [Blog](#) [Documentation](#) [Downloads](#)

The open-source, functional database with
Complex Event Processing in JavaScript.



Thank you!

Questions?

Comments?

Blame?

@Koenighotze

